

# Reinforcement Learning

## Chapter 3: Model-free RL

Ali Bereyhi

[ali.bereyhi@utoronto.ca](mailto:ali.bereyhi@utoronto.ca)

Department of Electrical and Computer Engineering  
University of Toronto

Fall 2025

# Control Loop via Temporal Difference

- + But still we are not fully **online**! We need to wait till **end of each episode!**
- Well! That's right! But, we could use TD!

Using TD in the control loop will make our algorithm **fully online**

- We update values after each **state-action** pair
- We then **improve** the policy

We should yet use  **$\epsilon$ -greedy improvement** to keep exploration

# SARSA: State-Action-Reward State-Action

**SARSA**  $\equiv$  State-Action Reward State-Action

SARSA algorithms use TD along with  $\epsilon$ -greedy update for the control loop

In general, we can develop various forms of SARSA

- We may use TD-0 for updating **action-values**
  - ↳ This is the basic SARSA
- We may use TD- $n$  for updating **action-values**
  - ↳ This is  $n$ -SARSA
- We may use  $TD_\lambda$  for updating **action-values**
  - ↳ This is SARSA( $\lambda$ )

# SARSA: First Try

Let's try to make a simple TD-based control loop

`TD_Control()`:

- 1: Initiate estimator as  $\hat{q}_\pi(s, a) = 0$  for all states and actions
- 2: **for** episode = 1 :  $K$  or until  $\pi$  stops changing **do**
- 3:   Initiate with a random state-action pair ( $S_0, A_0$ )
- 4:   **for**  $t = 0 : T - 1$  that is either terminal or terminated **do**
- 5:     Act  $A_t$  and observe  

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$$
- 6:     Update policy to  $\pi \leftarrow \epsilon\text{-Greedy}(\hat{q}_\pi)$
- 7:     Draw the new action  $A_{t+1}$  from  $\pi(\cdot | S_{t+1})$
- 8:     Compute  $\hat{v}_\pi(S_{t+1})$  from  $\hat{q}_\pi(S_{t+1}, a)$  and  $\pi(\cdot | S_{t+1})$
- 9:     Set  $G \leftarrow R_{t+1} + \gamma \hat{v}_\pi(S_{t+1})$
- 10:    Update  $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$
- 11:   **end for**
- 12: **end for**

## SARSA: Going On-Policy

In line 8 of our control algorithm: we compute  $\hat{v}_\pi(S_{t+1})$  as

$$\hat{v}_\pi(S_{t+1}) = \sum_{m=1}^M \pi(a^m | S_{t+1}) \hat{q}_\pi(S_{t+1}, a^m)$$

But, we do know that

- ① our estimates  $\hat{q}_\pi(S_{t+1}, a^m)$  are not that good, and also
- ② our policy has led us to next action  $A_{t+1}$

So, we could move on our policy and write

$$\pi(a | S_{t+1}) = \begin{cases} 1 & a = A_{t+1} \\ 0 & a \neq A_{t+1} \end{cases} \rightsquigarrow \hat{v}_\pi(S_{t+1}) = \hat{q}_\pi(S_{t+1}, A_{t+1})$$

We call this approach **on-policy**, since move on our policy

# SARSA: Basic Algorithm

SARSA() :

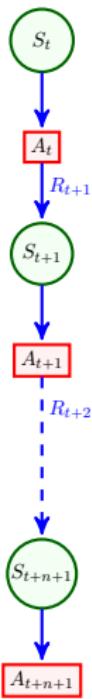
- 1: Initiate estimator as  $\hat{q}_\pi(s, a) = 0$  for all states and actions
- 2: for episode = 1 :  $K$  or until  $\pi$  stops changing do
- 3:   Initiate with a random state-action pair  $(S_0, A_0)$
- 4:   for  $t = 0 : T - 1$  that is either terminal or terminated do
- 5:     Act  $A_t$  and observe  

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$$
- 6:     Update policy to  $\pi \leftarrow \epsilon$ -Greedy( $\hat{q}_\pi$ )
- 7:     Draw the new action  $A_{t+1}$  from  $\pi(\cdot | S_{t+1})$  and move on policy  

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1}$$
- 8:     Set  $G \leftarrow R_{t+1} + \gamma \hat{q}_\pi(S_{t+1}, A_{t+1})$
- 9:     Update  $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$
- 10:   end for
- 11: end for

## SARSA: Deeper Return Samples

We can use a longer trajectory while we learn on-policy, i.e.,



$$G^n = \sum_{i=0}^n R_{t+i+1} + \gamma \hat{q}_\pi(S_{t+n+1}, A_{t+n+1})$$

*This will however add extra delay!*

*As a practice, you could*

re-write the basic SARSA with *n*-return ☺

## SARSA( $\lambda$ ): Tracing Eligibility of State-Action Pairs

We can extend SARSA to the case with  $\lambda$ -return: we have two options

- the case with forward-view
  - ↳ We know this is **not** practical! So, let's skip the details
- the case with backward-view and eligibility tracing
  - ↳ Let's look into this one

We first extend **eligibility tracing** to the case with **state-action pairs**

**ElgTrace**( $S_t, A_t, E(\cdot) | \lambda$ ):

- 1: Eligibility tracing function has  $N M$  components, i.e.,  $E(s, a)$  for all state-action pairs
- 2: **for** all state-action pairs ( $s, a$ ) **do**
- 3:   Update  $E(s, a) \leftarrow \gamma \lambda E(s, a)$
- 4: **end for**
- 5: Update  $E(S_t, A_t) \leftarrow E(S_t, A_t) + 1$

# SARSA: Alternative via TD- $\lambda$

SARSA( $\lambda$ ):

- 1: Initiate  $\hat{q}_\pi(s, a) = 0$  and  $E(s, a) = 0$  for **all states** and **actions**
- 2: **for** episode = 1 :  $K$  or until  $\pi$  stops changing **do**
- 3:   Initiate with a **random state-action pair** ( $S_0, A_0$ )
- 4:   **for**  $t = 0 : T - 1$  that is either **terminal** or **terminated** **do**
- 5:      $E(\cdot) \leftarrow \text{ElgTrace}(S_t, A_t, E(\cdot) | \lambda)$
- 6:     Act  $A_t$  and observe  $R_{t+1}$  and  $S_{t+1}$
- 7:     Update policy to  $\pi \leftarrow \epsilon\text{-Greedy}(\hat{q}_\pi)$
- 8:     Draw the new action  $A_{t+1}$  from  $\pi(\cdot | S_{t+1})$  and move on policy

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1}$$

- 9:     Set  $\Delta \leftarrow R_{t+1} + \gamma \hat{q}_\pi(S_{t+1}, A_{t+1}) - \hat{q}_\pi(S_t, A_t)$
- 10:    **for** **all state-action pairs** ( $s, a$ ) **do**
- 11:      Update  $\hat{q}_\pi(s, a) \leftarrow \hat{q}_\pi(s, a) + \alpha \Delta E(s, a)$
- 12:    **end for**
- 13:   **end for**
- 14: **end for**

## Going Off-Policy

Let's think about a fundamental question: while sampling the environment with a specific policy  $\pi$ , can we estimate the values of another policy  $\bar{\pi}$ ?

- + Why should this be a **fundamental** question?
- Well! There are several reasons
  - ↳ Maybe we sampled **environment** with our **bad** policy: can't we use our sample again?
  - ↳ Maybe we are **looking at other players**: can't we learn something about the environment from their samples?
    - ↳ Maybe they are **good** players: can't we use this fact to improve our policy?
    - ↳ Maybe they are **bad** players: can't we use this fact to avoid doing mistakes?

---

This is the idea of **off-policy control**

Let's start with some basics

# Importance Sampling

Consider following problem: we have random variable  $X$  drawn as  $\textcolor{blue}{X} \sim p(x)$  whose mean is

$$\mu_p = \mathbb{E}_p \{X\} = \sum_x p(x) x$$

We want to know how would be the expectation if we had  $\textcolor{red}{X} \sim q(x)$ : we write

$$\begin{aligned}\mu_q &= \mathbb{E}_q \{X\} = \sum_x q(x) x \\ &= \sum_x \textcolor{blue}{p}(x) \frac{q(x)}{p(x)} x = \mathbb{E}_{\textcolor{blue}{p}} \left\{ \frac{q(X)}{p(X)} X \right\}\end{aligned}$$

This gives us possibility to

estimate  $\mathbb{E}_q \{X\}$  using samples drawn from  $p(x)$

# Importance Sampling

Say we have drawn  $K$  samples from  $p(x)$ , i.e., we have

$$X_1, X_2, \dots, X_K$$

We can use Monte-Carlo to estimate  $\mu_p$  as

$$\hat{\mu}_p = \frac{1}{K} \sum_{k=1}^K X_k$$

We can also use Monte-Carlo to estimate  $\mu_q$  as

$$\hat{\mu}_q = \frac{1}{K} \sum_{k=1}^K \frac{q(X_k)}{p(X_k)} X_k$$

We call this method importance sampling

## Off-Policy Control via Importance Sampling

Now, let's get back to our problem: *assume we have played with policy  $\pi$  and collected  $K$  sample trajectories of length  $T$  all started at state  $S_0 = s$ , i.e.,*

$$s = S_0[k], A_0[k] \xrightarrow{R_1[k]} S_1[k], A_1[k] \xrightarrow{R_2[k]} \dots \xrightarrow{R_T[k]} S_T[k]$$

for  $k = 1 : K$ ; then, we could write

$$\hat{v}_\pi(s) = \frac{1}{K} \sum_{k=1}^K G[k]$$

This is the basic Monte-Carlo

# Off-Policy Control via Importance Sampling

But now, we want to use samples to evaluate another policy  $\bar{\pi}$

$$s = S_0[k], A_0[k] \xrightarrow{R_1[k]} S_1[k], A_1[k] \xrightarrow{R_2[k]} \dots \xrightarrow{R_T[k]} S_T[k]$$

We could also use importance sampling to write

$$\begin{aligned}\hat{v}_{\bar{\pi}}(s) &= \frac{1}{K} \sum_{k=1}^K \frac{\Pr\{\text{same action sequence with } \bar{\pi}\}}{\Pr\{\text{same action sequence with } \pi\}} G[k] \\ &= \frac{1}{K} \sum_{k=1}^K \frac{\bar{\pi}(A_0[k]|S_0[k]) \cdots \bar{\pi}(A_{T-1}[k]|S_{T-1}[k])}{\pi(A_0[k]|S_0[k]) \cdots \pi(A_{T-1}[k]|S_{T-1}[k])} G[k] \\ &= \frac{1}{K} \sum_{k=1}^K \prod_{\ell=0}^{T-1} \frac{\bar{\pi}(A_\ell[k]|S_\ell[k])}{\pi(A_\ell[k]|S_\ell[k])} G[k]\end{aligned}$$

# Off-Policy Control via Importance Sampling

We can further update the estimate in an *online fashion* from

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1} \xrightarrow{R_{t+2}} \dots \xrightarrow{R_T} S_T$$

by *online averaging* as

$$\hat{v}_{\bar{\pi}}(S_t) \leftarrow \hat{v}_{\bar{\pi}}(S_t) + \alpha \left( \prod_{\ell=t}^{T-1} \frac{\bar{\pi}(A_\ell | S_\ell)}{\pi(A_\ell | S_\ell)} G_t - \hat{v}_{\bar{\pi}}(S_t) \right)$$

So, we are evaluating  $\bar{\pi}$  via *Monte-Carlo*

*off* our policy  $\pi$

This is *off-policy control*

# Off-Policy Control via Importance Sampling

We can further apply off-policy control via TD

$$\hat{v}_{\bar{\pi}}(S_t) \leftarrow \hat{v}_{\bar{\pi}}(S_t) + \alpha \left( \frac{\bar{\pi}(A_t|S_t)}{\pi(A_t|S_t)} (R_{t+1} + \gamma \hat{v}_{\bar{\pi}}(S_{t+1})) - \hat{v}_{\bar{\pi}}(S_t) \right)$$

Note that for action-values estimate

$R_{t+1}$  does not depend any more on policy as we know action  $A_t$

Therefore, we have for action-value update

$$\hat{q}_{\bar{\pi}}(S_t, A_t) \leftarrow \hat{q}_{\bar{\pi}}(S_t, A_t) + \alpha \left( R_{t+1} + \gamma \frac{\bar{\pi}(A_t|S_t)}{\pi(A_t|S_t)} \hat{v}_{\bar{\pi}}(S_{t+1}) - \hat{q}_{\bar{\pi}}(S_t, A_t) \right)$$

# Q-Learning

## Q-Learning

Q-learning is an **off-policy TD control** algorithm, where we sample with  $\epsilon$ -greedy policy but update the action-values to evaluate greedy policy

This means in Q-learning  $\pi$  is  $\epsilon$ -greedy policy and  $\bar{\pi}$  is greedy. Let's consider basic TD evaluation: so, we can write

$$\hat{q}_{\bar{\pi}}(S_t, A_t) \leftarrow \hat{q}_{\bar{\pi}}(S_t, A_t) + \alpha(G - \hat{q}_{\bar{\pi}}(S_t, A_t))$$

where  $G$  should be

$$G = R_{t+1} + \gamma \hat{v}_{\bar{\pi}}(\bar{S}_{t+1})$$

Since we sample by  $\epsilon$ -greedy policy  $\pi$ , we use **importance sampling** and write

$$G = R_{t+1} + \gamma \frac{\bar{\pi}(A_t|S_t)}{\pi(A_t|S_t)} \hat{v}_{\bar{\pi}}(S_{t+1})$$

# Q-Learning

But, we really *don't need importance sampling*: we can simply observe that

$$\hat{v}_{\bar{\pi}}(S_{t+1}) = \sum_{m=1}^M \hat{q}_{\bar{\pi}}(S_{t+1}, a^m) \bar{\pi}(a^m | S_{t+1}) = \max_m \hat{q}_{\bar{\pi}}(S_{t+1}, a^m)$$

and we do know that

$$\frac{\bar{\pi}(A_t | S_t)}{\pi(A_t | S_t)} = \mathbf{1} \left\{ A_t = \operatorname{argmax}_a \hat{q}_{\bar{\pi}}(S_t, a) \right\}$$

So, we could directly update as

$$\hat{q}_{\bar{\pi}}(S_t, A_t) \leftarrow \hat{q}_{\bar{\pi}}(S_t, A_t) + \alpha \left( R_{t+1} + \gamma \max_m \hat{q}_{\bar{\pi}}(S_{t+1}, a^m) - \hat{q}_{\bar{\pi}}(S_t, A_t) \right)$$

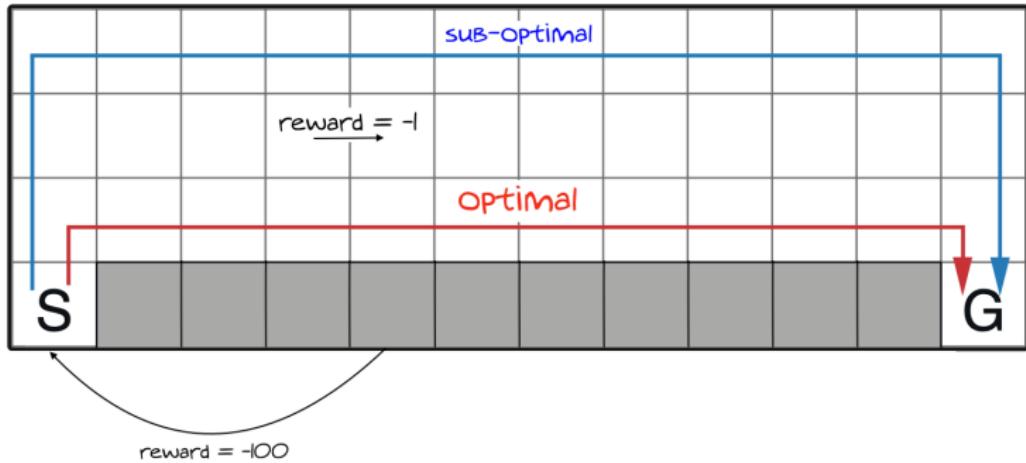
This concludes *Q-learning algorithm*

# Q-Learning: Basic Algorithm

**Q-Learning():**

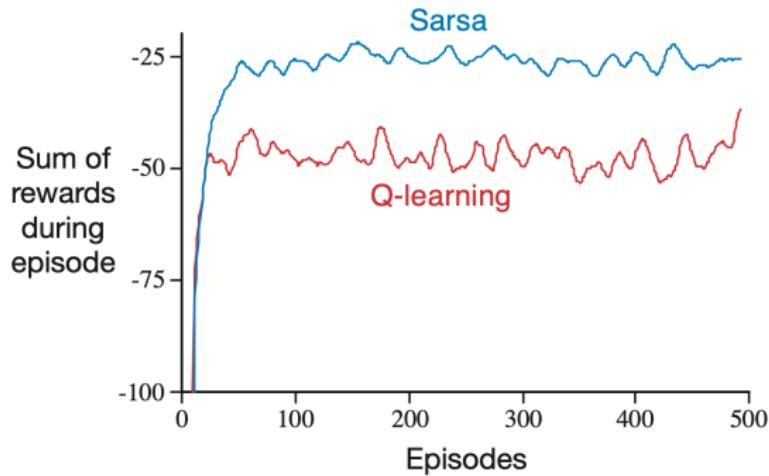
- 1: *Initiate estimator as  $\hat{q}_\star(s, a) = 0$  for all states and actions*
  - 2: **for** episode = 1 :  $K$  or until  $\pi$  stops changing **do**
  - 3:   *Initiate with a random state  $S_0$*
  - 4:   **for**  $t = 0 : T - 1$  that is either terminal or terminated **do**
  - 5:     *Update policy to  $\pi \leftarrow \epsilon\text{-Greedy}(\hat{q}_\star)$*
  - 6:     *Draw action  $A_t$  from  $\pi(\cdot | S_t)$  and observe*
- $$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$$
- 7:     *Set  $G \leftarrow R_{t+1} + \gamma \max_m \hat{q}_\star(S_{t+1}, a^m)$*
  - 8:     *Update  $\hat{q}_\star(S_t, A_t) \leftarrow \hat{q}_\star(S_t, A_t) + \alpha(G - \hat{q}_\star(S_t, A_t))$*
  - 9:   **end for**
  - 10: **end for**

## Example: Cliff Walking



Let's compare SARSA to Q-Learning algorithm!

## Example: Cliff Walking



Don't Mistake!

Q-learning collects less reward since it goes *off-policy*; however, it *estimates optimal* action-values: at some point it can start playing *optimally*

# Convergence of SARSA

## Recall: GLIE Algorithms

A GPI-type control loop is GLIE, if for any **state-action pair** ( $s, a$ ), we have the following asymptotic properties

- ① The number of visits to all **state-action pair** grows large

$$\lim_{K \rightarrow \infty} C_K(s, a) = \infty$$

- ② The **improved** policy in last episode converges to **greedy** policy

$$\lim_{K \rightarrow \infty} \pi_K(a^m | s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\pi_K}(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\pi_K}(s, a^m) \end{cases}$$

GLIE control algorithms converge to **optimal policy**

# Convergence of SARSA

- + But do we really have large number of episodes with SARSA?
- Not necessarily! We may have only one **infinitely long** trajectory
- + What should we do then?
- We can simply treat it as a large number of episodes of length 1

---

In (basic) SARSA, we only need one step in the trajectory

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$$

We could hence think of it as one episode

- ① each time step  $t$  we update the **action-values**
- ② each time step we **improve the policy**

# Convergence of SARSA

## Modification: GLIE Algorithms

An *online control loop is GLIE, if we have asymptotically in time t*

- ① The number of visits to all **state-action** pair grows large

$$\lim_{t \rightarrow \infty} C_t(s, a) = \infty$$

- ② The **improved** policy converges to **greedy** policy

$$\lim_{t \rightarrow \infty} \pi_t(a^m | s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\pi_t}(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\pi_t}(s, a^m) \end{cases}$$

# Convergence of SARSA: Make it GLIE

- + Can we guarantee that **both conditions** hold with SARSA?
- The **second one** is easy: we need to **scale  $\epsilon$  down with  $t$** , e.g.,  $\epsilon_t = 1/t$
- + What about the **first condition**?
- We should scale **the step-size  $\alpha$  according to Robbins-Monro**

## Robbins-Monro Sequence

Sequence  $\alpha_t$  is **Robbins-Monro** if we have

$$\sum_{t=0}^{\infty} \alpha_t = \infty \quad \text{and} \quad \sum_{t=0}^{\infty} \alpha_t^2 < \infty$$

For instance,  $\alpha_t = 1/t$  is a **Robbins-Monro sequence**

# Convergence of SARSA

## Convergence of SARSA

SARSA online control loop converges to the optimal action-values if

- ① Step-size is scheduled by a Robbins-Monro sequence
- ② Exploration factor  $\epsilon$  decays in time

In practice however

- $\epsilon$  is a **hyperparameter**
  - ↳ We know that we should **schedule** it
  - ↳ How we should do the **scheduling**? This is **hyperparameter tuning**
- $\alpha$  is a **hyperparameter**: some people call it **learning rate**
  - ↳ Its **scheduling** is again **hyperparameter tuning**

# Convergence of Q-Learning

## Convergence of Q-Learning

*Q-learning online control loop with exploration (non-zero  $\epsilon$ ) converges to the optimal action-values as  $t \rightarrow \infty$*

- + That's it?
- Yes!

Since we are evaluating off-policy, we don't care about behaving policy

# Q-Learning vs SARSA

- + So! Does it mean that Q-learning is always better?
- Not always!

In general Q-learning has several benefits

- Minimal convergence requirements
- It converges faster to the optimal policy
  - ↳ If we want to make SARSA that fast, we may get to a sub-optimal policy
- It has more flexibility and sample-efficiency

But, SARSA also has some benefits

- It is better suited for online control
  - ↳ Our behaving policy is the one going towards optimal one
  - ↳ In Q-learning, the behaving policy is not the optimal one
- It has lower complexity
  - ↳ We just deal with one policy

# End of Story!

## Model-Based RL

Bellman Equation

value iteration

policy iteration

## Model-free RL

on-policy methods

temporal difference

Monte Carlo

SARSA

off-policy methods

Q-learning

I would strongly suggest to start with **programming part** of Assignment 2!

There you solve Frozen Lake with **SARSA** and **Q-Learning**