

# ECE 1508: Reinforcement Learning

## Chapter 2: Model-based RL

Ali Bereyhi

[ali.bereyhi@utoronto.ca](mailto:ali.bereyhi@utoronto.ca)

Department of Electrical and Computer Engineering  
University of Toronto

Fall 2025

## Last Piece: Dynamic Programming

Right now, we know what to do when we know *MDP of environment*

- ① We can find optimal values from *Bellman optimality equations*
- ② We could then find the *optimal action*-values
- ③ We finally get the *optimal policy* from *optimal action*-values

---

The only remaining challenge is to find

*an algorithmic approach to solve Bellman optimality equations*

We complete this last piece using

*Dynamic Programming*  $\equiv$  DP

# Dynamic Programming: Basic Idea

Assume, we want to solve the problem of

$$x = f(x)$$

for some function  $f(x)$

We could solve it via *direct approach*:

- ① Rewrite it as  $f(x) - x = 0$
- ② Solve it via classic algorithms
  - ↳ Reduce it to a *known form*, e.g., a *polynomial*
  - ↳ Solve it via an *iterative method*, e.g., *Newton-Raphson* or *method of intervals*

# Dynamic Programming: Basic Idea

Assume, we want to solve the problem of

$$x = f(x)$$

for some function  $f(x)$

We could also solve it by **recursion**:

- ① Start with an  $x^0$  and set  $x^1 = f(x^0)$
- ② Until  $x^{k+1} \approx x^k$ , we do
  - ↳ Update **recursively** as  $x^{k+1} = f(x^k)$
  - ↳ Set  $k \leftarrow k + 1$

Under **some conditions on  $f(\cdot)$** , this approach can **converge**

# Dynamic Programming: Example

We want to solve

$$x = \frac{-1}{2+x}$$

- ① Start with an  $x^0 = 0$
- ② We now get into the recursion loop

$$\hookrightarrow x^1 = f(x^0) = -\frac{1}{2}$$

$$\hookrightarrow x^2 = f(x^1) = -\frac{2}{3}$$

$$\hookrightarrow x^3 = f(x^2) = -\frac{3}{4}$$

$\cdots$

$$\hookrightarrow x^k = f(x^{k-1}) = -\frac{k}{k+1}$$

We asymptotically converge to  $x^\infty = -1$  which is the solution

- $\hookrightarrow$  Note that we always converge no matter which point we start

# Dynamic Programming: Example

Now, let's write the *same equation* in a *different recursive form*

$$x = \frac{-1 - x^2}{2}$$

- ① Start with an  $x^0 = 0$
- ② We get into recursion loop

- ↳  $x^1 = f(x^0) = -0.5$
- ↳  $x^2 = f(x^1) = -0.625$
- ↳ ...
- ↳  $x^\infty = -1$

- ① Start with an  $x^0 = 5$
- ② We get into recursion loop

- ↳  $x^1 = f(x^0) = -13$
- ↳  $x^2 = f(x^1) = -85$
- ↳ ...
- ↳  $x^\infty = -\infty$

We can now diverge if we start with a wrong initial point!

**Not all recursive forms are always converging!**

# Dynamic Programming: Applications to Our Problem

Our problem has a similar form: we need to solve *Bellman equations*  
which are *recursive equations*

So, we could use *DP* to find the solution

---

There are two major *DP* approaches

- Policy Iteration that uses recursion to iterate between
  - ↳ Policy *Evaluation*
  - ↳ Policy *Improvement*
- Value Iteration which applies recursion on *optimal Bellman* equation

Let's look at these two approaches in detail

# Policy Evaluation: Step I

The first step is *policy evaluation*: we can formulate this problem as follows

## Ultimate Goal of Policy Evaluation

Given a *policy*  $\pi$ , we intend to *evaluate* values of *all states* by *recursion*

Before we start, let's recap a few definitions: recall *expected policy reward*

$$\bar{\mathcal{R}}_{\pi}(s) = \sum_{m=1}^M \bar{\mathcal{R}}(a^m, s) \pi(a^m | s)$$

For sake of compactness, we use the following notation

$$\bar{\mathcal{R}}_{\pi}(s) = \mathbb{E}_{\pi}\{\bar{\mathcal{R}}(A, s) | s\}$$

# Policy Evaluation: Step I

Similarly, we define the notation

$$\mathbb{E}_{\pi} \{ v_{\pi} (\bar{S}) | s, a \} = \sum_{n=1}^N v_{\pi} (s^n) p(s^n | s, a)$$

and also denote its expected form over the *action set* by

$$\begin{aligned} \mathbb{E}_{\pi} \{ v_{\pi} (\bar{S}) | s \} &= \sum_{n=1}^N v_{\pi} (s^n) p_{\pi} (s^n | s) \\ &= \sum_{m=1}^M \underbrace{\sum_{n=1}^N v_{\pi} (s^n) p(s^n | s, a^m) \pi(a^m | s)}_{\mathbb{E}_{\pi} \{ v_{\pi} (\bar{S}) | s, a^m \}} \\ &= \sum_{m=1}^M \mathbb{E}_{\pi} \{ v_{\pi} (\bar{S}) | s, a^m \} \pi(a^m | s) \end{aligned}$$

## Policy Evaluation: Step I

We can then write the *Bellman equations* compactly as

$$v_{\pi}(s) = \bar{\mathcal{R}}_{\pi}(s) + \gamma \mathbb{E}_{\pi} \{ v_{\pi}(\bar{s}) | s \}$$

for *value function* and also as

$$q_{\pi}(s, a) = \bar{\mathcal{R}}(s, a) + \gamma \mathbb{E}_{\pi} \{ v_{\pi}(\bar{s}) | s, a \}$$

for *action-value* function

Now, we are ready to *evaluate a policy* by *recursion*

## Policy Evaluation: Value Computation via Recursion

Recall our perspective on value computation:

values are  $N$  unknowns that we want to compute from Bellman equations

Now, if someone claims that the values

$$v_{\pi}(s^n) = v_n$$

for  $n = 1 : N$  are values of policy  $\pi$ , can we confirm it?

- + Shouldn't we simply use Bellman Equation?!
- Exactly!

# Policy Evaluation: Value Computation via Recursion

We could confirm

$$v_{\pi}(s^n) = v_n$$

by writing first finding for every state  $s$

$$\begin{aligned} \mathbb{E}_{\pi}\{v_{\pi}(\bar{S})|s\} &= \sum_{n=1}^N v_{\pi}(s^n) p_{\pi}(s^n|s) \\ &= \sum_{n=1}^N \sum_{m=1}^M v_{\pi}(s^n) p(s^n|s, a^m) \pi(a^m|s) \\ &= \sum_{n=1}^N \sum_{m=1}^M \underbrace{v_n}_{\text{claimed value}} \underbrace{p(s^n|s, a^m)}_{\text{transition model}} \underbrace{\pi(a^m|s)}_{\text{policy}} \end{aligned}$$

# Policy Evaluation: Value Computation via Recursion

We could confirm

$$v_\pi(s^n) = v_n$$

by writing first finding for every state  $s$

$$\mathbb{E}_\pi \{v_\pi(\bar{s}) | s\} = \text{computed from } v_n's := F(\{v_1, \dots, v_N\}, s)$$

and then checking if

$$\begin{aligned} v_\pi(s^n) &= v_n = \bar{\mathcal{R}}_\pi(s^n) + \gamma \mathbb{E}_\pi \{v_\pi(\bar{s}) | s^n\} \\ &= \bar{\mathcal{R}}_\pi(s^n) + \gamma F(\{v_1, \dots, v_N\}, s) \end{aligned}$$

holds for all  $n = 1 : N$

## Policy Evaluation: Value Computation via Recursion

If it happens that the claimed  $v_\pi(\cdot)$  is **not** a valid claim; then, we get out of Bellman equation

$$\bar{v}_\pi(s^n) = \bar{v}_n = \bar{\mathcal{R}}_\pi(s^n) + \gamma \mathbb{E}_\pi \{ v_\pi(\bar{s}) | s^n \}$$

which is different from the claimed  $v_\pi(\cdot)$ , i.e.,  $v_n \neq \bar{v}_n$

### Policy Evaluation

We iterate this procedure until we can confirm, i.e., we

- ① set  $v_\pi(\cdot) \leftarrow \bar{v}_\pi(\cdot)$
- ② repeat the same procedure and compute new  $\bar{v}_\pi(\cdot)$

We **stop** when  $v_\pi(\cdot) = \bar{v}_\pi(\cdot)$ , or at least it happens approximately

# Policy Evaluation

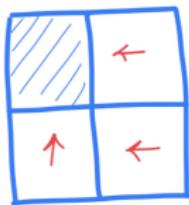
**PolicyEval**( $\pi, v_\pi^0$ ):

- 1: Initiate values with  $v_\pi^0$  and set  $k = 0$
  - 2: Make sure that  $v_\pi^0(s) = 0$  for **terminal states**  $s$
  - 3: Choose a **small** threshold  $\epsilon$  and initiate  $\Delta = +\infty$  # stopping criteria
  - 4: **for**  $n = 1 : N$  **do**
  - 5:   Compute  $\bar{\mathcal{R}}_\pi(s^n) = \mathbb{E}_\pi \{ \bar{\mathcal{R}}(s^n, a) \}$  # average rewards
  - 6:   **end for**
  - 7: **while**  $\Delta > \epsilon$  **do**
  - 8:   **for**  $n = 1 : N$  **do**
  - 9:     Update  $v_\pi^{k+1}(s^n) = \bar{\mathcal{R}}_\pi(s^n) + \gamma \mathbb{E}_\pi \{ v_\pi^k(\bar{S}) | s^n \}$  # DP update
  - 10:   **end for**
  - 11:    $\Delta = \max_n |v_\pi^{k+1}(s^n) - v_\pi^k(s^n)|$  # check convergence
  - 12:   Update  $k \leftarrow k + 1$
  - 13: **end while**
- Recursion Loop

## Attention

We should make sure that **terminal states** are all initiated with **zero** value

## Example: Dummy Grid World



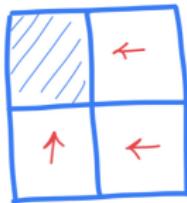
Let's try with our dummy grid world: we saw that

$$\bar{\mathcal{R}}_\pi(0) = 0 \quad \bar{\mathcal{R}}_\pi(1) = -1 \quad \bar{\mathcal{R}}_\pi(2) = -1 \quad \bar{\mathcal{R}}_\pi(3) = -1$$

Now let's evaluate its values by recursion: we first note that, if we have

$$\begin{array}{ll} \mathbb{E}_\pi \left\{ v_\pi^k (\bar{S}) | 0 \right\} = v_\pi^k(0) & \mathbb{E}_\pi \left\{ v_\pi^k (\bar{S}) | 1 \right\} = v_\pi^k(0) \\ \mathbb{E}_\pi \left\{ v_\pi^k (\bar{S}) | 2 \right\} = v_\pi^k(0) & \mathbb{E}_\pi \left\{ v_\pi^k (\bar{S}) | 3 \right\} = v_\pi^k(2) \end{array}$$

# Example: Dummy Grid World



PolicyEval( $\pi, v_\pi^0$ ) :

- 1: Initiate values with  $v_\pi^0(1), v_\pi^0(2)$  and  $v_\pi^0(3)$  at random and set  $v_\pi^0(0) = 0$
- 2: Set  $\epsilon = 0.001$ , and initiate  $\Delta = 1000$  # stopping criteria
- 3: while  $\Delta > \epsilon$  do
- 4:   Update  $v_\pi^{k+1}(1) = -1 + v_\pi^k(0)$  # DP update
- 5:   Update  $v_\pi^{k+1}(2) = -1 + v_\pi^k(0)$  # DP update
- 6:   Update  $v_\pi^{k+1}(3) = -1 + v_\pi^k(2)$  # DP update
- 7:    $\Delta = \max_{s \in \{1,2,3\}} |v_\pi^{k+1}(s) - v_\pi^k(s)|$  # check convergence
- 8:   Update  $k \leftarrow k + 1$
- 9: end while

It converges after only one recursion!

# Policy Improvement

Let us know recall **optimality constraint**: with optimal policy, we have

$$v_{\star}(s) = \max_m q_{\star}(s, a^m)$$

which can be achieved by policy

$$\pi^{\star}(a^m|s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\star}(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\star}(s, a^m) \end{cases}$$

This means that if  $\pi$  is not optimal, we would have

$$\pi(a^m|s) \neq \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\pi}(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\pi}(s, a^m) \end{cases}$$

# Policy Improvement

In other words, if we change our policy to

$$\bar{\pi}(a^m | s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\pi}(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\pi}(s, a^m) \end{cases}$$

Then, it should give us better values, i.e.,  $\bar{\pi} \geq \pi$ !

- + Are you sure?! I don't see it immediately
- We can actually show it!

This is what we call **policy improvement** theorem

# Policy Improvement

## Policy Improvement

Given (deterministic) policy  $\pi^k$ , we can always design a **better** policy  $\pi^{k+1}$  by setting it to

$$\pi^{k+1}(a^m|s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\pi^k}(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\pi^k}(s, a^m) \end{cases}$$

## Policy Improvement

**PolicyImprov**( $v_\pi$ ):

1: **for**  $n = 1 : N$  **do**

2:    **for**  $m = 1 : M$  **do**

3:      *Compute*  $\bar{\mathcal{R}}(\textcolor{blue}{s}^n, \textcolor{red}{a}^m)$

4:            $q_\pi$  (.

5:        end for

### 6: Compute an improved policy as

```
# policy improvement
```

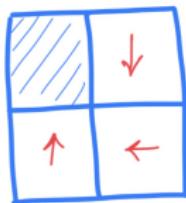
$$\bar{\pi}(a^m | s^n) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\pi}(s^n, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\pi}(s^n, a^m) \end{cases}$$

7: *end for*

## Attention

Here, we do **no** recursion

## Example: Dummy Grid World



Let's try dummy grid world with above non-optimal policy: here, we have

$$v_{\pi}(0) = 0 \quad v_{\pi}(1) = -3 \quad v_{\pi}(2) = -1 \quad v_{\pi}(3) = -2$$

We now look at *action-values* at the problematic state  $s = 1$

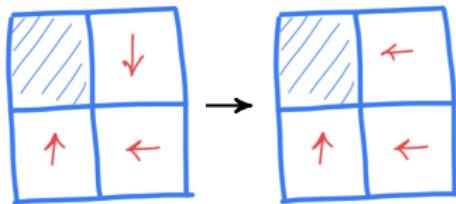
$$q_{\pi}(1, 0) = -1$$

$$q_{\pi}(1, 1) = -3$$

$$q_{\pi}(1, 2) = -3.5 \rightsquigarrow -3 = v_{\pi}(1) \neq \max_a q_{\pi}(1, a) = -1$$

$$q_{\pi}(1, 3) = -3.5$$

## Example: Dummy Grid World



Now if we improve the policy, we get

$$\bar{\pi}(a|1) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_{\pi}(1, a) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_{\pi}(1, a) \end{cases} = \begin{cases} 1 & a = 0 \\ 0 & a \neq 0 \end{cases}$$

which is actually **optimal**

## Policy Iteration: Improving Policy by Recursion

Looking at the **policy improvement** theorem, we see

If we plug in  $\pi^k = \pi^*$  into algorithm; then, after policy improvement

- ↳ we get  $\pi^{k+1} = \pi^*$
- ↳ say we **evaluate** values for  $\pi^{k+1} = \pi^*$  and plug back to algorithm
  - ↳ we get  $\pi^{k+2} = \pi^*$
  - ↳ say we **evaluate** values for  $\pi^{k+2} = \pi^*$  and plug back to algorithm
    - ↳ ...

So, **optimal policy** is a **fixed-point** for this **recursion**

## Policy Iteration

We can start with an **arbitrary policy**  $\pi^0$  and keep doing the above recursion until we see that  $\pi^{k+1} = \pi^k$  which indicates that we reached **optimal policy**

# Policy Iteration

```
PolicyItr():
```

- 1: Initiate with *random*  $v_\pi(s)$  for all **non-terminal** states  $s$

- 2: Set  $v_\pi(s) = 0$  for **terminal** states  $s$

- 3: Initiate two random policies  $\pi$  and  $\bar{\pi}$

- 4: **while**  $\pi \neq \bar{\pi}$  **do**

- 5:    $v_\pi = \text{PolicyEval}(\pi, v_\pi)$  and  $\pi \leftarrow \bar{\pi}$    **Recursion**

- 6:    $\bar{\pi} = \text{PolicyImprov}(v_\pi)$

- 7: **end while**

Recursion

Note that this is a **nested recursive computation**

- There is a loop for recursion inside the algorithm in which
  - ↳ at each iteration we evaluate the policy recursively
- But, we initiate each **policy evaluation** loop with the values of last iteration
  - ↳ this can improve the convergence speed

## Back-Tracking by Recursion

- + But wait a Moment! We already talked about back-tracking optimal policy from *Bellman optimality equation!* Don't we implement that?!
- Sure! We can do the same thing by recursion

We follow the same idea but we use recursion

- ① We can find optimal values from *Bellman optimality equations*
  - ↳ This is where we use *recursion*
- ② We could then find the *optimal action*-values
- ③ We finally get the *optimal policy* from *optimal action*-values

# Recall: Back-Tracking from Optimal Values

OptimBackTrack():

- 1: Solve Bellman equations # we use recursion!
- 2: **for**  $n = 1 : N$  **do**
- 3:   **for**  $m = 1 : M$  **do**
- 4:     Set  $q_*(s^n, a^m) = \bar{\mathcal{R}}(s^n, a^m) + \gamma \mathbb{E} \{v_*(\bar{S}) | s^n, a^m\}$  # action-values
- 5:   **end for**
- 6:   Compute optimal policy via optimality constraint

$$\pi^*(a^m | s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_*(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_*(s, a^m) \end{cases}$$

- 7: **end for**

# Recursion with Bellman Optimality

Recall Bellman optimality equation

$$v_\star(s) = \max_{\mathbf{a}} (\bar{\mathcal{R}}(s, \mathbf{a}) + \gamma \mathbb{E} \{ v_\star(\bar{s}) | s, \mathbf{a} \})$$

We can again solve it by recursion: we start with some  $v_\star^0(\cdot)$  and then for every state  $s$  and action  $a$ , we compute

$$\mathbb{E} \{ v_\star^k(\bar{s}) | s, \mathbf{a} \} = \sum_{n=1}^N \underbrace{v_\star^k(s^n)}_{\text{last computed value}} \underbrace{p(s^n | s, \mathbf{a})}_{\text{transition model}}$$

We then update the optimal value function as

$$v_\star^{k+1}(s) = \max_{\mathbf{a}} (\bar{\mathcal{R}}(s, \mathbf{a}) + \gamma \mathbb{E} \{ v_\star^k(\bar{s}) | s, \mathbf{a} \})$$

## Value Iteration vs Policy Iteration

Before we complete the value iteration algorithm: *it is interesting to put its recursion next to the one used for policy evaluation*

*With optimality equation, we iterate as*

$$v_{\star}^{k+1}(s) = \max_a [\bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \left\{ v_{\star}^k(\bar{S}) | s, a^m \right\}]$$

*With Bellman equation for a given policy  $\pi$ , we iterate as*

$$\begin{aligned} v_{\pi}^{k+1}(s) &= \bar{\mathcal{R}}_{\pi}(s) + \gamma \mathbb{E}_{\pi} \left\{ v_{\pi}^k(\bar{S}) | s \right\} \\ &= \sum_{m=1}^M \left( \bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \left\{ v_{\pi}^k(\bar{S}) | s, a^m \right\} \right) \pi(a^m | s) \end{aligned}$$

# Value Iteration vs Policy Iteration

With optimality equation, we iterate as

$$v_{\star}^{k+1}(s) = \max_{a^m} \left[ \bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \left\{ v_{\star}^k(\bar{S}) | s, a^m \right\} \right]$$

With Bellman equation for a given policy  $\pi$ , we iterate as

$$v_{\pi}^{k+1}(s) = \sum_{m=1}^M \left( \bar{\mathcal{R}}(s, a^m) + \gamma \mathbb{E} \left\{ v_{\pi}^k(\bar{S}) | s, a^m \right\} \right) \pi(a^m | s)$$

This indicates that for both recursive loops

- we compute  $M$  values per iteration per state
  - ↳ in policy iteration, we compute the average of these  $M$  via  $\pi$
  - ↳ in value iteration, we take the largest among these  $M$  values

# Value Iteration

**ValueItr():**

1: *Initiate with random  $v_*^0(s)$  for all states, and set  $v_*^0(s) = 0$  for terminal states*

2: *Choose a small threshold  $\epsilon$ , initiate  $\Delta = +\infty$  and  $k = 0$*

3: **while**  $\Delta > \epsilon$  **do**

4:   **for**  $n = 1 : N$  **do**

5:     **for**  $m = 1 : M$  **do**

6:       Compute  $q_*(s^n, a^m) = \bar{\mathcal{R}}(s^n, a^m) + \gamma \mathbb{E} \left\{ v_*^k(\bar{S}) | s^n, a^m \right\}$

7:     **end for**

8:       Update  $v_{\pi}^{k+1}(s^n) = \max_m q_*(s^n, a^m)$  # DP update

9:     **end for**

10:   Set  $\Delta = \max_n |v_{\pi}^{k+1}(s^n) - v_{\pi}^k(s^n)|$  and  $k \leftarrow k + 1$

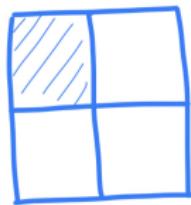
Recursion

11: **end while**

12: Compute an optimal policy as

$$\bar{\pi}(a^m | s) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} q_*(s, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} q_*(s, a^m) \end{cases}$$

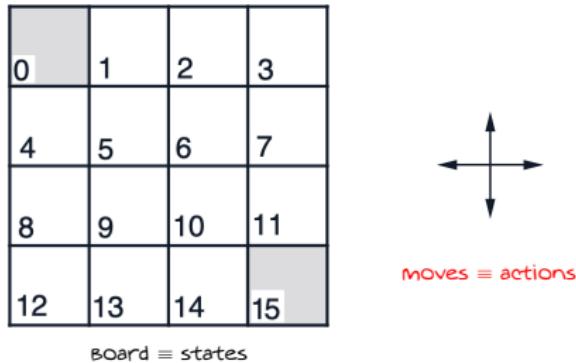
## Example: Dummy Grid World



You may try policy and value iteration for this problem at home!

Easy as Pie 😊

## Example: A Bit Larger Grid World<sup>1</sup>



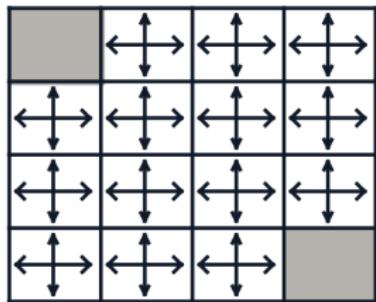
Let's do a bit of more serious example: we are now in a  $4 \times 4$  grid world

- We have two **terminal states** shown in gray
- Each move we do gets a -1 reward
  - ↳ We also get -1 reward if we hit a corner
  - ↳ We get zero reward at terminal state

In simple words: we are looking for **shortest path** to the **corners**

<sup>1</sup>This example is taken from Sutton and Barto's Book; Example 4.1 in Chapter 4

## Example: A Bit Larger Grid World



initial policy

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

initial values

Let's first try **policy iteration**: we start with

- a uniform random policy  $\pi^0$
- all values being zero, i.e.,  $v_{\pi^0}^0(s) = 0$  for all  $s$

## Example: A Bit Larger Grid World

Recall policy iteration:

```
PolicyItr():
```

- 1: Initiate with  $v_{\pi}$  ( $s$ ) for all non-terminal states  $s$

- 2: Set  $v_{\pi}(s) = 0$  for terminal states  $s$

- 3: Initiate two random policies  $\pi$  and  $\bar{\pi}$

- 4: **while**  $\pi \neq \bar{\pi}$  **do**

- 5:    $v_{\pi} = \text{PolicyEval}(\pi, v_{\pi})$  and  $\pi \leftarrow \bar{\pi}$    **Recursion**

- 6:    $\bar{\pi} = \text{PolicyImprov}(v_{\pi})$

Recursion

- 7: **end while**

We should start with  $v_{\pi_0}^0(\cdot)$  and do the red recursion first

- at the end of this recursion we have evaluated the random policy

# Example: A Bit Larger Grid World

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$v_{\pi^0}^0$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$v_{\pi^0}^1$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$v_{\pi^0}^2$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$v_{\pi^0}^3$

...

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

$v_{\pi^0}^\infty$

We now have **evaluated** the value of random policy  $v_{\pi^0} = v_{\pi^0}^\infty$

## Example: A Bit Larger Grid World

Recall policy iteration:

```
PolicyItr():
```

- 1: Initiate with  $v_{\pi}$  ( $s$ ) for all non-terminal states  $s$

- 2: Set  $v_{\pi}(s) = 0$  for terminal states  $s$

- 3: Initiate two random policies  $\pi$  and  $\bar{\pi}$

- 4: **while**  $\pi \neq \bar{\pi}$  **do**

- 5:    $v_{\pi} = \text{PolicyEval}(\pi, v_{\pi})$  and  $\pi \leftarrow \bar{\pi}$    **Recursion**

- 6:    $\bar{\pi} = \text{PolicyImprov}(v_{\pi})$

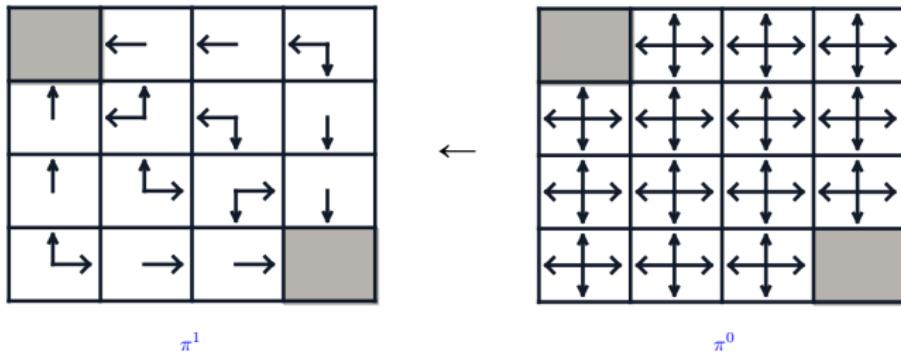
Recursion

- 7: **end while**

Next, we do the **outer recursion** recursion, i.e.,

- we **improve** the policy

## Example: A Bit Larger Grid World



We improve policy by taking actions with maximal action-values

- if we have multiple maximal action-values we can behave randomly

# Example: A Bit Larger Grid World

Recall policy iteration:

PolicyItr():

1: Initiate with  $v_\pi(s)$  for all non-terminal states  $s$

2: Set  $v_\pi(s) = 0$  for terminal states  $s$

3: Initiate two random policies  $\pi$  and  $\bar{\pi}$

4: while  $\pi \neq \bar{\pi}$  do

5:    $v_\pi = \text{PolicyEval}(\pi, v_\pi)$  and  $\pi \leftarrow \bar{\pi}$     Recursion

6:    $\bar{\pi} = \text{PolicyImprov}(v_\pi)$

7: end while

Recursion

We now start with  $v_{\pi^1}^0 = v_{\pi^0} = v_{\pi^0}^\infty$  and do the red recursion again

- at the end of this recursion we have evaluated the new policy  $\pi^1$

## Example: A Bit Larger Grid World

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

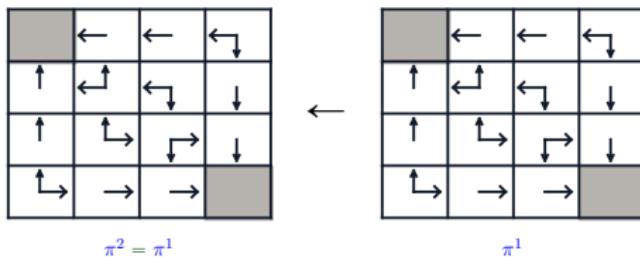
$v_{\pi^1}^0$

...

0.0	-1.0	-2.0	-3.0
-1.0	-2.0	-3.0	-2.0
-2.0	-3.0	-2.0	-1.0
-3.0	-2.0	-1.0	0.0

$v_{\pi^1}^{+\infty}$

After *evaluating* policy  $\pi^1$  as  $v_{\pi^1} = v_{\pi^1}^{\infty}$ , we do the next improvement



Well  $\pi^2 = \pi^1$  and we should *stop!*

## Example: A Bit Larger Grid World

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

initial values

Now we try **value iteration**: for start, we only need an initial value, so we set

- all values being zero, i.e.,  $v_\star^0(s) = 0$  for all  $s$

We keep recursion until we find the **optimal values**

# Example: A Bit Larger Grid World

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

 $v_*^0$ 

...

0.0	-1.0	-2.0	-3.0
-1.0	-2.0	-3.0	-2.0
-2.0	-3.0	-2.0	-1.0
-3.0	-2.0	-1.0	0.0

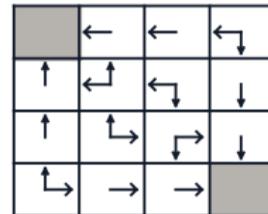
 $v_*^{+\infty}$ 

Now, we back-track the optimal policy  $\pi^*$

0.0	-1.0	-2.0	-3.0
-1.0	-2.0	-3.0	-2.0
-2.0	-3.0	-2.0	-1.0
-3.0	-2.0	-1.0	0.0

 $v_*$ 

action-values

 $q_*$  $\pi^*$

# Complexity of Policy and Value Iteration

- + It seems that value iteration has **less complexity!**
- Well, it is not in order, but yes! It usually converge faster

In our example with **policy iteration**, we had to **evaluate two policies**

- once for  $\pi^0$  and once for  $\pi^1$
- say the first recursion took  $K_1$  iterations and the second took  $K_2$ 
  - ↳ the **total number** of iterations is then  $K_1 + K_2$
  - ↳ in practice, it often happens that  $K_2 \ll K_1$ 
    - ↳ because we already start from **good values** with  $v_{\pi^1}^0 = v_{\pi^0}^{+\infty}$

With **value iteration**, we had to **only evaluate optimal policy**

- say it takes  $K_\star$  iterations: there is **no reason** that  $K_\star$  be same as  $K_1$  or  $K_2$ 
  - ↳ each evaluation has a **different initial** and **converging point**
  - ↳ in practice, it often happens that  $K_\star > K_1$  and  $K_\star \gg K_2$ 
    - ↳ so it **might be** that  $K_\star \approx K_1 + K_2$
    - ↳ but usually with **multiple policy improvements**, we see  $K_\star < K_1 + K_2 + \dots$

# Complexity of Policy and Value Iteration

- + If so, why should we use *policy iteration*?!
- Well, not all problems are like a dummy grid world

In practice, it might be computationally **hard** to get **very close to optimal values**

- in this case, we take non-converged values
  - ↳ we consider them *estimates* of optimal values
- in value iteration we **approximate** optimal policy with on these *estimates*
  - ↳ this might be a **loose** estimate

If we do the same **approximative** computation with policy iteration

- we often end up with a **better policy**

## Moral of Story

While **value iteration** typically show **faster convergence**, **policy iteration** can give **better policies** after convergence

# Generalized Policy Iteration

In practice, we can terminate or change the order of computation in policy iteration to reduce its complexity: for instance, we could have

GenPolicyItr():

- 1: Initiate with  $\text{random } v_\pi(s)$  for all **non-terminal** states  $s$
- 2: Set  $v_\pi(s) = 0$  for **terminal** states  $s$
- 3: Initiate two random policies  $\pi$  and  $\bar{\pi}$
- 4: **while**  $\pi \neq \bar{\pi}$  **do**
- 5:    $v_\pi = \text{TerminPolicyEval}(\pi, v_\pi)$  and  $\pi \leftarrow \text{changed}$
- 6:    $\bar{\pi} = \text{PolicyImprov}(v_\pi)$
- 7: **end while**

where  $\text{TerminPolicyEval}(\pi, v_\pi)$  evaluates *policy*  $\pi$  from starting value function  $v_\pi$  with a **terminating** recursion loop

# Generalized Policy Iteration: Terminating Evaluation

`TerminPolicyEval( $\pi, v_\pi^0$ ):`

```

1: Initiate values with  $v_\pi^0$  and set  $k = 0$ 
2: Make sure that  $v_\pi^0(s) = 0$  for terminal states  $s$ 
3: Choose a small threshold  $\epsilon$  and initiate  $\Delta = +\infty$            # stopping criteria
4: for  $n = 1 : N$  do
5:   Compute  $\bar{\mathcal{R}}_\pi(s^n) = \mathbb{E}_\pi \{ \bar{\mathcal{R}}(s^n, a) \}$            # average response
6:   end for
7:   while  $\Delta > \epsilon$  and  $k < K$  do changed
8:     for  $n = 1 : N$  do
9:       Update  $v_\pi^{k+1}(s^n) = \bar{\mathcal{R}}_\pi(s^n) + \gamma \mathbb{E}_\pi \{ v_\pi^k(\bar{S}) | s^n \}$            # DP update
10:    end for
11:     $\Delta = \max_n |v_\pi^{k+1}(s^n) - v_\pi^k(s^n)|$            # check convergence
12:    Update  $k \leftarrow k + 1$ 
13:   end while

```

Obviously,  $\text{TerminPolicyEval}(\pi, v_\pi)$  does **not** return the *exact values* of the policy  $\pi$ , but only an *estimate of them*

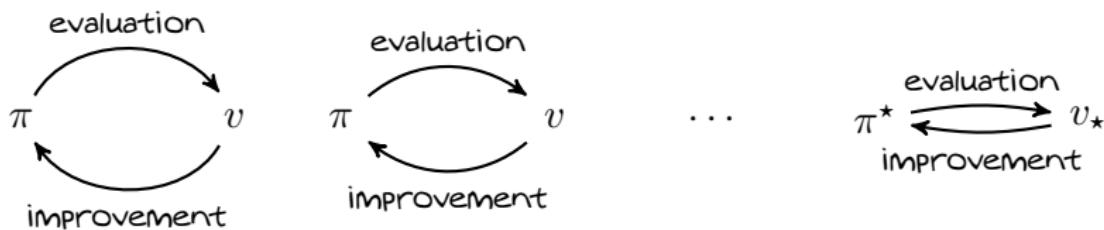
# Generalized Policy Iteration

We can come up with various such ideas: *these variants are often called*

*Generalized Policy Iteration  $\equiv$  GPI*

*These approaches all rely on*

*back-and-forth computation of policies and values*



*If designed properly, they all converge to optimal policy and optimal values*

## Some Final Remarks

- + We know the algorithms now, but how can we **guarantee** that they **converge**? You showed us an simple example that recursion could simply **diverge**!
- Well, we can show that what we discussed in this chapter converge: it comes from the nice properties of **Bellman equations**
  - ↳ There are several proofs; for instance see [a proof in Tom Mitchell's notes](#)

---

When it comes to practice, most known algorithms are proved to **converge to optimal policy and optimal values**; however, note that

- **Convergence guarantee** is different from the **speed of convergence**
  - ↳ An algorithm might **converge**, but **very slow**
- If you deal with an **unknown** algorithm; then, you should make sure that it **converges** to **optimal policy** and **optimal values**