

Reinforcement Learning

Chapter 3: Model-free RL

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Classical RL Methods: Recall

*Ultimate goal in an RL problem is to find the **optimal policy***

As mentioned, we have two **major challenges** in this way

- ① We need to compute **values** explicitly
- ② We often deal with settings with **huge state spaces**?

In this part of the course, we are going to handle the first challenge

- Previous chapter \rightsquigarrow Model-based methods
- This chapter \rightsquigarrow Model-free methods

Finally We Got Serious: Model-free RL

In **model-free** methods

we **do not have** an analytic model for the behavior of environment

We intend to compute **values** from **real data collected from environment**

Model-Based RL

Bellman Equation

value iteration

policy iteration

Model-free RL

on-policy methods

temporal difference

Monte Carlo

SARSA

off-policy methods

Q-learning

Model-free RL in Nutshell

- + If this is the **typical case** in RL problems, why did we spend so much time on learning **MDPs** and finding **optimal policy** there?
- Well! We need all those things, since we are going to do **the same thing** here **only without explicit model**

In a nutshell, we are going to find a way to apply

Generalized Policy Iteration \equiv GPI

But, now **without** knowing the **transition-rewarding function**

Let's take a look back at GPI

Generalized Policy Iteration

We wrote the pseudo-code for GPI as below

```
GenPolicyItr():
1: Initiate two random policies  $\pi$  and  $\bar{\pi}$ 
2: while  $\pi \neq \bar{\pi}$  do
3:    $v_\pi = \text{GenPolicyEval}(\pi)$  and  $\pi \leftarrow \bar{\pi}$ 
4:    $\bar{\pi} = \text{PolicyImprov}(v_\pi)$ 
5: end while
```

Let's recall where we had to use environment's model

- ① In **policy evaluation** phase when we compute values via **Bellman equations**
- ② In **policy improvement** when we compute action-values out of **values**

How can we do these tasks **without** knowing **transition-rewarding model**?

Computing Statistics from Data

Let's start with a very simple problem: assume we have an **unknown** signal generator which returns signals at **random**; this generator is connected to a device and we can **only see** the output of this device, i.e., we see

$$Y = f(\mathbf{X})$$

where \mathbf{X} is the **random** signal and $f(\cdot)$ denotes transform by the device

We want to know the **expected output** of our device, i.e.,

$$\mu_Y = \mathbb{E}\{Y\} = \mathbb{E}\{f(\mathbf{X})\}$$

If we knew the model of the **generator's model**, we could write

$$\mu_Y = \mathbb{E}\{f(\mathbf{X})\} = \sum_{\substack{x \in \mathcal{X} \\ \text{all outcomes}}} f(\mathbf{x}) \underbrace{p(x)}_{\text{model}}$$

Monte-Carlo Method

Now what can we do if we **don't know** the **model**

- + Well! Shouldn't we evaluate it by a **simple** numerical simulation?
- Exactly! This is what we call it **Monte-Carlo method**

In **Monte-Carlo method**, we sample our device K times **independently** as

$$Y_1, Y_2, \dots, Y_K$$

Then we **estimate** the **expected** value as

$$\hat{\mu}_Y = \frac{1}{K} \sum_{k=1}^K Y_k$$

Monte-Carlo Method

- + Why does Monte-Carlo work?
- Simply because of central limit theorem

Since the sequence Y_1, Y_2, \dots, Y_K contains independent samples of identical process, we could say that

$$\hat{\mu}_Y \sim \mathcal{N}\left(\mu_Y, \frac{\sigma^2}{K}\right)$$

when K is large enough: so we could think of it as

$$\hat{\mu}_Y \approx \mu_Y + \frac{\varepsilon}{\sqrt{K}}$$

for some random error term ε : this error vanishes as K goes large

Computing Values via Monte-Carlo

- + But, how can we apply this idea to RL? I don't see any connection!
- Well! Think of **rewards and transitions** as random signal and **value function** as device! We only need to take **enough samples** from the **environment**

Let's start with a very simple task: we want to compute the value of **state s** for **policy π** in an **episodic** environment. Monte-Carlo suggest that

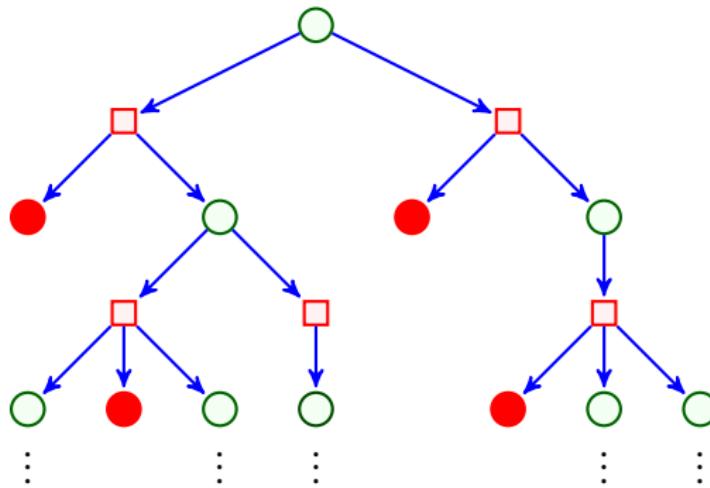
- ① We start at **state s** and play with **policy π** until we meet **terminal state**: say it happens at **time T**
- ② We compute the **sample** return as $G[1] = R_1 + \gamma R_2 + \dots + \gamma^{T-1} R_T$
- ③ We repeat this for K episodes and each episode, we collect $G[k]$

Then, we could estimate the value of **state s** as

$$\hat{v}_\pi(s) = \frac{1}{K} \sum_{k=1}^K G[k]$$

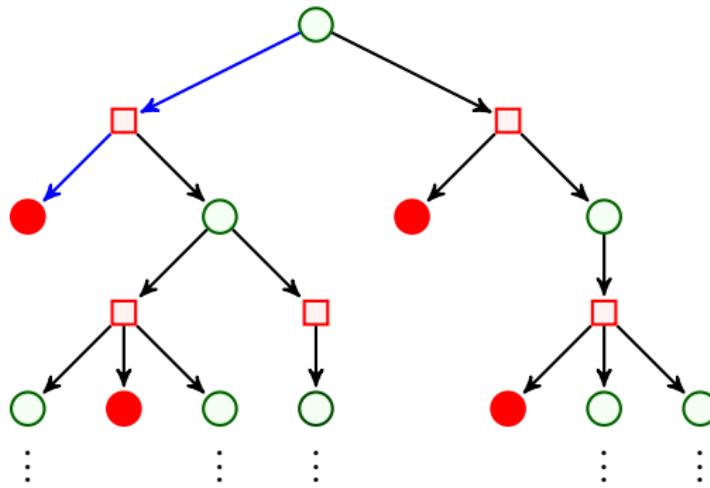
Values via Monte-Carlo: Trajectory Sampling

We can look at this approach as estimating values from *sample trajectories*: with known model, we can compute values by averaging over *possible trajectories*



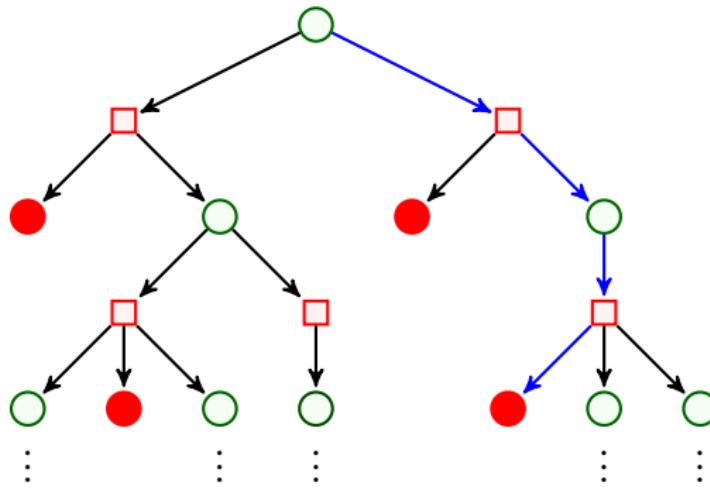
Values via Monte-Carlo: Trajectory Sampling

We can look at this approach as estimating values from *sample trajectories*: without known model, we can *sample* them and estimate values from them



Values via Monte-Carlo: Trajectory Sampling

We can look at this approach as estimating values from *sample trajectories*: without known model, we can *sample* them and estimate values from them



Computing Values via Monte-Carlo: Algorithm I

Let's put our estimation approach into an algorithm

MC_verI(π, s):

- 1: Initiate estimator of value as $\hat{v}_\pi(s) = 0$
- 2: **for** episode = 1 : K **do**
- 3: Initiate with state $S_0 = s$ and act via policy $\pi(a|s)$
- 4: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

- 5: Compute sample return $G = R_1 + \gamma R_2 + \dots + \gamma^{T-1} R_T$
- 6: Update estimate of value as $\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + G/K$
- 7: **end for**

Computing Values via Monte-Carlo

- + But, doesn't that take too long to compute a single value?
- Yes! This is in general a problem; however, in our **naive** algorithm it is too much delayed!

In our algorithm, we need to wait till very end of K episodes to access an estimate, but we rather prefer to have a bad estimate which gradually improves over episodes

We could use the idea of **online** averaging \equiv **incremental** averaging

Let's find out what it is!

Online Averaging

Say, we want to compute the average of K samples: we could write

$$\begin{aligned}\eta_K &= \frac{1}{K} \sum_{k=1}^K G_k = \frac{1}{K} \left(\sum_{k=1}^{K-1} G_k + G_K \right) = \frac{1}{K} ((K-1) \eta_{K-1} + G_K) \\ &= \left(1 - \frac{1}{K}\right) \eta_{K-1} + \frac{G_K}{K} \\ &= \eta_{K-1} + \frac{1}{K} (G_K - \eta_{K-1})\end{aligned}$$

But, we can define the previous average as

$$\eta_{K-1} = \frac{1}{K-1} \sum_{k=1}^{K-1} G_k \rightsquigarrow \sum_{k=1}^{K-1} G_k = (K-1) \eta_{K-1}$$

Online Averaging: Geometric Weights

Online Averaging

We can update the average in *online fashion* as

$$\eta_K = \eta_{K-1} + \frac{1}{K} \Delta_K$$

where $\Delta_K = G_K - \eta_{K-1}$ is the deviation in K -th episode

The above expression is given for *uniform averaging weights*, i.e., all samples have same weights: in more general form, we usually update

$$\eta_K = \eta_{K-1} + \alpha \Delta_K$$

for some $0 < \alpha \leq 1$ that can be *fixed* or *scaled with K*

- if it is *fixed* \equiv computing weighted average with *geometric weights*
- if it is *scaled linearly with K* \equiv computing *linear averaging*

Computing Values via Monte-Carlo: Algorithm II

Let's modify our earlier algorithm with online averaging

MC_verII(π, s):

- 1: Initiate estimator of value as $\hat{v}_\pi(s) = 0$
- 2: **for** episode = 1 : K **do**
- 3: Initiate with state $S_0 = s$ and act via policy $\pi(a|s)$
- 4: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

- 5: Compute sample return $G = R_1 + \gamma R_2 + \dots + \gamma^{T-1} R_T$
- 6: Update estimate of value as $\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + \alpha(G - \hat{v}_\pi(s))$
- 7: **end for**

Now after each episode, we have an estimate of value function at state s

Computing Values via Monte-Carlo: Improve Efficiency

In our algorithm: we go through the whole trajectory to compute the value on the state we started with! This does not sound *sample efficient*!

- + Well! What can we do **more**?! It seems to be the case!
- Not really! We can estimate values of **other states** down the **trajectory**!

In the following sample trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

it's not only S_0 whose sample return can be computed! We can also compute sample returns of S_1, \dots, S_{T-1}

All-Visit Monte-Carlo: Algorithm III

This concludes a policy evaluation algorithm based on Monte-Carlo

MC_Eval(π):

- 1: Initiate estimator of value as $\hat{v}_\pi(s^n) = 0$ for $n = 1 : N$
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a random state S_0 and act via policy $\pi(a|s)$
- 4: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

- 5: Initiate with $G = 0$
- 6: **for** $t = T - 1 : 0$ **do**
- 7: Update current return $G \leftarrow R_{t+1} + \gamma G$
- 8: Update estimate of value as $\hat{v}_\pi(S_t) \leftarrow \hat{v}_\pi(S_t) + \alpha(G - \hat{v}_\pi(S_t))$
- 9: **end for**
- 10: **end for**

All-Visit Monte-Carlo: Convergence

It's *intuitive* to say this algorithm converges to true values after *lots of episodes*

Asymptotic Convergence of Monte-Carlo

Let $\mathcal{C}_K(s)$ denote number of visits at *state s* during K Monte-Carlo episodes. Assume that the random state initialization is distributed such that $\mathcal{C}_K(s^n)$ grows large as K increases for $n = 1 : N$, i.e.,

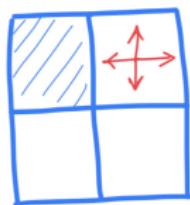
$$\lim_{K \rightarrow \infty} \mathcal{C}_K(s^n) = \infty$$

Then, as $K \rightarrow \infty$ the estimator of value function converges to its exact expression, i.e.,

$$\hat{v}_\pi(s) \xrightarrow{K \uparrow \infty} v_\pi(s)$$

for any *state s*

Example: Dummy Grid World with Random Walk



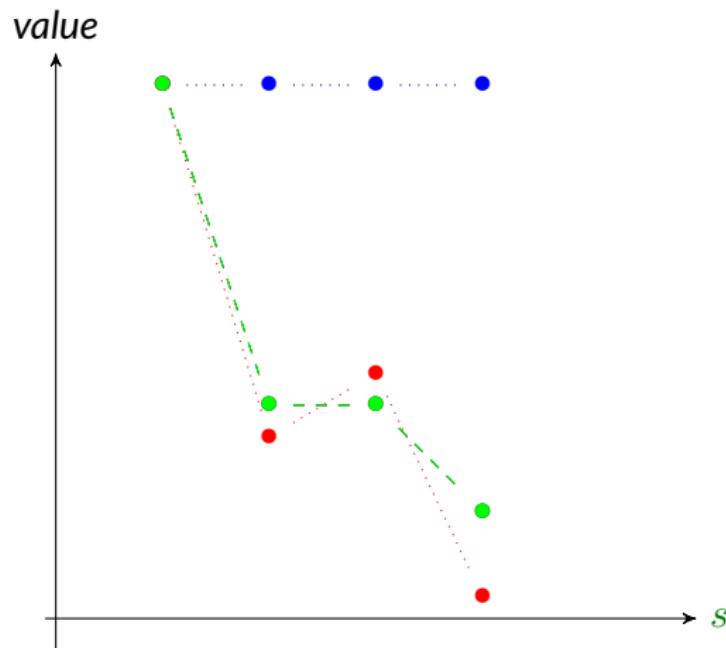
Let's get back to our **dummy world**: we now use **Monte-Carlo method** to compute the values for **uniform random policy**, i.e.,

$$\pi(a|s) = \frac{1}{4}$$

for all **actions** and **states**. From **Bellman equations**, we have

$$v_{\pi}(0) = 1 \quad v_{\pi}(1) = -4.5 \quad v_{\pi}(2) = -4.5 \quad v_{\pi}(3) = -6$$

Example: Dummy Grid World with Random Walk



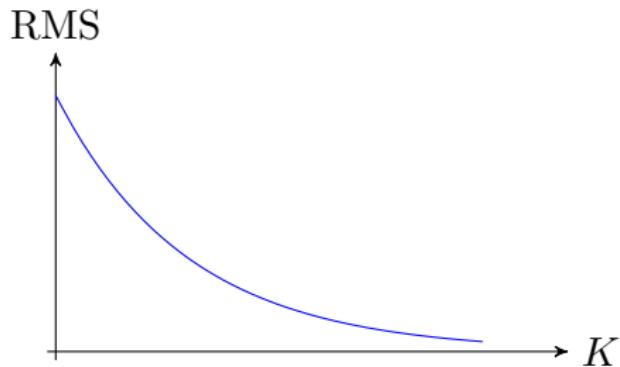
Typical Behavior: Variation Against Number of Episodes

We can compute the error of our estimation in each episode

$$\text{RMS} = \sqrt{\sum_{n=1}^N |\hat{v}_\pi(s^n) - v_\pi(s^n)|^2}$$

if we know the true value function, e.g., our random walk example

If we plot it against K ; then, we see



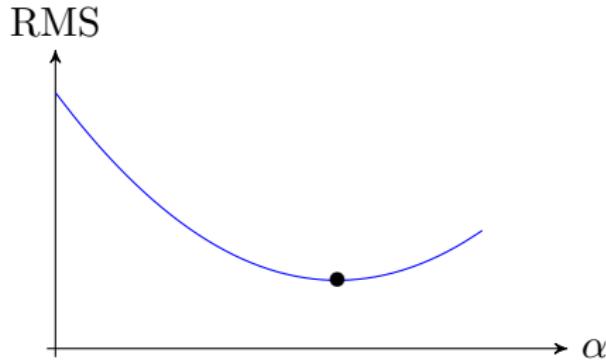
Typical Behavior: Variation Against Averaging Coefficient

We can compute the error of our estimation in each episode

$$\text{RMS} = \sqrt{\sum_{n=1}^N |\hat{v}_\pi(s^n) - v_\pi(s^n)|^2}$$

if we know the true value function, e.g., our random walk example

If we plot it against α ; then, we could see a minimum



Monte-Carlo Method: Action-Values

- + Now that we have **Monte-Carlo algorithm**, can we use it in GPI?
- Not yet! Remember that we need **action-values** for **policy improvement**

In GPI, we used to use **Bellman equation** for this

$$\begin{aligned}
 q_{\pi}(s, a) &= \bar{\mathcal{R}}(s, a) + \gamma \mathbb{E} \{ v_{\pi}(\bar{S}) | s, a \} \\
 &= \mathbb{E} \{ R_{t+1} | S_t = s, A_t = a \} + \gamma \mathbb{E} \{ v_{\pi}(S_{t+1}) | S_t = s, A_t = a \} \\
 &= \sum_{\ell=1}^L \sum_{n=1}^N \left(r^{\ell} + \gamma v_{\pi}(s^n) \right) \underbrace{p(r^{\ell}, s^n | s, a)}_{\text{transition-rewarding model}}
 \end{aligned}$$

But, now we cannot use it anymore!

Maybe, we can use **Monte-Carlo method** to estimate **action-values** directly

All-Visit Monte-Carlo: Action-Values

`MC_QEval(π):`

- 1: *Initiate estimator as $\hat{q}_\pi(s^n, a^m) = 0$ for $n = 1 : N$ and $m = 1 : M$*
- 2: **for** $episode = 1 : K$ **do**
- 3: *Initiate with a random state-action pair (S_0, A_0) and act via policy $\pi(a|s)$*
- 4: *Sample a trajectory*

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminal}$$

- 5: *Initiate with $G = 0$*
- 6: **for** $t = T - 1 : 0$ **do**
- 7: *Update current return $G \leftarrow R_{t+1} + \gamma G$*
- 8: *Update $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$*
- 9: **end for**
- 10: **end for**

We can now apply GPI using the Monte-Carlo method!

Policy Iteration with Monte-Carlo

We can use Monte-Carlo method to compute the *action-values*

- We then improve in each iteration by selecting best *action* for each *state*
 - ↳ This is what we typically call *greedy* improvement

```
MC_PolicyItr():
```

- 1: Initiate two random policies π and $\bar{\pi}$
- 2: **while** $\pi \neq \bar{\pi}$ **do**
- 3: $\hat{q}_\pi = \text{MC_QEval}(\pi)$ and $\pi \leftarrow \bar{\pi}$
- 4: $\bar{\pi} = \text{Greedy}(\hat{q}_\pi)$
- 5: **end while**

Policy Iteration with Monte-Carlo

Algorithmically, we can write the greedy update as

Greedy(\hat{q}_π):

- 1: **for** $n = 1 : N$ **do**
- 2: Improve the by taking **deterministically** the **best action**

$$\bar{\pi}(a^m | s^n) = \begin{cases} 1 & m = \underset{m}{\operatorname{argmax}} \hat{q}_\pi(s^n, a^m) \\ 0 & m \neq \underset{m}{\operatorname{argmax}} \hat{q}_\pi(s^n, a^m) \end{cases}$$

- 3: **end for**

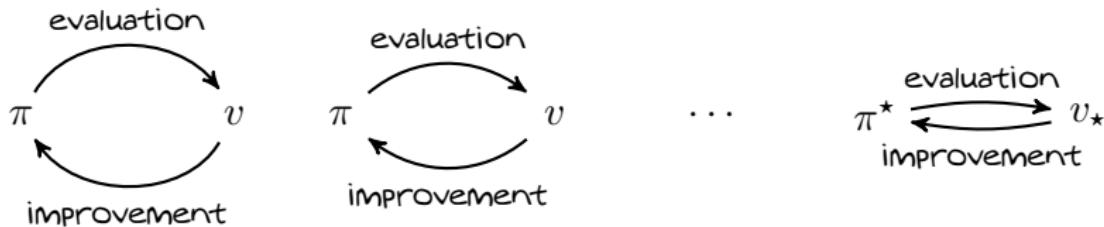
This is however not the **best** we could do!

We are going to have a whole lecture about it

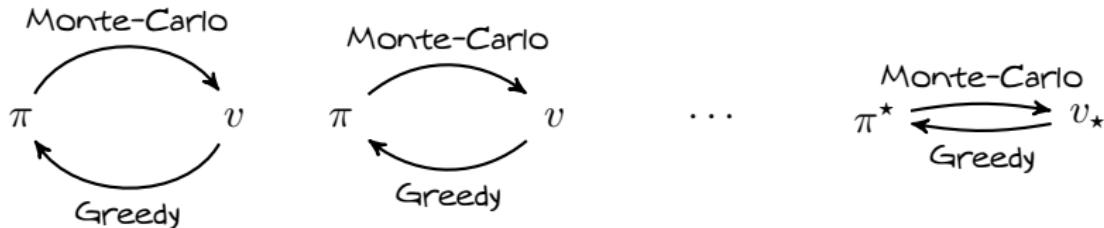
Stay tuned! We get back to this point in Section 4

GPI with Monte-Carlo

For any GPI, we said that we can think of

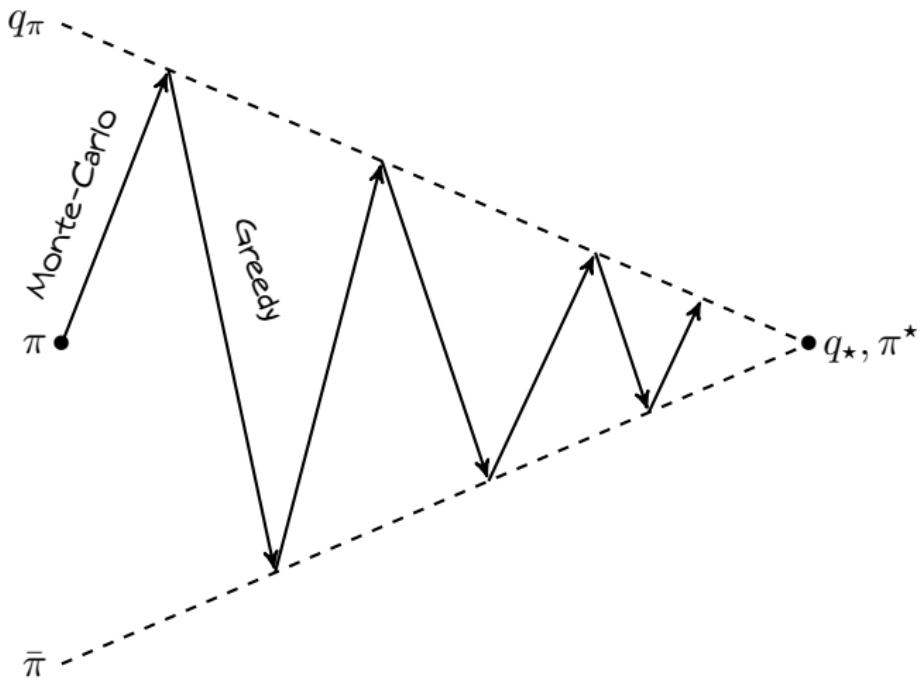


With **Monte-Carlo evaluation**, we can show this procedure as



GPI with Monte-Carlo

Another way to visualize this procedure is to think of following diagram



Non-Episodic Monte-Carlo: Terminating Trajectory

- + We only discussed *episodic scenarios!* Don't we use model-free RL in *non-episodic environment?*
- Sure we do! But, **Monte-Carlo** is not the best approach

A basic idea in this case is to *terminate sample trajectories*

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- With *long enough T* and $\gamma < 1$ the very later terms are *ineffective*
- But, we *cannot* use all the states in the trajectories
 - ↳ Sample returns of those who are *close to time T* are *not reliable!*

$$G_{T-1} = R_T + \underbrace{\gamma R_{T+1} + \dots}_{\text{we terminated them!}}$$

Terminating Monte-Carlo

TerminMC_Eval(π):

- 1: Initiate estimator of value as $\hat{v}_\pi(s^n) = 0$ for $n = 1 : N$
- 2: Choose **very large T and W** that satisfy $W < T$
- 3: **for** episode = 1 : K **do**
- 4: Initiate with a **random state S_0** and act via policy $\pi(a|s)$
- 5: Sample a trajectory and terminate after **T time steps**

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T: \text{terminated}$$

- 6: Initiate with $G = 0$
- 7: **for** $t = T - 1 : 0$ **do**
- 8: Update current return $G \leftarrow R_{t+1} + \gamma G$
- 9: **if** $t < T - W$ **then**
- 10: Update estimate of value as $\hat{v}_\pi(S_t) \leftarrow \hat{v}_\pi(S_t) + \alpha(G - \hat{v}_\pi(S_t))$
- 11: **end if**
- 12: **end for**
- 13: **end for**