

Reinforcement Learning

Chapter 3: Model-free RL

Ali Bereyhi


ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto


Fall 2025

Monte-Carlo vs TD-0: Spectrum

Temporal Difference



Monte-Carlo




TD-0 and Monte-Carlo are two extreme sides of a spectrum

- In TD-0, we use sample trajectory only for **one step**
- In Monte-Carlo, we use the **complete sample trajectory** for each state

Can we draw a solution between these two extreme points?

First Solution: Bootstrapping with More Steps

A primary approach to find such a **balanced** solution is to *extend the idea of bootstrapping to a larger number of steps*



- + How can we do it?
 - Well! We could simply expand the recursive property of value function

n-Bootstrapping

Looking at a sample return, we can simply write

$$\begin{aligned}
 G_t &= R_{t+1} + \gamma G_{t+1} \\
 &= R_{t+1} + \gamma R_{t+2} + \gamma^2 G_{t+2} \\
 &\quad \vdots \\
 &= R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{\textcolor{blue}{n}} R_{t+1+\textcolor{blue}{n}} \gamma^{\textcolor{blue}{n}+1} G_{t+1+\textcolor{blue}{n}}
 \end{aligned}$$

Now, assuming that *at time step t we are at state s, i.e.,*

$$S_t = s, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1} \xrightarrow{R_{t+2}} \cdots \xrightarrow{R_{t+1+n}} S_{t+1+n}$$

we can bootstrap over a longer part of sample trajectory

n-Bootstrapping

$$S_t = s, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1} \xrightarrow{R_{t+2}} \dots \xrightarrow{R_{t+n}} S_{t+n}$$

In this trajectory, we can expand Bellman equation with deeper recursion

$$v_\pi(s) = \sum_{i=0}^n \gamma^i \mathbb{E}\{R_{t+i+1}|s\} + \gamma^{n+1} \mathbb{E}_\pi\{v_\pi(S_{t+n+1})|s\}$$

So, if we have K sample trajectory, we can estimate value of state s by n steps of bootstrapping, i.e.,

$$\hat{v}_\pi(s) = \frac{1}{K} \sum_{k=1}^K \underbrace{\left(\sum_{i=0}^n \gamma^i R_{t+i+1}[k] + \hat{v}_\pi(S_{t+n+1}[k]) \right)}_{\text{computed on sample } k}$$

n -Bootstrapping \equiv TD- n

Let's formulate this approach: for a given *sample trajectory* and value function estimator $\hat{v}_\pi(\cdot)$, we define *n -bootstrapping return* at time t as

$$G_t^n = \sum_{i=0}^n \gamma^i R_{t+i+1} + \gamma^{n+1} \hat{v}_\pi(S_{t+n+1})$$

Given that the *sample trajectory* was started at state $S_t = s$, we can use *online averaging* and update the value estimator of state s as


$$\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + \alpha(G_t^n - \hat{v}_\pi(s))$$

This is what we call TD- n method of learning

This is obviously more general than TD-0! In TD-0, we had simply $n = 0$


Backup Diagram: Sampling with n -Bootstrapping

With *n*-bootstrapping we sample ($n + 1$)-step trajectories from action-state tree of environment




Backup Diagram: Sampling with n -Bootstrapping

With n -bootstrapping we sample $(n + 1)$ -step trajectories from action-state tree of environment



Backup Diagram: Sampling with n -Bootstrapping

With n -bootstrapping we sample $(n + 1)$ -step trajectories from action-state tree of environment



TD- n : Policy Evaluation

We can extend our TD-0 evaluation algorithm to TD- n

`TDn_Eval(π):`

- 1: Initiate estimator of value as $\hat{v}_\pi(s) = 0$ for all states
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a random state S_0 and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a terminal stated or some terminating T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - n - 1$ **do**
- 6: Compute $G = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n+1} v_\pi(S_{t+n+1})$
- 7: Update as $\hat{v}_\pi(S_t) \leftarrow \hat{v}_\pi(S_t) + \alpha(G - \hat{v}_\pi(S_t))$
- 8: **end for**
- 9: **end for**

★

TD- n : Policy Q-Evaluation

Same-wise, we can extend our algorithm for action-value computation

`TDn_QEval(π):`

- 1: Initiate estimator of value as $\hat{q}_\pi(s, a) = 0$ for all states
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a random state-action pair (S_0, A_0) and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a terminal stated or some terminating T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - n - 1$ **do**
- 6: Compute $G = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n+1} v_\pi(S_{t+n+1})$
- 7: Update as $\hat{q}_\pi(S_t, A_t) \leftarrow \hat{q}_\pi(S_t, A_t) + \alpha(G - \hat{q}_\pi(S_t, A_t))$
- 8: **end for**
- 9: **end for**

★

TD- ∞ : Going Back to Monte-Carlo

- + You told us that we look for a solution between TD-0 and Monte-Carlo! I see TD-0 is TD- n with $n = 0$, but where does Monte-Carlo stand?
- Well! You may see already that Monte-Carlo is TD- ∞

Say the environment is episodic: we can say that if we bootstrap very deep; then, at some point we hit a terminal state, i.e.,

$$\lim_{n \rightarrow \infty} v_\pi(S_{t+n+1}) = 0$$

This concludes that at any time t in an episodic environment

$$G_t^\infty = \sum_{i=0}^{\infty} \gamma^i R_{t+i+1} = G_t$$

and hence TD- ∞ will update its estimator by


$$\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + \alpha(G_t - \hat{v}_\pi(s))$$

TD- n : A Discrete Spectrum

We can look at TD- n as a *discrete* spectrum between Monte-Carlo and TD-0




Example: Dummy Grid World with Random Walk



Once you got home, try to get back to our **dummy world** and use **TD- n** method for **multiple n** to compute the values for **uniform random policy** ☺

Typical Behavior: Variation Against Depth

If you do some practice with random walk, you will see following curves for different choices of n



We observe a **minimum** against n : this is a **typical** behavior! Any illustration?

Averaging Different Depth

- + How **deep** we should then bootstrap?
- Well! We could try to find the **best** one **for each setting**, or we could **average** the result of **multiple bootstrapping depths**

For example, we can

```
1: Initiate ...
2: for episode = 1 : K do
3:   ...
4:   for t = 0 : T - 3 do
5:     Compute  $G_0 = R_{t+1} + \gamma v_\pi(S_{t+1})$ 
6:     Compute  $G_2 = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 v_\pi(S_{t+3})$ 
7:     Set  $G = (G_0 + G_2) / 2$ 
8:     Update as  $\hat{v}_\pi(S_t) \leftarrow \hat{v}_\pi(S_t) + \alpha(G - \hat{v}_\pi(S_t))$ 
9:   end for
10: end for
```

Averaging Different Depth: λ -Return


- + Why should that be a *good idea*?
- Because if the *good one* is within the average; then, it could *dominate* and *improve* estimation
- + Then, *which ones* we should take? We still have no clue!
- Let's take *all of them!* We can do it by *geometric weights*!

λ -Return

For any $0 \leq \lambda \leq 1$, the λ -return at time t over L steps is defined as

$$G_t^\lambda = (1 - \lambda) \sum_{n=0}^{L-1} \lambda^n G_t^n + \lambda^L G_t^L$$

Averaging Different Depth: λ -Return



For each state in the *sample trajectory*, we can

- Compute 0-bootstrapping return
 - ↳ Give it weight $1 - \lambda$
- Compute 1-bootstrapping return
 - ↳ Give it weight $(1 - \lambda) \lambda$
- ⋮
- Compute $(L - 1)$ -bootstrapping return
 - ↳ Give it weight $(1 - \lambda) \lambda^{L-1}$
- Compute L -bootstrapping return
 - ↳ Give it weight λ^L

TD $_{\lambda}$: Policy Evaluation

We can evaluate policy by averaging its λ -returns over multiple episodes

TD_Eval $_{\lambda}$ (π):

- 1: Initiate estimator of value as $\hat{v}_{\pi}(s^n) = 0$ for $n = 1 : N$
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a random state S_0 and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a terminal stated or some terminating T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: Set $G \leftarrow G_t^{\lambda}$ towards end of trajectory
- 7: Update as $\hat{v}_{\pi}(S_t) \leftarrow \hat{v}_{\pi}(S_t) + \alpha(G - \hat{v}_{\pi}(S_t))$
- 8: **end for**
- 9: **end for**

TD- λ : Special Cases

It is easy to see that $\lambda = 0$ and $\lambda = 1$ are again two extreme cases: assume we are dealing with an *episodic environment*

- with $\lambda = 0$

↳ All weights are zero but that of G_t^0 which is weighted one

$$G_t^\lambda = R_{t+1} + \gamma G_{t+1}$$

↳ this is basic bootstrapping: TD $_0$ is hence simply TD-0

- with $\lambda = 1$


↳ All weights are zero but that of G_t^{T-t} which is weighted one

$$G_t^\lambda = R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{T-t-1} R_T$$

↳ this basic Monte-Carlo: TD $_1$ is Monte-Carlo


TD $_{\lambda}$: A Continuous Spectrum

We can look at TD $_{\lambda}$ as a **continuous** spectrum between Monte-Carlo and TD-0



Typical Behavior: Variation Against λ


If you do some practice with random walk, you will see following curves for different choices of λ



We observe a **minimum** against λ : this is **analog** to TD- n behavior!

TD $_{\lambda}$: Weighting Function

Let's look at the weights decay in λ -return¹



¹From Sutton and Barto's book in Chapter 12

Assigning Credit for Future by Weighting

- + What is the **intuition** behind computing the λ -return with **those weights**?
- This is a very **valid question!** Let's see!

Let's look back at TD_λ approach: at each time t in the trajectory, we compute

$$G_t^\lambda = \sum_{n=1}^{T-t-1} w_t^n G_t^n$$


and then update the value of the **current state** S_t based on these **future returns**

the more **far away** this **future** is ↵ the **less** weight its return gets

In fact, we are **assigning credit** to our **current state** which will impact our **future** update: the more we go **forward** in time, the **less** this impact will be

TD $_{\lambda}$: Forward View

We can imagine this time advancement as *building impact* towards *future*²



But this approach is hard to be implemented *online*

we need to wait till *the end of episode* to compute the λ -returns!

²This figure is taken from Sutton and Barto's book in Chapter 12

TD $_{\lambda}$: Backward View

- + But, it does not seem to be another way for *credit assignment*?
- Well maybe we could apply the same idea **backward**
- + How can we do it?
- We can invoke the idea of *eligibility tracing*

Eligibility Tracing in Nutshell

At each time, when we update the value of a **state** in a **sample trajectory**, we also update the value of **previous states** we already met, with a weight **decaying** as we **go back in time**

Let's make an algorithm for that!

Eligibility Tracing

`ElgTrace($S_t, E(\cdot)$):`

- 1: Eligibility tracing function has N components, i.e., $E(s)$ for all states
- 2: **for** $n = 1 : N$ **do**
- 3: Update $E(s^n) \leftarrow \gamma\lambda E(s^n) \leftarrow \leftarrow$ choosing $\gamma\lambda$ for equivalency to forward view
- 4: **end for**
- 5: Update $E(S_t) \leftarrow E(S_t) + 1$

Say we initiate $E(s) = 0$ for all states and get to the following trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

Assume that we set $\gamma\lambda = 0.1$; then we have

- ① At $t = 0$, we only change $E(S_0) = 1 \leftarrow$ fresh memory \equiv high impact
- ② At $t = 1$, we change $E(S_0) = 0.1$ and $E(S_1) = 1$

Eligibility Tracing

`ElgTrace($S_t, E(\cdot)$):`

- 1: Eligibility tracing function has N components, i.e., $E(s)$ for all states
- 2: **for** $n = 1 : N$ **do**
- 3: Update $E(s^n) \leftarrow \gamma\lambda E(s^n) + \dots$ choosing $\gamma\lambda$ for equivalency to forward view
- 4: **end for**
- 5: Update $E(S_t) \leftarrow E(S_t) + 1$

Say we initiate $E(s) = 0$ for all states and get to the following trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

Now say that we see at $t = 2$ the same state as in $t = 0$, i.e., $S_0 = S_2$

- ③ At $t = 2$, we change $E(S_1) = 0.1$ and $E(S_0) \leftarrow 0.1E(S_0) + 1 = 1.01$

Updating with Eligibility Tracing: Intuition

- + What is the **use** of this algorithm?
- We can simply update **all previous states** each time t weighted by **eligibility traces**

We can 0-bootstrap at each time, i.e., compute **error** as

$$\Delta_t = \underbrace{R_{t+1} + \gamma \hat{v}_\pi(S_{t+1}) - \hat{v}_\pi(S_t)}_{G_t^0}$$

and update any state $s = S_0, \dots, S_t$ that has non-zero trace of eligibility as

$$\hat{v}_\pi(s) \leftarrow \hat{v}_\pi(s) + \alpha \Delta_t E_t(s)$$

with $E_t(s)$ denoting the **eligibility trace** that we have updated up to time t

TD $_{\lambda}$ vs Eligibility Tracing

- + Is there any **concrete reason** beside **simple intuition** that this is a good idea?
- We can actually see that this is an online form of basic TD $_{\lambda}$

In fact, we can show by telescopic sum that

$$\sum_{t=0}^{T-1} \Delta_t E_t(s) = \sum_{t=0}^{T-1} \left(G_t^{\lambda} - \hat{v}_{\pi}(S_t) \right) \mathbf{1}\{S_t = s\}$$

This means that at the end of episode, we are updating the same!

- If we set $\lambda = 0$ in the **eligibility tracing**
 - ↳ all **eligibility trace** remains 0 for all states: only we have $E(S_t) = 1$
 - ↳ we are back to TD-0 as it was with TD $_0$
- If we set $\lambda = 1$ in the **eligibility tracing**
 - ↳ we are back to **Monte-Carlo** approach as in TD $_1$

TD $_{\lambda}$ with Eligibility Tracing: Policy Evaluation

We can evaluate policy by averaging its λ -returns over multiple episodes

ElgTD_Eval $_{\lambda}(\pi)$:


- 1: Initiate value estimator and eligibility traces as $\hat{v}_{\pi}(s) = 0$ and $E(s) = 0$ for all s
- 2: **for** episode = 1 : K **do**
- 3: Initiate with a random state S_0 and act via policy $\pi(a|s)$
- 4: Sample a trajectory until either a terminal stated or some terminating T

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: $E(\cdot) \leftarrow \text{ElgTrace}(S_t, E(\cdot))$
- 7: Compute $G = R_{t+1} + \gamma \hat{v}_{\pi}(S_{t+1})$ and find error $\Delta = G - \hat{v}_{\pi}(S_t)$
- 8: **for** $n = 1 : N$ **do**
- 9: Update $\hat{v}_{\pi}(s^n) \leftarrow \hat{v}_{\pi} + \alpha E(s^n) \Delta$
- 10: **end for**
- 11: **end for**
- 12: **end for**

Eligibility Tracing: Backward View

We can imagine eligibility tracing as *backward assignment of credits*³



Now we can *update values* each time and don't need to wait till *end of episode!*

³This figure is taken from Sutton and Barto's book in Chapter 12