

Reinforcement Learning

Chapter 4: Function Approximation

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Review: Where Are We Now?

Model-Based RL

Bellman Equation

value iteration

policy iteration

Model-free RL

on-policy methods

temporal difference

Monte Carlo

SARSA

off-policy methods

Q-learning

We need to overcome one last challenge

we need to learn how to deal with large-scale problems

Tabular RL: What We Had Already

What we have studied up to now is usually called **tabular RL**

- + Why we call it **tabular**?
 - Because we think of **value** and **action-value** function as a **table**

Let's consider a mode-free control loop with **finite number of states** and **actions**

- We initiate all **action-values** with zero
 - ↳ We initiate a **table of zeros**
- But, we do know that these **action-values** are some specific values
 - ↳ We are dealing with a **table of unknowns**
- We try **estimating** them from samples

Tabular RL: Schematic

	a^1	a^2	...	a^M
s^1	$\hat{q}_\pi(s^1, a^1)$	$\hat{q}_\pi(s^1, a^2)$		
s^2				
\vdots				
s^N				$\hat{q}_\pi(s^N, a^M)$

Q-table

s^1	$\hat{v}_\pi(s^1)$
s^2	$\hat{v}_\pi(s^2)$
\vdots	\vdots
s^N	$\hat{v}_\pi(s^N)$

value-table

Computational Complexity of Tabular RL

A tabular approach needs estimation of NM values from samples

- Say we use Monte-Carlo

- ↳ We need to visit all the state-action pairs enough times
 - ↳ Say enough is C for us; then, we need

$$\# \text{ sample pairs in all episodes} \approx CNM$$

- Same thing with temporal difference

Moral of Story

In tabular RL methods, the number of required sample interactions with the environment scales with number of states and number of actions

Complexity Examples: Backgammon




There are roughly 10^{20} possible states for backgammon

- Say we can get over with each state-action pair and update in only 1 Picosec = 10^{-12} sec; then, we need

$$\# \text{ time} \approx 10^8 \text{ CM sec} \approx 3.2 \text{ CM years}$$

Complexity Examples: Chess



In the fun part of Assignment 1, you saw that Shannon found out about roughly 10^{120} possible states for chess, and later on some people came out with some approximations for legal positions: let's take Tromp's number, i.e., $\approx 4 \times 10^{44}$

- Say again we need only 1 Picosec per step; then, we need

$$\# \text{ time} \approx 4 \times 10^{32} \text{ CM sec} \approx 12 \times 10^{24} \text{ CM years}$$

Milky Way is about **13.6 billion years old!**

Complexity Examples: Game Go




There are roughly 10^{170} possible states for board game Go

- *Say again we need only 1 picosec*

time \approx do we really need computing it? 😊

Computational Complexity: Continuous State Space



We saw the Cart-Pole problem: if we want to put it into a tabular form

- We need to put it into a **table**
 - ↳ We have to make a **discrete grid** of states
 - ↳ Say we have L state parameters and we grid each one with B bins
 - ↳ We have B^L grid points in total
- This grows **exponentially** large in number of state parameters!

Computational Complexity of Tabular RL

In most interesting problems, we are dealing with **scaling problem**

- *We have an exponentially large number of states*
 - *We cannot even visit a subset of them!*
 - *Our tabular RL algorithms will not give us any meaningful results!*
-
- + *What then?! Should we always play 4×4 Frozen Lake with RL?!*
 - **No!** We try to **approximate our table** from our limited observations

Approximating Value Function

Let's start with a simple case: say we want to evaluate a policy, i.e.,

We are given with a policy π and want to find some estimate $\hat{v}_\pi(\cdot)$

In the *tabular* RL, we assume that

$$v_\pi(s^1) = v^1 \quad \dots \quad v_\pi(s^N) = v^N$$

for some unknown v^1, \dots, v^N and try to estimate them

Approximating Value Function

Let us now approach the problem differently: we assume that

$$v_{\pi}(s) = f(s, w)$$

for some weights in w and a known function $f()$

- $f()$ is a **approximation model** that we assume for the value function
 - w contains a set of **learnable parameters**
-
- + How on earth we know such a thing?!
 - We don't really know it! We just assume it; however, sometimes it really makes sense

Approximating Value Function

If we assume an **approximation model**: we use our observations to find weight vector w^* that fits this **approximator** best to our observations. We then set

$$\hat{v}_\pi(s) = f(s, w^*)$$

- + How is it better than tabular then?!
 - Well! For lots of reasons
- ① We can use every single sample to update the estimate for all states
 - ↳ We use samples to fit w : this impacts on the whole value function
- ② We get some **updating** estimates for states that have not been visited
 - ↳ If we get the right w : we get the estimator for **all** states
- ③ We can capture the impact of one state on the others
 - ↳ If we update w after seeing S_t : we change estimated values of **all** states

Approximating Value Function

- + But how can we find such approximators?!
- There are various types of them
 - Linear function approximators
 - Deep neural networks (DNNs)
 - Eigen-transforms, e.g., Fourier or Wavelets
 - ...

In this course, we are focusing only on *parametric approximators* which include

- Linear function approximators
- DNNs

Because they are *differentiable*: we will see why this is important!

But before using them, let's look at them and understand how they *work*

Function Approximation: Formulation

In function approximation, we have

- A set of sample **input**-**outputs** of a function that we **do not know**

$$\mathbb{D} = \{(\mathbf{x}_i, \mathbf{y}_i) : i = 1, \dots, I\}$$

- \mathbf{x}_i is the **input** and \mathbf{y}_i is the **output**
- Let's denote the unknown function with $g(\cdot)$, i.e., $\mathbf{y}_i = g(\mathbf{x}_i)$
- We assume an **approximating model** for this **unknown** function

$$\hat{\mathbf{y}} = \mathbf{f}(\mathbf{x} | \mathbf{w})$$


- In this **approximator** \mathbf{w} is learnable, i.e., we are free to tune it as we wish
- We treat output of this **approximator** as estimate of function outputs

Example: Linear Function Approximation

A simple example of a function approximator is the *linear approximator*


$$\hat{y} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

- \mathbf{W} is the matrix of weights
- \mathbf{b} is the vector of biases



Example: Deep Neural Networks

A deep neural network is also a parameterized approximator



Example: Deep Neural Networks

- + Why should we rely on these approximators?
- They are known to be very powerful

Universal Approximation Theorem

Given a family for neural networks: for any function g from its function space, there exists a sequence of configurations that approximates g arbitrarily precise

Function Approximation: Formulation

With **parametric** approximators: we want to find weight \mathbf{w} so that we have

- a **good approximator**: it fits best the **dataset**, i.e.,

$$\hat{\mathbf{y}}_i = \mathbf{f}(\mathbf{x}_i | \mathbf{w}) \approx \mathbf{y}_i$$

↳ We need to define a notion for \approx

- an **approximator** that **generalizes**: if we get a new sample input \mathbf{x}_{new} , we somehow can make sure that

$$\hat{\mathbf{y}}_{\text{new}} = \mathbf{f}(\mathbf{x}_{\text{new}} | \mathbf{w}) \approx g(\mathbf{x}_{\text{new}})$$


↳ We should find a way to check this

Visualizing Function Approximation

Let's assume one-dimensional inputs and outputs: *this assumption helps us visualize the function approximator*

$$y = f(\textcolor{green}{x} | \mathbf{w})$$

With scalar input, we can visualize the model as



As *learnable* parameters change, approximator sketches different functions

Training: Empirical Risk Minimization

Empirical Risk

Let \mathbf{w} includes all learnable parameters, and the dataset be

$$\mathbb{D} = \{(\mathbf{x}_i, \mathbf{y}_i) : i = 1, \dots, I\}$$

for loss function \mathcal{L} , the empirical risk is defined as

$$\hat{R}(\mathbf{w}) = \frac{1}{I} \sum_{i=1}^I \mathcal{L}(f(\mathbf{x}_i | \mathbf{w}), \mathbf{y}_i)$$

The training is performed by minimizing the empirical risk

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \hat{R}(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \frac{1}{I} \sum_{i=1}^I \mathcal{L}(f(\mathbf{x}_i | \mathbf{w}), \mathbf{y}_i) \quad (\text{Training})$$

Gradient Descent

Let's use **gradient descent** for function approximation: we want to minimize

$$\hat{R}(\mathbf{w}) = \frac{1}{I} \sum_{i=1}^I \mathcal{L}(f(\mathbf{x}_i | \mathbf{w}), \mathbf{y}_i)$$

`GradientDescent()`:

- 1: Initiate with some initial $\mathbf{w}^{(0)}$ and set a learning rate η
- 2: **while** weights not converged **do**
- 3: **for** $i = 1, \dots, I$ **do**
- 4: Compute gradient for $\nabla_i = \nabla \mathcal{L}\left(f\left(\mathbf{x}_i | \mathbf{w}^{(t-1)}\right), \mathbf{y}_i\right)$
- 5: Update gradient as $\nabla \hat{R}(\mathbf{w}^{(t-1)}) \leftarrow \nabla \hat{R}(\mathbf{w}^{(t-1)}) + \nabla_i / I$
- 6: **end for**
- 7: Update weights as $\mathbf{w}^{(t)} \leftarrow \mathbf{w}^{(t-1)} - \eta \nabla \hat{R}(\mathbf{w}^{(t-1)})$
- 8: **end while**

We call this form of gradient descent **full-batch** which can be too complex

Stochastic Gradient Descent

We use **gradient descent** for function approximation: we want to minimize

$$\hat{R}(\mathbf{w}) = \frac{1}{I} \sum_{i=1}^I \mathcal{L}(f(\mathbf{x}_i|\mathbf{w}), \mathbf{y}_i)$$

SGD():

- 1: Initiate with some initial $\mathbf{w}^{(0)}$ and set a learning rate η
- 2: **while** weights not converged **do**
- 3: **for** a random subset of batch $b = 1, \dots, B$ **do**
- 4: Compute gradient for $\nabla_b = \nabla \mathcal{L}\left(f\left(\mathbf{x}_b|\mathbf{w}^{(t-1)}\right), \mathbf{y}_b\right)$
- 5: Update gradient as $\nabla \hat{R}(\mathbf{w}^{(t-1)}) \leftarrow \nabla \hat{R}(\mathbf{w}^{(t-1)}) + \nabla_b / I$
- 6: **end for**
- 7: Update weights as $\mathbf{w}^{(t)} \leftarrow \mathbf{w}^{(t-1)} - \eta \nabla \hat{R}(\mathbf{w}^{(t-1)})$ **← unbiased estimator**
- 8: **end while**

This is what we call **stochastic mini-batch** gradient descent