

Reinforcement Learning

Chapter 4: Function Approximation

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Back to Model-free RL

We now get back to model-free RL: *this time however*

we use a parameterized estimator

This means that

- to estimate the **value** function, we use

$$\hat{v}_{\mathbf{w}}(s)$$

- to estimate the **action-value** function, we use

$$\hat{q}_{\mathbf{w}}(s, a)$$

↳ These functions are **parameterized** by some \mathbf{w}

↳ They could be as **simple** as linear approximators or **advanced** like DNNs

Back to Model-free RL

- + But, we have in general ***prediction*** and ***control*** problems. In which one are we going to use function ***approximation***?
- Well, we can use in both

Recall


We have two major problems in ***model-free*** RL

- ***Prediction*** in which for a given policy π we evaluate values by sampling the environment
- ***Control*** in which after each interaction, we improve our policy aiming to converge to the ***optimal policy***

Let's start with the ***prediction***

Basic Prediction via Approximator

Say we want to **evaluate** the value function of a policy π : with **function approximation** we assume that the **value approximator** is




- We sample the environment by policy π
 - ↳ We draw sample trajectories using policy π

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- We want to use these samples to train our **approximation model**
 - ↳ We want to find w that approximates the best $v_\pi(\cdot)$

Building Simple Approximator

- + Well you have talked a lot about **approximators**, but can you give us a concrete example?!
- Sure! Let's look at **linear approximator**: say our approximator is a linear function of the input
- + Does it mean the following diagram?!



- + But it doesn't seem to work with only **one parameter**!
- Oops! We haven't yet defined the **feature representation of states**!

Feature Representation

In general, we could represent a particular state in *various* forms

Say we have N states s^1, \dots, s^N : we can represent them by a *scalar*, e.g.

$$s^n \mapsto \mathbf{x}(s^n) = n$$

Or a *one-sparse* vector, i.e.,

$$s^n \mapsto \mathbf{x}(s^n) = \begin{bmatrix} 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix} \leftarrow \text{entry } n$$

Feature Representation

Feature Representation of States

Feature representation maps each **state** into a vector of **features** that corresponds to that **state**, i.e.,

$$\mathbf{x}(\cdot) : \mathbb{S} \mapsto \mathbb{R}^J$$

for some integer J that is the **feature dimension**

Example: We can represent the feature by **tokenization**

$$s^n \mapsto \mathbf{x}(s^n) = \begin{bmatrix} 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix} \leftarrow \text{entry } n$$

Feature Representation

Feature Representation of States

Feature representation maps each **state** into a vector of features that corresponds to that **state**, i.e.,

$$\mathbf{x}(\cdot) : \mathcal{S} \mapsto \mathbb{R}^J$$

for some integer J that is the **feature dimension**

In practice, we use more **advanced problem-specific** features

- The state of a robot can be specified by its **geometric features**
 - ↳ its distance to the edges of the room, its direction, etc
- The state of a computer game is completely explained by its **frames**
 - ↳ we collect all frames from a state to another

Example: Mountain Car



Let's consider the famous *mountain car example*

- A car is stuck in a valley
 - ↳ it can *accelerate to left, accelerate to right, or do nothing*
- It observes its *velocity v and location x* on the horizontal axis
 - ↳ they are *continuous* variables

The goal is to train the car to get out of this valley as quick as possible

- The car is *rewarded by -1* after each unsuccessful trial

Example: Mountain Car



This car can be in **infinite** number of **states**; however, we can represent its state by its **velocity and location**


$$\mathbf{x}(s) = \begin{bmatrix} v \\ x \end{bmatrix}$$

This is much **more feasible** to work with!

Back to Simple Approximator

- + How is this **feature** related to our discussion on **function approximation**?!
 - Well! Approximation model maps the **features** to **value**

So, in our problem, we in fact assume that



Example: Linear Approximator

Linear approximation model maps the state features to its value as

$$v_w(s) = \mathbf{w}^T \mathbf{x}(s)$$

Example: Mountain Car




In mountain car example, we could have a two-dimensional linear approximator

$$\begin{aligned}v_{\mathbf{w}}(s) &= \mathbf{w}^T \mathbf{x}(s) = [w_1 \quad w_2] \begin{bmatrix} v \\ x \end{bmatrix} \\&= w_1 v + w_2 x\end{aligned}$$

Back to Simple Approximator

- + How is this feature related to our discussion on function approximation?!
- Well! Approximation model maps the features to value

So, in our problem, we in fact assume that



Example: Deep Approximator


Deep approximation model maps the state features to its value via a DNN

$$v_{\mathbf{w}}(s) = \text{DNN}(\mathbf{x}(s) | \mathbf{w})$$

Example: Mountain Car



We may also use a NN to map the **feature** to its value



Here, the weights are $\mathbf{w} = [w_1, \dots, w_6, b_1, b_2, b_3]^T$

Training Approximation Model

- + How can we train a given approximation model?
- Let's again get help from a **genie**

Assume that a **genie** could tell us the exact value of all states: in this case we want to find w such that for each s^n

$$v_w(s^n) \stackrel{!}{=} v_\pi(s^n) \rightsquigarrow |v_w(s^n) - v_\pi(s^n)| \stackrel{!}{=} 0$$

This is **equivalent** to say that

$$\frac{1}{N} \sum_{n=1}^N |v_w(s^n) - v_\pi(s^n)|^2 \stackrel{!}{=} 0$$

Training Approximation Model

But, we cannot necessarily find such \mathbf{w} : we could instead find

$$\mathbf{w}^* = \min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N |v_{\mathbf{w}}(s^n) - v_{\pi}(s^n)|^2$$

One may also think about a more general weighted average: we can assume that under policy π each state s^n happens with a probability $p_{\pi}(s^n)$ and write

$$\begin{aligned}\mathbf{w}^* &= \min_{\mathbf{w}} \sum_{n=1}^N p_{\pi}(s^n) |v_{\mathbf{w}}(s^n) - v_{\pi}(s^n)|^2 \\ &= \min_{\mathbf{w}} \mathbb{E}_{\pi} \{ |v_{\mathbf{w}}(S) - v_{\pi}(S)|^2 \}\end{aligned}$$

We train by minimizing residual sum of squares \equiv least-squares (LS) method

LS Training: Gradient Descent

We use gradient descent to solve this problem: *we are minimizing the risk*

$$\mathcal{L}(\mathbf{w}) = \mathbb{E}_{\pi} \left\{ |v_{\mathbf{w}}(S) - v_{\pi}(S)|^2 \right\}$$

`GradientDescent()`:

- 1: Initiate with some initial \mathbf{w} and learning rate η
- 2: **while** weights not converged **do**
- 3: Compute gradient $\nabla \mathcal{L}(\mathbf{w})$
- 4: Update weights as $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla \mathcal{L}(\mathbf{w}^{t-1})$
- 5: **end while**

Let's compute the gradient: we can use the chain rule

$$\begin{aligned}\nabla \mathcal{L}(\mathbf{w}) &= \frac{\partial \mathcal{L}}{\partial v_{\mathbf{w}}(S)} \nabla v_{\mathbf{w}}(S) \\ &= 2 \mathbb{E}_{\pi} \{ (v_{\mathbf{w}}(S) - v_{\pi}(S)) \nabla v_{\mathbf{w}}(S) \}\end{aligned}$$

LS Training: Gradient Descent

Now we set the learning rate to $\eta = 0.5\alpha$ for some α ; then, we have

$$\begin{aligned}\mathbf{w}^{(t)} &\leftarrow \mathbf{w}^{(t-1)} - \alpha \mathbb{E}_\pi \{(v_\mathbf{w}(S) - v_\pi(S)) \nabla v_\mathbf{w}(S)\} \\ &\leftarrow \mathbf{w}^{(t-1)} + \alpha \mathbb{E}_\pi \{(v_\pi(S) - v_\mathbf{w}(S)) \nabla v_\mathbf{w}(S)\}\end{aligned}$$

So, we the gradient descent based evaluation reduces to

GD_Eval():

- 1: Initiate with some initial $\mathbf{w}^{(0)}$ and learning rate η
- 2: **while** weights not converged **do**
- 3: Update weights as $\mathbf{w} \leftarrow \mathbf{w} + \alpha \mathbb{E}_\pi \{(v_\pi(S) - v_\mathbf{w}(S)) \nabla v_\mathbf{w}(S)\}$
- 4: **end while**

- + That's nice! But, how in earth we know $v_\pi(S)$!?
- Well! We may use what we learned in model-free RL again!

LS Training via Monte Carlo

The key challenge is to find out *an estimator for*

$$\mathbb{E}_{\pi} \{(v_{\pi}(S) - v_w(S)) \nabla v_w(S)\}$$

We can do it by Monte Carlo: *say we have an episodic environment with terminal state and sampled an episode*

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

We can convert this trajectory into

$$(S_0, G_0) \rightarrow (S_1, G_1) \rightarrow \dots \rightarrow (S_{T-1}, G_{T-1})$$

with G_t being the sample return, i.e.,

$$G_t = \sum_{i=0}^{T-t-1} \gamma^i R_{t+1+i}$$

LS Training via Monte Carlo

The key challenge is to find out *an estimator for*

$$\mathbb{E}_\pi \{(v_\pi(S) - v_w(S)) \nabla v_w(S)\}$$

We do know that G_t is an estimator of $v_\pi(S_t)$: so we could say

$$\mathbb{E}_\pi \{(v_\pi(S) - v_w(S)) \nabla v_w(S)\} \approx \frac{1}{T} \sum_{t=0}^{T-1} (G_t - v_w(S_t)) \nabla v_w(S_t)$$

We can compute this one from our **observations!**

LS Training via Monte Carlo

We can then write this approximate gradient descent approach as

`GD_MC_Eval()`:

- 1: Initiate with some initial \mathbf{w} and learning rate α
- 2: **while** weights not converged **do**
- 3: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 4: Compute the sequence

$$(S_0, G_0) \rightarrow (S_1, G_1) \rightarrow \dots \rightarrow (S_{T-1}, G_{T-1})$$

- 5: Update weights as

$$\mathbf{w} \leftarrow \mathbf{w} + \frac{\alpha}{T} \sum_{t=0}^{T-1} (G_t - v_{\mathbf{w}}(S_t)) \nabla v_{\mathbf{w}}(S_t)$$

- 6: **end while**

LS Training via Monte Carlo

This is not practical to wait for exact convergence: we try couple of episodes

GD_MC_Eval():

- 1: Initiate with some initial \mathbf{w} and learning rate α
- 2: **for** episode $k = 1 : K$ **do**
- 3: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 4: Compute the sequence

$$(S_0, G_0) \rightarrow (S_1, G_1) \rightarrow \dots \rightarrow (S_{T-1}, G_{T-1})$$

- 5: Update weights as

$$\mathbf{w} \leftarrow \mathbf{w} + \frac{\alpha}{T} \sum_{t=0}^{T-1} (G_t - v_{\mathbf{w}}(S_t)) \nabla v_{\mathbf{w}}(S_t)$$

- 6: **end for**

LS Training via Monte Carlo

- + But, should we really update *once an episode*?!
 - Not really!

We can look at this algorithm as batch training with the batch being

$$(S_0, G_0), (S_1, G_1), \dots, (S_{T-1}, G_{T-1})$$

We can also use mini-batches, and we can reduce the *mini-batch size* to 1: in this case, the estimator of the gradient will be a single sample, i.e.,

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha (G_t - v_{\mathbf{w}}(S_t)) \nabla v_{\mathbf{w}}(S_t)$$

LS Training via Monte Carlo

We can train the approximator via SGD and Monte Carlo

`SGD_MC_Eval()`:

- 1: Initiate with some initial \mathbf{w} and learning rate α
- 2: **for** episode $k = 1 : K$ **do**
- 3: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 4: Compute the sequence

$$(S_0, G_0) \rightarrow (S_1, G_1) \rightarrow \dots \rightarrow (S_{T-1}, G_{T-1})$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: Update weights as

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha (G_t - v_{\mathbf{w}}(S_t)) \nabla v_{\mathbf{w}}(S_t)$$

- 7: **end for**
- 8: **end for**

LS Training via Temporal Difference

- + So can't we use also TD to acquire an estimate?
- Sure!

We can evaluate an estimator by TD: say we sample

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

We can then compute an estimator of $v_\pi(S_t)$ via TD as

$$G_t^0 = R_{t+1} + \gamma v_w(S_{t+1})$$

LS Training via Temporal Difference

We can alternatively train the approximator via temporal difference

`SGD_TD_Eval()`:

- 1: Initiate with some initial w and learning rate α
- 2: **for** episode $k = 1 : K$ **do**
- 3: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 4: **for** $t = 0 : T - 1$ **do**
- 5: Update weights as

$$w \leftarrow w + \alpha (R_{t+1} + \gamma v_w(S_{t+1}) - v_w(S_t)) \nabla v_w(S_t)$$

- 6: **end for**
- 7: **end for**

You can extend it to TD- n or TD $_{\lambda}$ as well

Back to Tabular RL

- + These algorithm look like what we had before!
- Well! This is actually an **extended version** of it!

Let's consider a special case in which

- ① We use **tokenization** for feature representation

Tokenization

We represent the feature vector as

$$\mathbf{x}(s^n) = \begin{bmatrix} 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix} \leftarrow \text{entry } n = \mathbf{1}\{s = s^n\}$$

Back to Tabular RL

Let's consider a special case in which

- ① We use **tokenization** for feature representation
- ② We use a **linear** model for approximation

Linear Model

We approximate the value function via a linear transform

$$v_{\mathbf{w}}(s) = \mathbf{w}^T \mathbf{x}(s)$$

With linear model and tokenization, our estimate of value at state s^n is

$$\hat{v}_{\pi}(s^n) = v_{\mathbf{w}}(s^n) = \mathbf{w}^T \mathbf{1}\{s = s^n\} = w_n$$

Also we have

$$\nabla v_{\mathbf{w}}(s^n) = \mathbf{x}(s^n) = \mathbf{1}\{s = s^n\}$$

Back to Tabular RL

Let's consider a special case in which

- ① We use **tokenization** for feature representation
- ② We use a **linear** model for approximation

Let's now look at the SGD update with TD

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} + \alpha (R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t)) \nabla v_{\mathbf{w}}(S_t) \\ &\leftarrow \mathbf{w} + \alpha \left(R_{t+1} + \gamma \underbrace{v_{\mathbf{w}}(S_{t+1})}_{\hat{v}_{\pi}(S_{t+1})} - \underbrace{v_{\mathbf{w}}(S_t)}_{\hat{v}_{\pi}(S_t)} \right) \underbrace{\mathbf{x}(S_t)}_{\mathbf{1}\{s=S_t\}} \end{aligned}$$

Say $S_t = s^n$: then, it is only entry n which gets updated

$$w_n \leftarrow w_n + \alpha (R_{t+1} + \gamma \hat{v}_{\pi}(S_{t+1}) - \hat{v}_{\pi}(S_t))$$

Back to Tabular RL

Let's consider a special case in which

- ① We use **tokenization** for feature representation
- ② We use a **linear** model for approximation

Say $S_t = s^n$: then, it is only entry n which gets updated

$$w_n \leftarrow w_n + \alpha (R_{t+1} + \gamma \hat{v}_\pi(S_{t+1}) - \hat{v}_\pi(S_t))$$

and we know that $w_n = \hat{v}_\pi(s^n) = \hat{v}_\pi(S_t)$, so we can write

$$\hat{v}_\pi(S_t) \leftarrow \hat{v}_\pi(S_t) + \alpha (R_{t+1} + \gamma \hat{v}_\pi(S_{t+1}) - \hat{v}_\pi(S_t))$$

Bingo! This is the **tabular TD update!**

Back to Tabular RL

Moral of Story

Tabular RL is a special case of RL with function approximation when

- ① We use **tokenization** for feature representation
- ② We use a **linear** model for approximation

*So, we expect learning **better** when we go for other **feature representation and approximation models***

Training Action-Value Approximator

- + In practice however we always need **action-values!** Right?
- Well! We can simply extend everything to state-action pairs

Feature Representation of State-Actions


Feature representation maps each state-action pair into a vector of features that correspond to that state and action, i.e.,

$$\mathbf{x}(\cdot) : \mathcal{S} \times \mathcal{A} \mapsto \mathbb{R}^J$$

for some integer J that is the feature dimension


Action-Value Approximator: Form I

We can further consider an approximation model: *it maps the feature vector of each state-action pair (s, a) into its action-value*




Action-Value Approximator: Form II

We may stack this approximator for various actions: *we can look at the end-to-end setting as a general approximator*



Action-Value Approximator: Form II

We may stack this approximator for various actions: *maybe, we can consider a general approximation model in this case*



Let's make an agreement

$\hat{Q}_\pi(s, \mathbf{a})$ and $Q_{\mathbf{w}}(s, \mathbf{a})$ represent the complete **vector** of action-values
and $\hat{Q}_\pi(s, a)$ and $Q_{\mathbf{w}}(s, a)$ denote the **entry corresponding to a**

Training Approximation Model

The LS training can further be applied here: *with the help of genie, we train the action-value approximator as*

$$\mathbf{w}^* = \min_{\mathbf{w}} \mathbb{E}_{\pi} \left\{ |Q_{\mathbf{w}}(S, A) - Q_{\pi}(S, A)|^2 \right\}$$

which is iteratively solved via gradient descent using update rule

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \mathbb{E}_{\pi} \left\{ (Q_{\pi}(S, A) - Q_{\mathbf{w}}(S, A)) \nabla Q_{\mathbf{w}}(S, A) \right\}$$

Again, we could use Monte Carlo or TD to find an estimator of $Q_{\pi}(S, A)$

LS Training of Action-Value Approximator

Say we have sampled a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

We can convert this trajectory into a sequence of (S_t, A_t, G_t)

- Using Monte Carlo we can update by

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha(G_t - Q_{\mathbf{w}}(S_t, A_t)) \nabla Q_{\mathbf{w}}(S_t, A_t)$$

- Using TD we can update as

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha(R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - Q_{\mathbf{w}}(S_t, A_t)) \nabla Q_{\mathbf{w}}(S_t, A_t)$$

and we can find $v_{\mathbf{w}}(S_{t+1})$ as

$$v_{\mathbf{w}}(s) = \sum_{m=1}^M \pi(a^m|s) Q_{\mathbf{w}}(s, a^m)$$

LS Training of Action-Value Approximator

We can use all approaches we developed before

- only **difference** is that we replace the simple update by gradient descent
- it gets back to tabular RL if we use **tokenization** and **linear approximation**

Let's see a simple example: remember the backward view of TD_λ

- we **trace** the eligibility of each **state-action**
- we **propagate** the update to previous state-action pairs

ElgTrace($S_t, A_t, E(\cdot) | \lambda$):

- 1: Eligibility tracing function has NM components, i.e., $E(s, a)$ for all **state-action** pairs
- 2: **for** all state-action pairs (s, a) **do**
- 3: Update $E(s, a) \leftarrow \gamma \lambda E(s, a)$
- 4: **end for**
- 5: Update $E(S_t, A_t) \leftarrow E(S_t, A_t) + 1$

LS Training of Action-Value Approximator

`ElgTrace($S_t, A_t, E(\cdot) | \lambda$):`

- 1: Eligibility tracing function has NM components, i.e., $E(s, a)$ for all state-action pairs
- 2: for all state-action pairs (s, a) do
- 3: Update $E(s, a) \leftarrow \gamma \lambda E(s, a)$
- 4: end for
- 5: Update $E(S_t, A_t) \leftarrow E(S_t, A_t) + 1$

This was with tabular RL which uses **tokenization** and **linear approximation**: let's see if we can represent it in terms of approximation model components

- with **tokenization**, we have NM -dimensional feature

$$\mathbf{x}(s^n, a^m) = \mathbf{1}\{s = s^n, a = a^m\}$$

- with **linear** model, we have \mathbf{w} which is of the same dimension
- with **linear** model, we have

$$\nabla Q_{\mathbf{w}}(s, a) = \mathbf{x}(s, a)$$

LS Training of Action-Value Approximator

ElgTrace($S_t, A_t, E(\cdot) | \lambda$):

- 1: Eligibility tracing function has NM components, i.e., $E(s, a)$ for all state-action pairs
- 2: for all state-action pairs (s, a) do
- 3: Update $E(s, a) \leftarrow \gamma \lambda E(s, a)$
- 4: end for
- 5: Update $E(S_t, A_t) \leftarrow E(S_t, A_t) + 1$

Let's define a vector E_w which has the NM entries of $E(S_t, A_t)$

- we can write the update rule in line 4 as $E_w \leftarrow \gamma \lambda E_w$
- and the update rule in line 5 as

$$\begin{aligned} E_w &\leftarrow E_w + \mathbf{1}\{s = S_t, a = A_t\} \\ &\leftarrow E_w + \mathbf{x}(S_t, A_t) \\ &\leftarrow E_w + \nabla Q_w(S_t, A_t) \end{aligned}$$

LS Training of Action-Value Approximator

`ElgTrace($S_t, A_t, E(\cdot) | \lambda$):`

- 1: Eligibility tracing function has NM components, i.e., $E(s, a)$ for all state-action pairs
- 2: for all state-action pairs (s, a) do
- 3: Update $E(s, a) \leftarrow \gamma \lambda E(s, a)$
- 4: end for
- 5: Update $E(S_t, A_t) \leftarrow E(S_t, A_t) + 1$

Merging the two lines, we get into

$$E_w \leftarrow \gamma \lambda E_w + \nabla Q_w(S_t, A_t)$$

This is the more general form of eligibility tracing

- we can use it for any approximation
- our trace is of the size of the feature vector

LS Training via TD_λ

So, we could evaluate via training as

SGD_TD_QEval(λ):

- 1: Initiate with some initial w and learning rate α
 - 2: **for** episode $k = 1 : K$ **do**
 - 3: Sample a trajectory

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- ```

4: for $t = 0 : T - 1$ do
5: Compute $\Delta = R_{t+1} + \gamma v_w(S_{t+1}) - Q_w(S_t, A_t)$ # forward propagation
6: Compute $\nabla = \nabla Q_w(S_t, A_t)$ # backpropagation
7: $E_w \leftarrow \lambda \gamma E_w + \nabla$
8: Update weights as
 $w \leftarrow w + \alpha \Delta E_w$
9: end for
10: end for

```

## Example: Mountain Car



We can compare **tabular RL** against the one with **function approximation**

