

Reinforcement Learning

Chapter 5: RL via Policy Gradient

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Policy Gradient Theorem: Point of Departure

Policy Gradient Theorem

For a policy network with non-zero probabilities, the gradient of the average trajectory return is always given by

$$\nabla \mathcal{J}(\pi_{\theta}) = \mathbb{E}_{S \sim d_{\pi_{\theta}}, A|S \sim \pi_{\theta}} \{q_{\pi_{\theta}}(S, A) \nabla \log \pi_{\theta}(A|S)\}$$

Policy gradient theorem is the base of

Policy Gradient Methods \equiv PGM

It gives a **feasible approach** for **training** a policy network; however, depending on how we use it we can end up with various PGMs

PGMs in Nutshell

Policy Gradient Theorem

For a policy network with non-zero probabilities, the gradient of the average trajectory return is always given by

$$\nabla \mathcal{J}(\pi_{\theta}) = \mathbb{E}_{S \sim d_{\pi_{\theta}}, A | S \sim \pi_{\theta}} \{ q_{\pi_{\theta}}(S, A) \nabla \log \pi_{\theta}(A|S) \}$$

PGMs can roughly divided into three classes

- ① **Vanilla** PGM estimates $q_{\pi_{\theta}}(S, A)$ and $\nabla \log \pi_{\theta}(A|S)$ via **sampling**
- ② **Baseline** PGM that reduces estimation variance by temporal **unbiasing trick**
- ③ **Trust region** PGM enables **reuse** of older samples to improve **efficiency**
 - ↳ This is what we learn in the next section of this chapter

We are going through them in the same order!

Vanilla PGM: Basic SGD

Vanilla PGM is pretty straightforward: *sample environment with a trajectory and train policy network via SGD using result of policy gradient theorem*

- Use SGD to update θ in each time, i.e., update as $\theta \leftarrow \theta + \alpha \nabla \mathcal{J}(\pi_\theta)$
- Compute gradient using policy gradient theorem

$$\nabla \mathcal{J}(\pi_\theta) = \mathbb{E}_{S \sim d_{\pi_\theta}, A | S \sim \pi_\theta} \{ q_{\pi_\theta}(S, A) \nabla \log \pi_\theta(A|S) \}$$

- Estimate the gradient via *individual samples*, i.e.,

$$\hat{\nabla} \mathcal{J}(\pi_\theta) = Q_t \nabla \log \pi_\theta(A_t|S_t)$$

where Q_t is an estimator of *action-value* at pair (S_t, A_t) in sample

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

Vanilla PGM: Generic Form

VanillaPGM():

- 1: *Initiate with θ and learning rate α*
- 2: *Use a Q-estimator $\text{QEst}()$*
- 3: **while interacting do**
- 4: *Sample the environment with policy π_θ*

$$S_0, A_0 \xrightarrow{R_1} S_{t+1}, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: *Set $Q_t = \text{QEst}(S_t, A_t)$*
- 7: *Update policy network $\theta \leftarrow \theta + \alpha Q_t \nabla \log \pi_\theta(A_t | S_t)$*
- 8: **end for**
- 9: **end while**

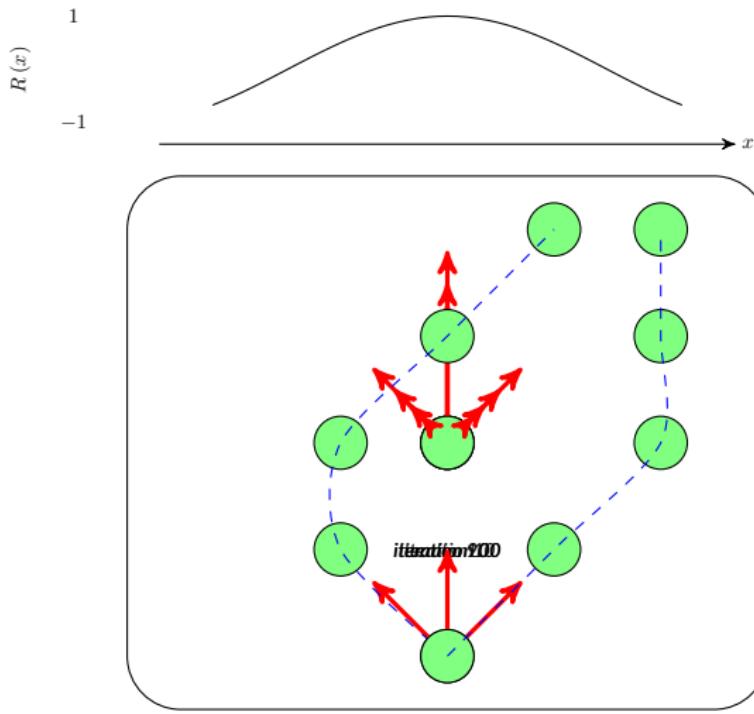
Revisiting REINFORCE

It is easy to see that REINFORCE() is a **vanilla PGM**: here, we set estimator of **action-value** to be

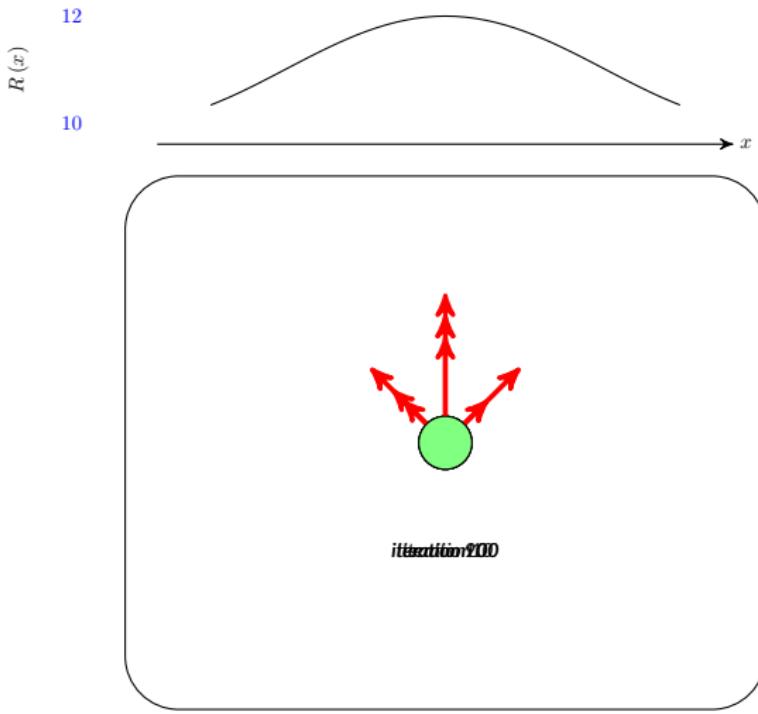
$$Q_t = G_t = \sum_{i=t}^{T-1} \gamma^i R_{i+1}$$

- + But in our initial derivation, we saw derived G_0 instead of G_t !
- Well! If we replace in the policy gradient theorem, we could see that it would be still an estimator if we replace G_t with G_0
- + So, we have many estimators! How can we choose among them?!
- This is what we do in **baseline PGM**
- + What about using TD to estimate **action-values**?
- We could do it! But there will be a bit of complications. We will see it soon!

Example: Controlling Moving Particle - Case I



Example: Controlling Moving Particle - Case II



Vanilla PGM: Bias Issue

This is a **crucial observation**: with a simple shift in value, **vanilla** PGM slows significantly in convergence!

VanillaPGM():

- 1: Initiate with θ and learning rate α
- 2: Use a **Q-estimator** QEst()
- 3: **while interacting do**
- 4: Sample the environment with policy π_θ

$$S_0, A_0 \xrightarrow{R_1} S_{t+1}, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: Set $Q_t = \text{QEst}(S_t, A_t)$
- 7: Update as $\theta \leftarrow \theta + \alpha [Q_t] \nabla \log \pi_\theta(A_t | S_t)$ ← here is the trouble
- 8: **end for**
- 9: **end while**

Vanilla PGM: Bias Issue

This is a **crucial observation**: with a simple shift in value, **vanilla** PGM slows **significantly** in convergence!

Let's look at this update rule

$$\theta \leftarrow \theta + \alpha Q_t \nabla \log \pi_\theta (A_t | S_t)$$

With larger values, Q_t becomes larger, hence

- if $\nabla \log \pi_\theta (A_t | S_t)$ becomes a **small** positive
 - ↳ θ increases **largely**
- if $\nabla \log \pi_\theta (A_t | S_t)$ becomes a **small** negative
 - ↳ θ drops **largely**

We need to have Q_t concentrated around **zero**

Vanilla PGM: Bias Issue

We need to have Q_t concentrated around zero

- + But, wait a moment! We derived this expression from **policy gradient theorem**! If we change Q_t to something else, we are deviating from it!
- Well! This is **not necessarily** true!

Let's try an experiment: in the gradient term given by policy gradient algorithm, we replace the **action-value** term with a **shifted one**, i.e., replace $q_{\pi_\theta}(S, A)$ with

$$q'_{\pi_\theta}(S, A) = q_{\pi_\theta}(S, A) - u(S)$$

The term $u(S)$ can change with **state**, but it is **fixed** in terms of **actions**

Unbiasing Policy Gradient

Replacing $q'_{\pi_\theta}(S, A)$ into the gradient expression, we have

$$\begin{aligned}
 \mathcal{E}(\pi_\theta) &= \mathbb{E}_{S \sim d_{\pi_\theta}, A|S \sim \pi_\theta} \{ q'_{\pi_\theta}(S, A) \nabla \log \pi_\theta(A|S) \} \\
 &= \mathbb{E}_{S \sim d_{\pi_\theta}, A|S \sim \pi_\theta} \{ (q_{\pi_\theta}(S, A) - u(S)) \nabla \log \pi_\theta(A|S) \} \\
 &= \mathbb{E}_{S \sim d_{\pi_\theta}} \{ \mathbb{E}_{\pi_\theta} \{ (q_{\pi_\theta}(S, A) - u(S)) \nabla \log \pi_\theta(A|S) | S \} \} \\
 &= \underbrace{\mathbb{E}_{S \sim d_{\pi_\theta}} \{ \mathbb{E}_{\pi_\theta} \{ q_{\pi_\theta}(S, A) \nabla \log \pi_\theta(A|S) | S \} \}}_{\nabla \mathcal{J}(\pi_\theta)} \\
 &\quad - \underbrace{\mathbb{E}_{S \sim d_{\pi_\theta}} \{ u(S) \mathbb{E}_{\pi_\theta} \{ \nabla \log \pi_\theta(A|S) | S \} \}}_?
 \end{aligned}$$

So, we have

$$\mathcal{E}(\pi_\theta) = \nabla \mathcal{J}(\pi_\theta) - \mathbb{E}_{S \sim d_{\pi_\theta}} \{ u(S) \mathbb{E}_{\pi_\theta} \{ \nabla \log \pi_\theta(A|S) | S \} \}$$

To compute the second term, we can use a *simple trick*

Gradient Averaging Trick

Assume $X \sim p_{\theta}(x)$: X is distributed by a distribution that is parameterized by some θ . We can then write

$$\begin{aligned}\mathbb{E}_{p_{\theta}} \{ \nabla_{\theta} \log p_{\theta}(X) \} &= \mathbb{E}_{p_{\theta}} \left\{ \frac{\nabla p_{\theta}(X)}{p_{\theta}(X)} \right\} = \int_x \frac{\nabla p_{\theta}(x)}{p_{\theta}(x)} p_{\theta}(x) \\ &= \int_x \nabla p_{\theta}(x) = \nabla \underbrace{\int_x p_{\theta}(x)}_1 = \nabla 1 = 0\end{aligned}$$

Lemma: Gradient Averaging

For any parameterized distribution $p_{\theta}(x)$, we have

$$\mathbb{E}_{p_{\theta}} \{ \nabla_{\theta} \log p_{\theta}(X) \} = 0$$

Unbiasing Policy Gradient

This concludes that

$$\begin{aligned}\mathcal{E}(\pi_{\theta}) &= \nabla \mathcal{J}(\pi_{\theta}) - \mathbb{E}_{S \sim d_{\pi_{\theta}}} \left\{ u(S) \underbrace{\mathbb{E}_{\pi_{\theta}} \{ \nabla \log \pi_{\theta}(A|S) | S \}}_0 \right\} \\ &= \nabla \mathcal{J}(\pi_{\theta})\end{aligned}$$

*In other words: we can add any term that is **fixed in terms of actions** to our **value estimator** without any harm to policy gradient theorem*

- This **fixed term** is often called **baseline**
- It can **improve** convergence behavior
- It's something to be **engineered** in general
 - ↳ But no worries! We will see an obvious choice shortly ☺

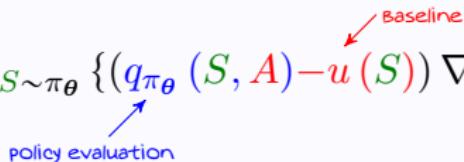
Policy Gradient Theorem with Baseline

Policy Gradient with Baseline

For a policy network with non-zero probabilities, the gradient of the average trajectory return is always given by

$$\nabla \mathcal{J}(\pi_{\theta}) = \mathbb{E}_{S \sim d_{\pi_{\theta}}, A | S \sim \pi_{\theta}} \{(q_{\pi_{\theta}}(S, A) - u(S)) \nabla \log \pi_{\theta}(A|S)\}$$

for any **baseline** $u(\cdot)$



Baseline PGM: General Form

BaselinePGM():

- 1: *Initiate with θ and learning rate α*
- 2: *Use a Q-estimator $\text{QEst}()$*
- 3: **while interacting do**
- 4: *Sample the environment with policy π_θ*

$$S_0, A_0 \xrightarrow{R_1} S_{t+1}, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: Set $Q_t = \text{QEst}(S_t, A_t)$
- 7: Compute **baseline estimator** $B_t = u(S_t)$
- 8: Update policy network $\theta \leftarrow \theta + \alpha (Q_t - B_t) \nabla \log \pi_\theta (A_t | S_t)$
- 9: **end for**
- 10: **end while**

Value Function: Obvious Choice of Baseline

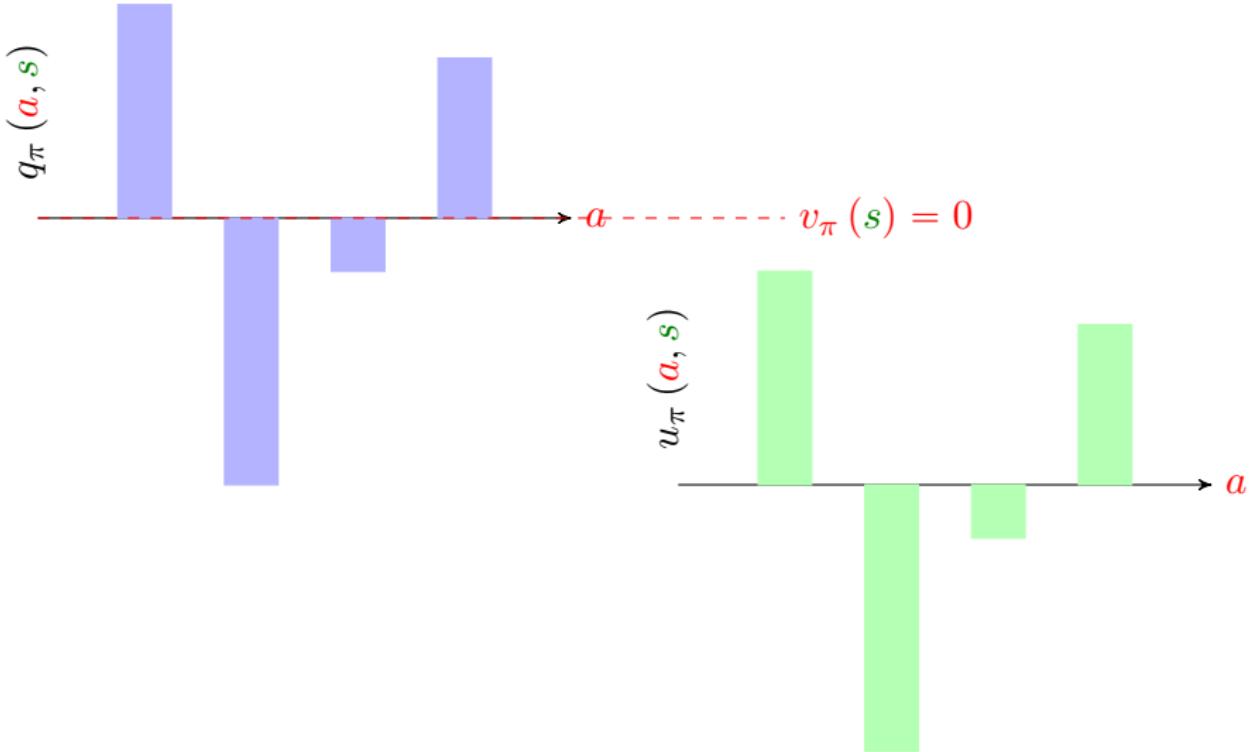
- + What is the obvious choice of baseline?!
- Value function $v_{\pi_\theta}(s)$!
- + How is it obvious?!
- In this case, shifted action-value represents the co-called advantage

Advantage

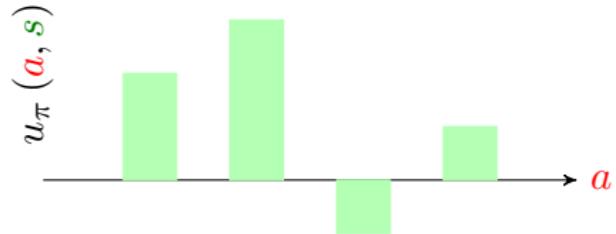
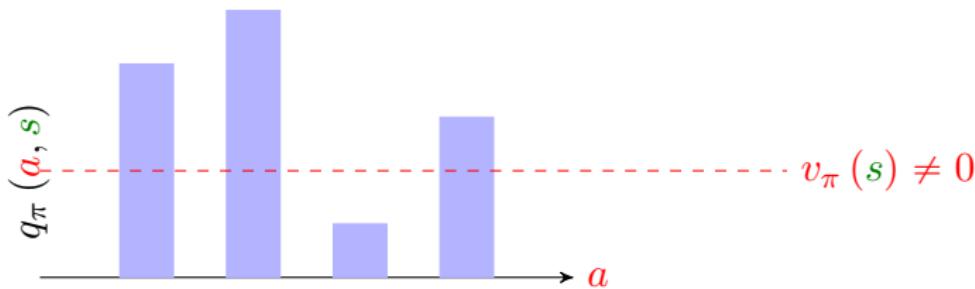
Given policy π , the advantage of action a at state s is defined as

$$u_\pi(a, s) = q_\pi(a, s) - v_\pi(s)$$

Advantage: Visualization



Advantage: Visualization



Advantage at any state concentrates around zero

Baseline PGM: Advantage Optimization

Policy Gradient with Advantage

For a policy network with non-zero probabilities, the gradient of the average trajectory return is also given by

$$\begin{aligned}\nabla \mathcal{J}(\pi_{\theta}) &= \mathbb{E}_{S \sim d_{\pi_{\theta}}, A|S \sim \pi_{\theta}} \{u_{\pi_{\theta}}(S, A) \nabla \log \pi_{\theta}(A|S)\} \\ &= \mathbb{E}_{S \sim d_{\pi_{\theta}}, A|S \sim \pi_{\theta}} \{(q_{\pi_{\theta}}(S, A) - v_{\pi_{\theta}}(S)) \nabla \log \pi_{\theta}(A|S)\}\end{aligned}$$

- + But how can we find an estimator for advantage?
- If we can estimate action-values, we can obviously use

$$v_{\pi_{\theta}}(s) = \mathbb{E}_{\pi_{\theta}} \{q_{\pi_{\theta}}(s, A)\} = \int_a q_{\pi_{\theta}}(s, a) \pi_{\theta}(a|s)$$

Baseline PGM: Advantage Optimization

AdvantagePGM() :

- 1: Initiate with θ and learning rate α
- 2: Use a **Q-estimator** QEst()
- 3: **while interacting do**
- 4: Sample the environment with policy π_θ

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 5: **for** $t = 0 : T - 1$ **do**
- 6: Set $Q_t = \text{QEst}(S_t, A_t)$
- 7: Compute **value** $V_t = \mathbb{E}_{\pi_\theta} \{Q_t | S_t\}$ and **sample advantage** $U_t = Q_t - V_t$
- 8: Update policy network $\theta \leftarrow \theta + \alpha U_t \nabla \log \pi_\theta (A_t | S_t)$
- 9: **end for**
- 10: **end while**

PGM with Temporal Difference Estimate

- + We have only considered Monte Carlo approach to estimate values! Why don't we use TD?!
- Well! If we only work with a policy network, it could be challenging

Say we have a particular sample trajectory that looks at time t as

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$$

If we use TD-0 to estimate $\hat{q}_{\pi_\theta}(S_t, A_t)$, we would write

$$\hat{q}_{\pi_\theta}(S_t, A_t) = R_{t+1} + \gamma \hat{v}_{\pi_\theta}(S_{t+1})$$

↑
where do we get this estimate?!

Estimating via TD in PGM: Main Challenge

In value-based RL, we gradually find an *estimate* for values

- In tabular RL, we make a *Q-table* and update it
- In *deep* RL, we train a *value network*, e.g., a DQN

In *pure PGM*, we have neither of them!

-
- + Then what can we do? We cannot always use *Monte Carlo*! What if the environment is not *episodic*?
 - Well there are three solutions with *only one of them working!*

Estimating via TD in PGM: Possible Solutions

- ① We may stick to **Monte Carlo** approach
 - ↳ It has high variance and is hard to use in **non-episodic environments**
- ② We may try to **evaluate** the policy in each iteration
 - ↳ This is in practice **computationally infeasible**
- ③ We may train a **value network** in addition to the **policy network**
 - ↳ This describes the class of **actor-critic** methods
 - ↳ An **actor** who plays with the **policy network** and update it via PGM
 - ↳ A **critic** who evaluates the **policy** with a **value network** and update it with DQL
 - ↳ We will get to these methods in the **next chapter**

For now, let's make an **idealistic** assumption: we assume that we can really evaluate a policy, i.e., given π_θ for any θ

we can compute $v_{\pi_\theta}(s)$ and $q_{\pi_\theta}(s, A)$

We will later get rid of this **idealistic** assumption by the help of **value networks**

Estimating via TD in PGM: Idealistic Case

With this assumption, we can rewrite our algorithms in an online form, e.g.,

AdvantagePGM():

- 1: Initiate with θ and learning rate α
- 2: **while interacting do**
- 3: Sample the environment with policy π_θ

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 4: **for** $t = 0 : T - 1$ **do**
- 5: Compute $q_{\pi_\theta}(S_t, A_t)$ and $v_{\pi_\theta}(S_t) = \mathbb{E}_{\pi_\theta}\{q_{\pi_\theta}(S_t, A_t) | S_t\}$
- 6: Compute sample advantage $U_t = q_{\pi_\theta}(S_t, A_t) - v_{\pi_\theta}(S_t)$
- 7: **end for**
- 8: Update policy network

$$\theta \leftarrow \theta + \alpha \sum_{t=0}^{T-1} U_t \nabla \log \pi_\theta(A_t | S_t)$$

- 9: **end while**

TD Error as Advantage Estimator

$$S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$$

Let's look at the *sample advantage*: we have

$$U_t = q_{\pi_\theta}(S_t, A_t) - v_{\pi_\theta}(S_t)$$

Using *Bellman's equation*, we can write

$$U_t = \mathbb{E}\{R_{t+1}\} + \gamma \mathbb{E}_{\pi_\theta}\{v_{\pi_\theta}(S_{t+1}) | S_t, A_t\} - v_{\pi_\theta}(S_t)$$

which we can be estimated by

$$\hat{U}_t = R_{t+1} + \gamma v_{\pi_\theta}(S_{t+1}) - v_{\pi_\theta}(S_t)$$

This is the **TD-0 error**: TD error is an estimator of advantage!

Advantage PGM: Online via TD Estimate

AdvantagePGM():

- 1: Initiate with θ and learning rate α
- 2: **while interacting do**
- 3: Sample the environment with policy π_θ

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 4: **for** $t = 0 : T - 1$ **do**
- 5: Compute **sample advantage** $U_t = R_{t+1} + \gamma v_{\pi_\theta}(S_{t+1}) - v_{\pi_\theta}(S_t)$
- 6: **end for**
- 7: Update policy network

$$\theta \leftarrow \theta + \alpha \sum_{t=0}^{T-1} U_t \nabla \log \pi_\theta(A_t | S_t)$$

- 8: **end while**

- + And, we do **not** need **action-values!**
- Right! **Value function** is **enough**

Advantage PGM: Online via TD Estimate

Obviously, we can find a more robust estimator via TD- n

AdvantagePGM(n) :

- 1: Initiate with θ and learning rate α
- 2: **while interacting do**
- 3: Sample the environment with policy π_θ

$$S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

- 4: **for** $t = 0 : T - n - 1$ **do**

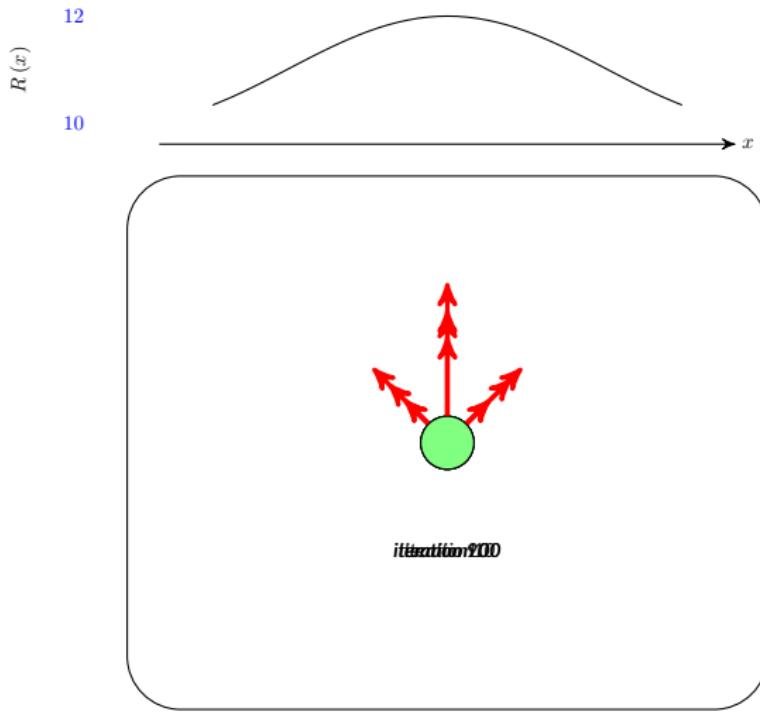
- 5: Compute $U_t = \sum_{i=0}^n \gamma^i R_{t+i+1} + \gamma^{n+1} v_{\pi_\theta}(S_{t+n+1}) - v_{\pi_\theta}(S_t)$

- 6: **end for**

$$\theta \leftarrow \theta + \alpha \sum_{t=0}^{T-n-1} U_t \nabla \log \pi_\theta(A_t | S_t)$$

- 7: **end while**

Example: Controlling Moving Particle - Case II



Crucial Challenge in PGM: Sample Inefficiency

After implementing AdvantagePGM(), we see: though using **baseline**, the **variance** reduces, it still needs **long time** to converge

the main reason is that PGM is **sample inefficient**

In all above algorithms

- We sample $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_T} S_T$ and use it for update
- We never get back to this sample

This is generally a big challenge in PGMs!

- + Can't we do what we did in DQL?!
- You mean **experience reply**??!
- + Right! Just **keep previous samples in a buffer** and reuse them again
- Well! The issue is that those samples were collected by **other policies**, i.e., π_θ for **old θ** . Through time, we have gone far away from them!

Solution: Trust Region PGM

- + So was it with DQL! How did we get rid of that?!
- We were playing off-policy, so we did not need to have sample with the target policy
- + So, isn't there any way to improve the sample efficiency?
- There is one and we do know it!

The solution to this challenge is to use the idea of importance sampling: recall that if $X \sim p(x)$ we could write

$$\mathbb{E}_q\{X\} = \int_x x q(x) = \int_x x \frac{q(x)}{p(x)} p(x) = \mathbb{E}_p\left\{X \frac{q(X)}{p(X)}\right\}$$

This leads to PGMs with trust region that we will learn next!