

Reinforcement Learning

Chapter 6: Actor Critic Methods


Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025


Deep RL: Sort of Division



Deep RL: Sort of Division

In *actor-critic approaches* we have both networks

- an *actor* has a *policy network*
 - ↳ This network enables it to *act* at each *particular state*
- a *critic* has a *value network*
 - ↳ This network enables it to *evaluate* its *policy*
 - ↳ The *evaluation* will help *improving* the *policy*



Deep RL: Sort of Division

Attention

For many people *actor-critic* \equiv PGM: they usually argue that

- to implement a PGM we need to estimate values
- we should do it by a value network

So, any PGM is at the end actor-critic

That's practically true; however, in principle, we can

implement PGMs via basic Monte Carlo

So, we could also have a pure PGM, e.g., REINFORCE!

Implementing PGMs

Let's get back to **PGMs**: say we want to implement a PGM

- We usually use **sample advantages**, i.e.,

$$U_t = R_{t+1} + \gamma v_{\pi_\theta}(S_{t+1}) - v_{\pi_\theta}(S_t)$$

So, we need to know the value function $v_{\pi_\theta}(\cdot)$ of our **policy** π_θ

- + Well, why don't we **evaluate** it once and use it forever?
- **Attention!** We need this **evaluation each time** we update **policy** π_θ !
- + How exactly we do it then? You promised to tell us!
- Sure! Let's use what we have learned up to now

Advantage PGM: Implementing

Let's look at the classic *advantage optimization* PGM

AdvantagePGM():

- 1: Initiate with θ and learning rate α
- 2: **while interacting do**
- 3: Set $\hat{\nabla} = \mathbf{0}$
- 4: **for mini-batch $b = 1 : B$ do**
- 5: Sample $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$ with policy π_θ
- 6: **for $t = 0 : T - 1$ do**
- 7: Compute *sample advantage* $U_t = R_{t+1} + \gamma v_{\pi_\theta}(S_{t+1}) - v_{\pi_\theta}(S_t)$
- 8: Compute sample gradient $\hat{\nabla} \leftarrow \hat{\nabla} + U_t \nabla \log \pi_\theta(A_t | S_t) / B$
- 9: **end for**
- 10: **end for**
- 11: Update policy network $\theta \leftarrow \theta + \alpha \hat{\nabla}$
- 12: **end while**

To implement, we need to estimate $v_{\pi_\theta}(S_t)$ for all trajectories in **mini-batch**

Estimating Values: Monte-Carlo

Say we are looking into one trajectory τ

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

We know how to use this trajectory to compute value estimates: for each t

$$\hat{V}_t = \text{estimate of value for } S_t = G_t = \sum_{i=t}^T \gamma^i R_{i+1}$$

If we the same state happens multiple times in the trajectory: we count the number of times $S_t = S$ appears in the trajectory and average estimates, i.e.,

$$\hat{v}_{\pi_\theta}(S) = \frac{1}{\mathcal{N}(S \in \tau)} \sum_{t=0}^{T-1} \mathbf{1}\{S_t = S\} \hat{V}_t$$

where $\mathcal{N}(S \in \tau)$ is the number of times S has appeared in τ

Estimating Values: Monte-Carlo

If we have a *mini-batch* \mathbb{B} of trajectories

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

We use the same approach

$$\hat{V}_t [\tau] = G_t [\tau] = \sum_{i=t}^T \gamma^i R_{i+1} [\tau]$$

We count the number of times $S_t = S$ appears in all trajectories and average the sample estimates, i.e.,

$$\hat{v}_{\pi_\theta}(S) = \frac{1}{\mathcal{N}(S \in \mathbb{B})} \sum_{\tau \in \mathbb{B}} \sum_{t=0}^{T-1} \mathbf{1}\{S_t[\tau] = S\} \hat{V}_t[\tau]$$

where $\mathcal{N}(S \in \mathbb{B})$ the number of times S has appeared in \mathbb{B}

Advantage PGM: With Value Estimates

EstAdvantagePGM():

- 1: Initiate with θ and learning rate α
- 2: **while interacting do**
- 3: **for mini-batch $b = 1 : B$ do**
- 4: Sample $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$ with policy π_θ
- 5: **end for**
- 6: Estimate value of all observed states in the **mini-batch** as $\hat{v}_{\pi_\theta}(S_t)$
- 7: Set $\hat{\nabla} = 0$
- 8: **for $b = 1 : B$ do**
- 9: **for $t = 0 : T - 1$ do**
- 10: Compute **sample advantages** $U_t = R_{t+1} + \gamma \hat{v}_{\pi_\theta}(S_{t+1}) - \hat{v}_{\pi_\theta}(S_t)$
- 11: Update sample gradient $\hat{\nabla} \leftarrow \hat{\nabla} + U_t \nabla \log \pi_\theta(A_t | S_t) / B$
- 12: **end for**
- 13: **end for**
- 14: Update policy network $\theta \leftarrow \theta + \alpha \hat{\nabla}$
- 15: **end while**

Advantage PGM: With Value Estimates

We could guess that this algorithm is **not** going to perform very **impressive!**


- + And why is that?!
- For the exact same reasons we said at the beginning of Chapter 4
 - ↳ We have **lots of states**
 - ↳ Many of them are **rarely observed** in a **small mini-batch**
 - ↳ The estimates can hence be very **high variance**
 - ↳ Also we need to wait for the **whole** mini-batch to be ready
 - ↳ ...
- + So, what is the solution?
- You tell me!
- + We go for function approximation via **value networks!**
- You got it right!

Recall: Value Network

Let's keep our trajectories here

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

What we need is a simple v-network, as we only need the state values



In Chapter 4, we saw that we could train it via sample returns, i.e.,

$$\text{Dataset} = \left\{ (S_t[\tau], \hat{V}_t[\tau]) : \forall t \text{ and } \tau \right\}$$

and we train the network by minimizing the least-square loss

Value Network: Training

Let's keep our trajectories here

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

This means that we compute the loss function

$$\mathcal{L}^v(\mathbf{w}) = \sum_{\tau} \sum_t \left(v_{\mathbf{w}}(S_t[\tau]) - \hat{V}_t[\tau] \right)^2$$

*and update the weights of the **v-network** as*

$$\mathbf{w} \leftarrow \operatorname{argmin}_{\mathbf{w}} \mathcal{L}^v(\mathbf{w})$$

which we approximately solve using gradient descent

Basic Actor-Critic

This is going to end us with a basic *actor-critic algorithm*:

AC_v1():

- 1: Initiate with θ and w , as well as a learning rate α
- 2: **while interacting do**
- 3: **for** mini-batch $b = 1 : B$ **do**
- 4: Sample $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$ with policy π_θ
- 5: **for** $t = 0 : T - 1$ **do**
- 6: Compute **value estimate** \hat{V}_t
- 7: Compute **sample advantages** $U_t = R_{t+1} + \gamma v_w(S_{t+1}) - v_w(S_t)$
- 8: Update sample gradient $\hat{\nabla} \leftarrow \hat{\nabla} + U_t \nabla \log \pi_\theta(A_t | S_t) / B$
- 9: **end for**
- 10: **end for**
- 11: Update policy network $\theta \leftarrow \theta + \alpha \hat{\nabla}$
- 12: Update w by SGD using **value estimates** \hat{V}_t
- 13: **end while**

Training Value Network: TD Estimates

$$\tau : S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$$

But now that we have a **value network**, we could also use TD: at step t , we set

$$\hat{V}_t = \text{estimate of value for } S_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1})$$

- We can estimate the **advantage** using the current **value network**

$$U_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_{t+1})$$

- We then use least-squares update **value network** by TD sample estimates

$$\mathcal{L}^v(\mathbf{w}) = \sum_{\tau} \sum_{t=0}^{T-1} \left(v_{\mathbf{w}}(S_t[\tau]) - \hat{V}_t[\tau] \right)^2$$

Training Value Network: TD Estimates

Let's write the update rule of the value network: we compute the gradient of loss and move in that direction

$$\nabla \mathcal{L}^v(\mathbf{w}) = 2 \sum_{t=0}^{T-1} \underbrace{\left(v_{\mathbf{w}}(S_t) - \hat{V}_t \right)}_{-\Delta_t} \nabla v_{\mathbf{w}}(S_t)$$

We used to call Δ_t the TD error, and set the learning rate to some $\beta/2$ to get

$$\mathbf{w} \leftarrow \mathbf{w} + \beta \sum_{t=0}^{T-1} \Delta_t \nabla v_{\mathbf{w}}(S_t)$$

Training Value Network: TD Estimates

Let's look at **TD error**: recall that our **labels**, i.e., sample estimates of values, are

$$\hat{V}_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1})$$

So the **TD error** is given by

$$\begin{aligned}\Delta_t &= \hat{V}_t - v_{\mathbf{w}}(S_t) \\ &= R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t) \\ &= U_t\end{aligned}$$

This leads us to what we **observed** in Chapter 5

Recall: Advantage vs TD Error

Advantage is an estimator of TD error

A2C: Basic Version

A2C():

- 1: Initiate with θ and w , as well as learning rates α and β
- 2: **while interacting do**
- 3: Start with zero gradients $\hat{\nabla}_w = \hat{\nabla}_\theta = \mathbf{0}$
- 4: Sample $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$ with policy π_θ
- 5: Compute sample advantages $U_t = R_{t+1} + \gamma v_w(S_{t+1}) - v_w(S_t)$
- 6: **for** $t = 0 : T - 1$ **do**
- 7: Compute sample policy gradient $\hat{\nabla}_\theta \leftarrow \hat{\nabla}_\theta + U_t \nabla \log \pi_\theta(A_t | S_t)$
- 8: Compute sample value gradient $\hat{\nabla}_w \leftarrow \hat{\nabla}_w + U_t \nabla v_w(S_t)$
- 9: **end for**
- 10: Update policy network $\theta \leftarrow \theta + \alpha \hat{\nabla}_\theta$
- 11: Update value network $w \leftarrow w + \beta \hat{\nabla}_w$
- 12: **end while**

This is the **single-trajectory** form of

Advantage Actor Critic \equiv A2C

A2C: Online Version

Since we use TD, we can also update *online*, i.e., in each time step

OnlineA2C():

- 1: Initiate with θ and w , a random state S_0 , $t = 0$ and learning rates α and β
- 2: **while interacting do**
- 3: Sample A_t from $\pi_\theta(\cdot|S_t)$
- 4: Sample single step $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$ from environment
- 5: Compute sample advantage $U_t = R_{t+1} + \gamma v_w(S_{t+1}) - v_w(S_t)$
- 6: Update policy network $\theta \leftarrow \theta + \alpha U_t \nabla \log \pi_\theta(A_t|S_t)$
- 7: Update value network $w \leftarrow w + \beta U_t \nabla v_w(S_t)$
- 8: **if** S_{t+1} is terminal **then** draw a random S_{t+1}
- 9: Set $t \leftarrow t + 1$
- 10: **end while**

But, that would be too noisy and hence quite *unstable*

A2C: Mini-Batch Version

We can further extend to **mini-batch** learning

`miniBatchA2C():`


- 1: Initiate with θ and w , as well as learning rates α and β
- 2: **while interacting do**
- 3: Start with zero gradients $\hat{\nabla}_w = \hat{\nabla}_\theta = \mathbf{0}$
- 4: **for mini-batch $b = 1 : B$ do**
- 5: Sample $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$ with policy π_θ
- 6: Compute sample advantages $U_t = R_{t+1} + \gamma v_w(S_{t+1}) - v_w(S_t)$
- 7: **for $t = 0 : T - 1$ do**
- 8: Compute sample policy gradient $\hat{\nabla}_\theta \leftarrow \hat{\nabla}_\theta + U_t \nabla \log \pi_\theta(A_t | S_t)$
- 9: Compute sample value gradient $\hat{\nabla}_w \leftarrow \hat{\nabla}_w + U_t \nabla v_w(S_t)$
- 10: **end for**
- 11: **end for**
- 12: Update policy network $\theta \leftarrow \theta + \alpha \hat{\nabla}_\theta$
- 13: Update value network $w \leftarrow w + \beta \hat{\nabla}_w$
- 14: **end while**

Actor-Critic via Shared-Network

There is one extra **obvious** fact: *the policy and values that we learn are very much mutually related!*

- + So, why don't we learn them **together**?!
- Actually we can!


We can consider an **actor-critic** model, i.e.,



and train it all together!

- This model can be simply a DNN
- The DNN's output contains both **policy** and **value**

Actor-Critic via Shared-Network: Visualization



Here, **value** and **policy** share same layers except the few **last layers**

Actor-Critic via Shared-Network: Loss

- + But how can we train the loss in this network?
- We could let it to be proportional to sum of our both objectives

We could define a new loss as

$$\begin{aligned}\mathcal{L}(\mathbf{w}) &= -\mathcal{J}(\pi_{\mathbf{w}}) + \xi \mathcal{L}^v(\mathbf{w}) \\ &= \sum_{\tau} \sum_{t=0}^{T-1} -U_t[\tau] \log \pi_{\mathbf{w}}(A_t[\tau] | S_t[\tau]) + \xi \left(v_{\mathbf{w}}(S_t[\tau]) - \hat{V}_t[\tau] \right)^2\end{aligned}$$

for some hyperparameter ξ : it's easy to see that in this case

$$\nabla \mathcal{L}(\mathbf{w}) = - \sum_{\tau} \sum_{t=0}^{T-1} U_t[\tau] [\nabla \log \pi_{\mathbf{w}}(A_t[\tau] | S_t[\tau]) + \xi \nabla v_{\mathbf{w}}(S_t[\tau])]$$

A2C: Shared-Network Version

sharedNetA2C() :

- 1: *Initiate shared network (π_w, v_w) with w*
- 2: *Choose potentially scheduled value-weight ξ and learning rate α*
- 3: **while interacting do**
- 4: *Start with zero gradients $\hat{\nabla}_w = \mathbf{0}$*
- 5: **for mini-batch $b = 1 : B$ do**
- 6: *Sample $S_0, A_0 \xrightarrow{R_1} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$ with policy π_w*
- 7: *Compute sample advantages $U_t = R_{t+1} + \gamma v_w(S_{t+1}) - v_w(S_t)$*
- 8: **for $t = 0 : T - 1$ do**
- 9: *Compute sample gradient $\hat{\nabla} \leftarrow \hat{\nabla} + U_t [\nabla \log \pi_w(A_t | S_t) + \xi \nabla v_w(S_t)]$*
- 10: **end for**
- 11: **end for**
- 12: *Update shared network $w \leftarrow w + \alpha \hat{\nabla}$*
- 13: **end while**