

Reinforcement Learning

Chapter 6: Actor Critic Methods

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

Fall 2025

Learning Deterministic Policy

From model-based RL we know that: the *optimal policy can be deterministic*


Why don't we train a *policy network* that learns a *deterministic policy*?

-
- + You are contradicting yourself! You said that *stochastic policy* is a general case that includes *deterministic policies* as well! Now you want to get back to a *deterministic policy*!?
 - Well! You're *right!* But there will be no harm in learning a *deterministic policy*! It might only be *less effective!*
 - + Why we should do it then?
 - It could give us some *benefits*, especially when we have *continuous action-space*

Learning Deterministic Policy

With *continuous action-space*, policy is a density *function*

- With *discrete action-space*, we can show policy by a finite vector¹



- Unlike *discrete action-space*, we *cannot* do this with *continuous actions*
 - We should learn a *function* from *state feature*, e.g.,

$$\pi_\theta(a|s) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(a - \text{DNN}(x(s)|\theta))}{2\sigma^2} \right\}$$

- We then sample from this learned *density function*, e.g.,
 - Draw a sample from Gaussian distribution with *mean* $\text{DNN}(x(s)|\theta)$ and *variance* σ^2

¹We have seen this in Assignment 3

Learning Deterministic Policy

There are several things that could go wrong


- *What if the generated sample is out of accepted range?*
 - ↳ Our sample from Gaussian distribution is **extremely large**
 - ↳ In **sensitive control settings**, this could harm the system
- *What if we only try a few samples?*
 - ↳ We don't see the probabilities as opposed to **continuous actions**
 - ↳ We then **cannot** really **reject** too many samples

*With **continuous actions**, we usually prefer to learn a **deterministic policy**: its main feature is that it can be represented by a **single action a^****

$$\pi_{\theta}(a|s) = \begin{cases} 1 & a = a^* \\ 0 & a \neq a^* \end{cases}$$

Deterministic Policy Network

Considering a **deterministic** policy, we **only** need to learn an estimate of **optimal action** in each state: we can revise our policy network into a **deterministic policy network**



Deterministic Policy Network

Deterministic policy network maps a **state**-features into a **single action** and can be realized by a DNN with input being the state feature representation and a **single output**


Deterministic Policy Network: Training

- + How should we train *these networks*? Similar to other *policy networks*?
- We can look at them as a special form of *policy networks* and do the same thing, yes! However, it turns out that in this special case, *there are better ways to do it!* Especially as we always implement them *actor-critic*
- + So, we should start all over *again*!?
- Not really! We should basically use what we learned for *DQL*

Deterministic Policy Network: Training

Recall the property of *optimal policy*: it gives *maximum value* and *action-value*

If we have a Q-network that estimates *optimal action-value*, we can say



The output satisfies

$$\hat{q}_*(s, \hat{a}^*) = \max_a \hat{q}_*(s, a)$$

Deterministic Policy Network: Training

So, in *actor-critic* form with a *Q*-network, we could train the deterministic policy network as

$$\theta^* = \underset{\theta}{\operatorname{argmax}} Q_w(s, \mu_{\theta}(s))$$

which we can solve using *gradient ascent* by updating as

$$\begin{aligned} \theta &\leftarrow \theta + \alpha \nabla_{\theta} Q_w(s, \mu_{\theta}(s)) \\ &\leftarrow \theta + \alpha \underbrace{\frac{\partial}{\partial a} Q_w(s, a) |_{a=\mu_{\theta}(s)}}_{\text{Backprop over } Q\text{-Net}} \underbrace{\nabla \mu_{\theta}(s)}_{\text{Backprop over policy}} \end{aligned}$$

Moral of Story

As long as we have an estimator of optimal *action-value function*, we can train deterministic policy network very easily!

Deterministic Policy Gradient

- + But how can we find such an estimator?
- Well! We have done this **before!**
- + You mean in **DQL**!?
- Exactly! In **Q-learning** we use **Bellman's optimality equation** to estimate **optimal action-value** function: we can do the same here

Recall that **Bellman's optimality equation** indicate that

$$q_{\star}(S_t, A_t) = R_{t+1} + \gamma \mathbb{E}_{S_{t+1} \sim p(\cdot | S_t, A_t)} \left\{ \max_a q_{\star}(S_{t+1}, a) \right\}$$

and if we know the action $a^{\star} = \operatorname{argmax}_a q_{\star}(S_{t+1}, a)$, we could write

$$q_{\star}(S_t, A_t) = R_{t+1} + \gamma \mathbb{E}_{S_{t+1} \sim p(\cdot | S_t, A_t)} \{ q_{\star}(S_{t+1}, a^{\star}) \}$$

Deterministic Policy Gradient

If we use our **deterministic** policy network in one time step we sample

$$S_t, \mu_{\theta}(S_{t+1}) \xrightarrow{R_{t+1}} S_{t+1}$$

We can then sample an estimator of **optimal action-value** at $a = \mu_{\theta}(S_{t+1})$

$$\hat{Q}_t = R_{t+1} + \gamma Q_w(S_{t+1}, \mu_{\theta}(S_{t+1}))$$

Once we are over with **sample trajectory**: we update **Q-network** to minimize loss

$$\mathcal{L}(w) = \frac{1}{T} \sum_{t=0}^{T-1} \left(Q_w(S_t, \mu_{\theta}(S_t)) - \hat{Q}_t \right)^2$$

And the life is **much easier** as compared to **TRPO and PPO** 😊

Deterministic Policy Gradient

We do very well know how to do this

$$\mathbf{w} \leftarrow \mathbf{w} - \beta \nabla \mathcal{L}(\mathbf{w})$$

$$\leftarrow \mathbf{w} + \beta \text{mean} \left[\left(\hat{Q}_t - Q_{\mathbf{w}}(S_t, \mu_{\theta}(S_t)) \right) \nabla Q_{\mathbf{w}}(S_t, \mu_{\theta}(S_t)) \right]$$

$$\leftarrow \mathbf{w} + \beta \text{mean} [\Delta_t \nabla Q_{\mathbf{w}}(S_t, \mu_{\theta}(S_t))]$$

where $\Delta_t = \hat{Q}_t - Q_{\mathbf{w}}(S_t, \mu_{\theta}(S_t))$ is the TD error

Alternating between the two update rules, we end up with a

Deterministic Policy Gradient \equiv DPG

algorithm: there are various DPG algorithms; we take a look into the famous one

A Basic DPG Algorithm

We can use these updates to write a simple *online* DPG algorithm

DPG_v1():

- 1: Initiate with θ and w , as well as factor $\alpha < 1$ and learning rate β
- 2: Initiate some *initial state* S_0 and draw *action* A_0 as $A_0 \leftarrow \mu_\theta(S_0)$
- 3: **while interacting do**
- 4: Sample a *time step* $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$
- 5: Draw the next optimal action as $A_{t+1} \leftarrow \mu_\theta(S_{t+1})$
- 6: Compute $\Delta = R_{t+1} + \gamma Q_w(S_{t+1}, A_{t+1}) - Q_w(S_t, A_t)$
- 7: Update *value network* as $w \leftarrow w + \beta \Delta \nabla Q_w(S_t, A_t)$
- 8: Update *policy network* as $\theta \leftarrow \theta + \alpha \frac{\partial}{\partial a} Q_w(S_t, a) |_{a=A_t} \nabla \mu_\theta(S_t)$
- 9: Go for next state $S_t \leftarrow S_{t+1}$
- 10: **if** S_t is terminal **then**
- 11: Draw a new *random state* S_0 and $A_0 \leftarrow \mu_\theta(S_0)$
- 12: **end if**
- 13: **end while**

Basic DPG: Practical Challenges

Using our knowledge, we can easily detect **challenges** of this basic algorithm

- Lack of **exploration** → ϵ -greedy improvement
 - ↳ We follow **blindly** the **deterministic policy network**
 - ↳ We do not give any chance for **exploration**
 - ↳ This can quickly stick us to a **bad locally-optimal** deterministic policy
- **High-variance** gradient estimators → **experience replay**
 - ↳ It's **online** and hence update the networks with **single** time step **samples**
 - ↳ We need more samples to compute **better estimators**
 - ↳ We would like to have **independent samples**
- Variation of **training labels** → **target network**
 - ↳ Each time we update, we **change** the label in the **training batch**
 - ↳ This can severely deteriorate the **convergence of algorithm**

DPG: ϵ -Greedy Improvement

To have sufficient **exploration** of environment: we can follow ϵ -greedy approach

- + But how does it work here? You said we have **continuous actions**!
- Well! We can add **continuous randomness** to our policy

Say we get $A_t \leftarrow \mu_\theta(S_t)$ at time t : then we replace our action with

$$A_t \leftarrow A_t + \sqrt{\epsilon} Z_t$$

where Z_t is **random noise** with **mean zero** and **variance one**

Classical choice of Z_t is **zero-mean unit-variance Gaussian variable**, i.e.,

$$Z_t \sim \mathcal{N}(0, 1)$$

Note that the noise term $\sqrt{\epsilon} Z_t$ is then **zero-mean** with **variance ϵ**

DPG: ϵ -Greedy Improvement

- + But what if after adding $\sqrt{\epsilon}Z_t$, the **action** gets out of its **allowed range**?
For instance, we get $A_t = 5$ and $\sqrt{\epsilon}Z_t = 3$, but we should have **all actions between 2 and 6**
- That's a valid question! We usually **clip** the action in this case

To avoid **out-of-range actions**, we replace apply ϵ -greedy approach as

$$A_t \leftarrow \text{Clip}(\mu_\theta(S_t) + \sqrt{\epsilon}Z_t, a_{\min}, a_{\max})$$

where a_{\min} and a_{\max} are minimum and maximum **allowed actions** and

$$\text{Clip}(x, a_{\min}, a_{\max}) = \begin{cases} a_{\min} & x < a_{\min} \\ x & a_{\min} \leq x \leq a_{\max} \\ a_{\max} & x > a_{\max} \end{cases}$$

DPG: Replay Buffer


To reduce estimator's variance and enhance sample-efficiency, we can use experience replay as in DQL

- + Wait a moment! But we talked about the fact that *experience replay* can increase *estimator's variance* in PGMs due to *importance sampling* argument! Now, we just ignore all those discussions?!
- With *stochastic* policy yes! But here we have a *deterministic policy*

Deterministic policy returns *only one action*

- For each choice of θ , our policy chooses *only one action*
 - ↳ If we change θ we only change this action
- Policy update *does not change* the probability of *all actions*
 - ↳ This is in fact why Q-learning does not suffer from *high estimate variance*


Recall: Experience Replay in DQL



We update DQN by randomly sampled mini-batches using **TD error**

$$\Delta_b \leftarrow R_{b+1} + \gamma \max_m Q_w(S_{b+1}, a^m) - Q_w(S_b, A_b)$$

DPG: Experience Replay



We now use the *deterministic* policy network to compute *TD error*

$$\Delta_b \leftarrow R_{b+1} + \gamma Q_w(S_{b+1}, \mu_\theta(S_{b+1})) - Q_w(S_b, A_b)$$

DPG: Target Network

The last thing to handle is to keep training dataset fixed for a while

- After each **mini-batch**, we change both **policy** and **Q-network**
 - ↳ We update w and θ
- If we use the same networks to compute the estimate

$$\hat{Q}_b = R_{b+1} + \gamma Q_w(S_{b+1}, \mu_\theta(S_{b+1}))$$

then our next iteration runs over a **different** dataset

- ↳ This can cause our training loop to **diverge**


We have dealt with this in **DQL** using **target network**

Target Network in DPG


Copy **DQN** and **policy** network into the exactly same target networks

- Use these target networks to compute the **estimates**
- Update them every multiple iterations by new copies of **online** networks

Recall: Target Network in DQL



DPG: Target Networks



DDPG: Deep DPG Algorithm

DDPG() :

- 1: Initiate with $\theta = \bar{\theta}$ and $w = \bar{w}$, as well as factor $\alpha < 1$ and learning rate β
- 2: Initiate state S_0 and draw $A_0 \leftarrow \text{Clip}(\mu_\theta(S_0) + \sqrt{\epsilon}\mathcal{N}(0, 1))$
- 3: **while interacting do**
- 4: Sample a time step $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$ and save in replay buffer
- 5: **for** multiple iterations **do**
- 6: Sample $S_b, A_b \xrightarrow{R_{b+1}} S_{b+1}$ from replay buffer
- 7: Draw $A_{b+1} \leftarrow \text{Clip}(\mu_{\bar{\theta}}(S_{b+1}) + \sqrt{\epsilon}\mathcal{N}(0, 1))$
- 8: Compute $\Delta = R_{b+1} + \gamma Q_{\bar{w}}(S_{b+1}, A_{b+1}) - Q_w(S_b, A_b)$
- 9: Update value network as $w \leftarrow w + \beta \Delta \nabla Q_w(S_t, A_t)$
- 10: Update policy network as $\theta \leftarrow \theta + \alpha \frac{\partial}{\partial a} Q_w(S_t, a) |_{a=A_t} \nabla \mu_\theta(S_t)$
- 11: **if** H iterations passed **then**
- 12: Copy $\bar{\theta} \leftarrow \theta$ and $\bar{w} \leftarrow w$
- 13: **end if**
- 14: **end for**
- 15: **end while**

DDPG: Overestimate Issue

DDPG has been the key DPG algorithm and widely used for

- Dealing with continuous action spaces
- Enabling *off-policy* learning with a deterministic version of PGM

DDPG has shown a key issue; namely, **overestimate** of **values**

Value Overestimate of DDPG

Action-values estimated by DDPG can have significant biases

- + We had it also in DQL and you said "it's not a big deal in general"! So, do we care about it here?!
- Well! Here is more important! Because, we use those estimates to update the policy and large biases would explode TD error!
- + So, shall we do double DQL then?!
- Pretty much yes!

TD3: Twin Delayed DDPG

Value overestimate has been addressed in the extended version of DDPG

Twin Delayed DDPG \equiv TD3

In TD3, we add three extra tricks to DDPG

- ① We use double DQN to suppress undesired bias
 - ↳ Remember that **bias** was mainly coming out of max operator
 - ↳ We came with a remedy called **double Q-learning**
- ② We delay the policy update, i.e., update the policy less frequent
 - ↳ We update **DQN** after each mini-batch, but
 - ↳ We update **policy network** once every couple of mini-batches
- ③ We add extra noise to actions when we use them in target networks
 - ↳ This way we make the **estimation error** somehow independent
 - ↳ This leads to **less bias**

From DPG to PGM

Both ideas of learning **deterministic** or **stochastic** policy have pros and cons

- For **deterministic** policy we could say
 - ↳ We can efficiently use **off-policy** learning
 - ↳ We can get better sample efficiency
 - ↳ **But** we get less chance of being optimal
 - ↳ We do not search among possible random optimal policies
- For **stochastic** policy we could say
 - ↳ We search among much larger set of policies
 - ↳ We can converge to a better policy
 - ↳ **But** we have troubles with sample efficiency
 - ↳ We cannot easily learn **off-policy** due to limits of **importance sampling**
- + Is there any way to get good things of both worlds?
- Soft actor-critic approaches actually do this

Recall: Information Content and Entropy

To understand the idea behind soft actor-critic, let's recap some definitions

Information Content

The information content of random variable $X \sim p(x)$ is

$$i(X) = \log \frac{1}{p(X)}$$

The information contents have some interesting properties

- It's always **non-negative**, since $0 \leq p(x) \leq 1$
- The **less likely** outcome $X = x$ is, the **more** will be its **information content**
 - ↳ Think about it! You will find it very intuitive

Recall: Information Content and Entropy

Entropy

For random variable $X \sim p(x)$, entropy is its average **information content**, i.e.,

$$H_p(X) = \mathbb{E}_p\{i(X)\} = \mathbb{E}_p\left\{\log \frac{1}{p(X)}\right\} = \int_x p(x) \log \frac{1}{p(x)}$$

Entropy quantifies how much **confusion** we have about X

- If X is **highly random**, e.g., uniformly or Gaussian distributed,
 - ↳ Then $H_p(X)$ is very large
- If X is **deterministic**
 - ↳ Then $H_p(X) = 0$

Redefining Value Function

After dealing with both **deterministic** and **stochastic** policies we might formulate the best policy as follows

- It's globally **deterministic**
 - ↳ If one action gives better reward, it should go for it
- It's locally **stochastic**
 - ↳ Among actions with same rewards, it chooses one at random

We could capture both these behaviors by looking into a new metric

Say we play with policy π : at time t , we are interested in

$$\tilde{R}_{t+1} = R_{t+1} + \xi H_\pi(A_t | S_t)$$

for some ξ , where $H_\pi(A_t | S_t)$ is entropy of **action** $A_t \sim \pi(\cdot | S_t)$

Redefining Value Function

Say we play with policy π : at time t , we are interested in

$$\tilde{R}_{t+1} = R_{t+1} + \xi H_\pi(A_t | S_t)$$

for some ξ , where $H_\pi(A_t | S_t)$ is entropy of action $A_t \sim \pi(\cdot | S_t)$

This new **modified reward** incorporates both desires

- Being globally **deterministic**
 - ↳ For actions with larger R_{t+1} , the modified \tilde{R}_{t+1} is also larger
- Being locally **stochastic**
 - ↳ For actions with same R_{t+1} , policy with higher randomness has larger \tilde{R}_{t+1}

Well! This might be a **better reward!**

SAC: Soft Actor-Critic

We can use either DPG or PGM to develop an actor-critic method for this new reward, i.e., we could define

$$\begin{aligned} v_{\pi}(s) &= \mathbb{E}_{\pi} \left\{ \sum_{i=0}^{\infty} \gamma^i \tilde{R}_{t+i+1} | S_t = s \right\} \\ &= \mathbb{E}_{\pi} \left\{ \sum_{i=0}^{\infty} \gamma^i [R_{t+i+1} + \xi H_{\pi}(A_{t+i} | S_{t+i})] | S_t = s \right\} \end{aligned}$$

Interestingly, we end up in both cases with the *same* policy and value gradients!

The derived actor critic method is referred to as

Soft Actor-Critic \equiv SAC

DRL Algorithms

Most DRL algorithms used in practice are *actor-critic*

We already discussed all main classes of *actor-critic approaches*

To each class, there are various extensions

- You are able now to *follow all those extensions*
- If necessary, you could *come up with your own particular extension!*

A rule-of-thumb is

- If you deal with *discrete actions* and have no concern on *sample efficiency*
 - ↳ Use *stochastic-policy actor-critic* approaches
- If you deal with *continuous actions* and/or need *sample efficiency*
 - ↳ Use *DPG-like actor-critic* approaches

OpenAI: Spinning Up in DRL

Congratulations! You are now *Deep RL experts!*



Looking for some mini-projects for further practices? Take a look at [OpenAI page](#)