

Reinforcement Learning

Chapter 7: Applications and Advancements in DRL

Ali Bereyhi

ali.bereyhi@utoronto.ca

Department of Electrical and Computer Engineering
University of Toronto

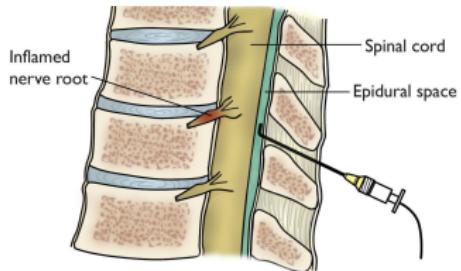
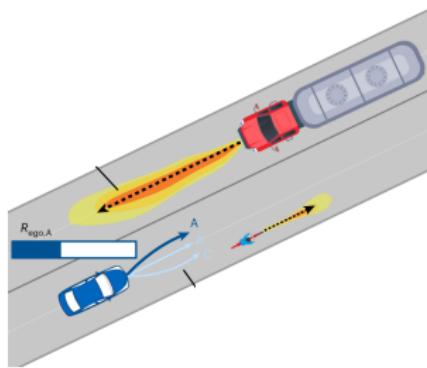
Fall 2025

Challenge of Sparse Rewarding

In many RL problems, we deal with **sparse** rewards

- In Frozen Lake game, we only get the reward at the end!

Can this be tolerated in many practical problems?



Waiting for those sparse **rewards** could be **very dangerous!**

A Solution: Reward Shaping

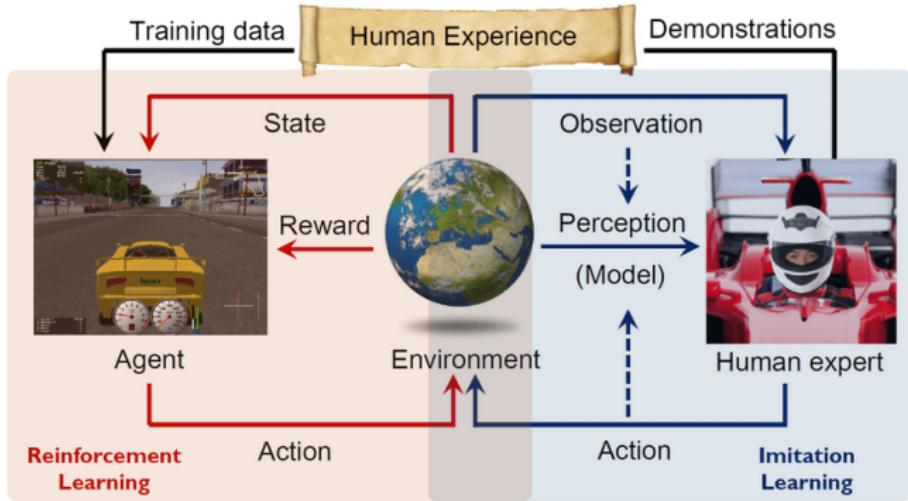
If we could formulate what we want: maybe, we could *shape the reward*

- In Assignment 1, we did it for simple Frozen Lake game
 - ↳ Just need to add some *negative rewards* in between
- By this approach, we make *much denser reward*
- + Can we always do it?
- Not really!

In practice, we cannot necessarily formulate the types of rewards we want

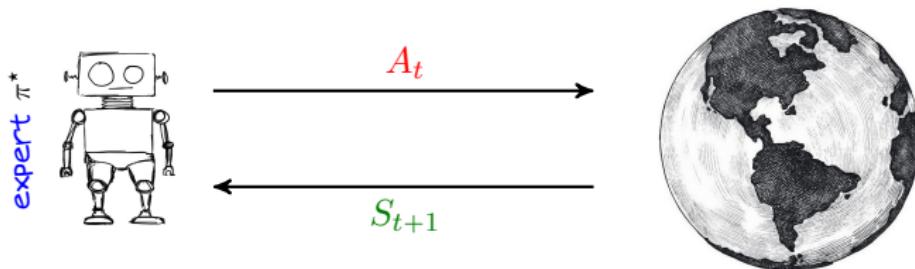
- We often have access only to a *human expert*
 - ↳ a *professional driver* or an *expert surgeon*

Imitation Learning: Learning by Expert



*In practice, we want to employ an **expert** to learn **safely** and **efficiently***

Behavior Cloning



We are looking for mimicking the *expert*

$$\min_{\theta} \mathcal{L}(\pi_{\theta}, \pi^*)$$

and of course, we do it by *sampling*!

Behavior Cloning: Not Always Efficient

Algorithm 1: DAgger: Dataset Aggregation

Data: π^*

Result: $\hat{\pi}^*$

$\mathcal{D} \leftarrow 0$

Initialize $\hat{\pi}$

for $i = 1$ **to** N **do**

$$\pi_i = \beta_i \pi^* + (1 - \beta_i) \hat{\pi}$$

Rollout policy π_i to sample trajectory $\tau = \{x_0, x_1, \dots\}$

Query expert to generate dataset $\mathcal{D}_i = \{(x_0, \pi^*(x_0)), (x_1, \pi^*(x_1)), \dots\}$

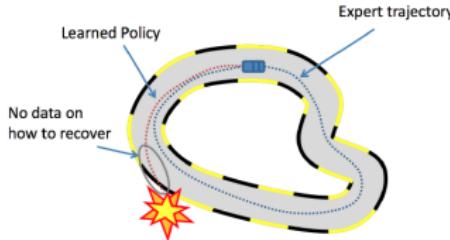
Aggregate datasets, $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_i$

Retrain policy $\hat{\pi}$ using aggregated dataset \mathcal{D}

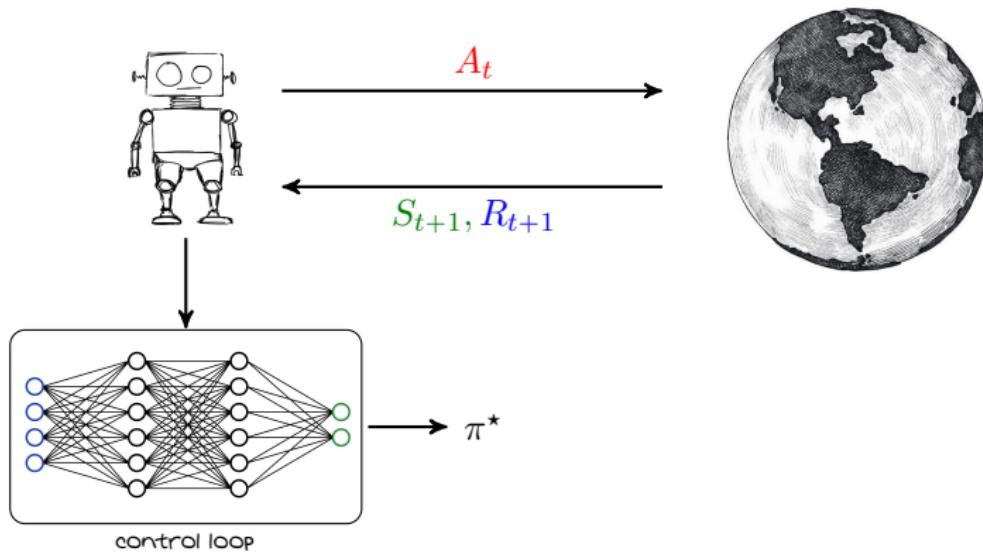
return $\hat{\pi}$

DAgger: famous efficient
algorithm for behavior cloning

But, we are always limited in collecting expert samples



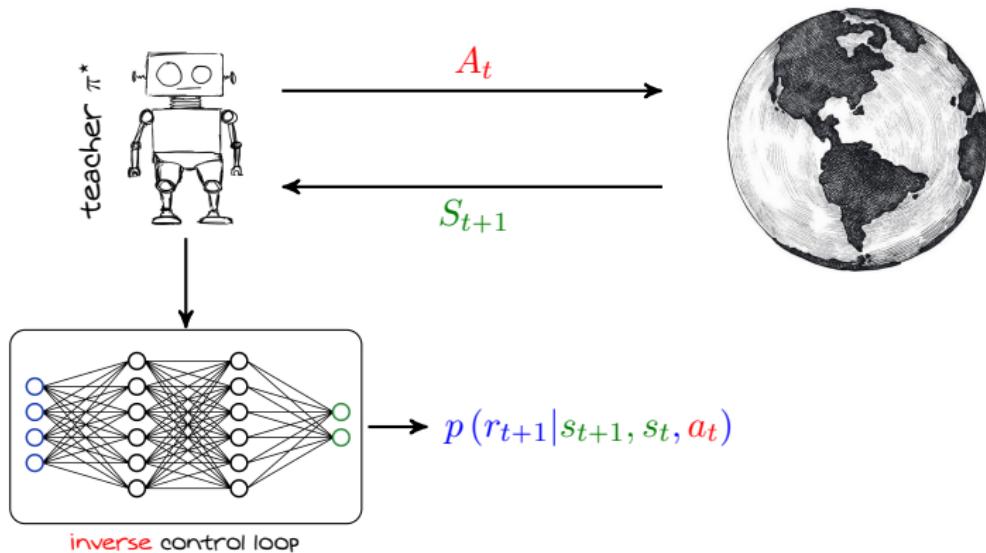
Inverse Reinforcement Learning



In classical RL setting

we collect rewards \rightsquigarrow we try to learn optimal policy

Inverse Reinforcement Learning: Learning Reward Function



In inverse RL setting

a **teacher** plays \rightsquigarrow we try to learn **rewarding system**

Inverse Reinforcement Learning: Some Notes

This problem can be cast using same formulation: we can write the optimal value function again

$$\begin{aligned} v_{\star}(s) &= \mathbb{E}_{\pi^{\star}} \{G_t | S_t = s\} \\ &= \mathbb{E}_{\pi^{\star}} \left\{ \sum_{i=0}^{\infty} \gamma^i R_{t+i+1} | S_t = s \right\} \end{aligned}$$

In RL, we know R_{t+i+1} and look for π^{\star}

- We use **function approximation**, e.g., $\pi^{\star}(\cdot) \equiv f_w(\cdot)$

In inverse RL, we know π^{\star} and look for R_{t+i+1}

- We can again use **function approximation**, e.g., $p(r_{t+1}|\cdot) \equiv f_w(\cdot)$

Pieter Abbeel has a nice lecture on inverse RL: take a look at [this link](#)

Also, you may find [this imitation learning lecture-note](#) useful to learn a bit more!

Using Expert Demonstration

Not all *experts* can formulate their *policy*

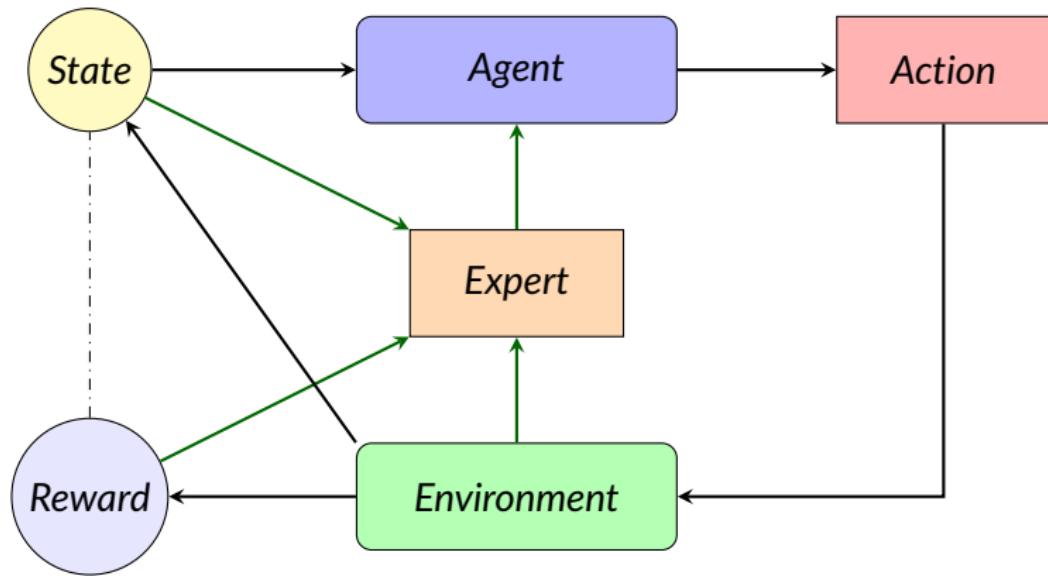
- It might be *not really easy* to formulate it
 - ↳ A board-game player may act based on experience and intuition
- It might be *much easier to demonstrate* the policy
 - ↳ A professional driver could explain what they do in various conditions

In this case learning the *rewarding system* is not really feasible

- We don't exactly know optimal policy
- We can only describe it using *expert demonstrations*

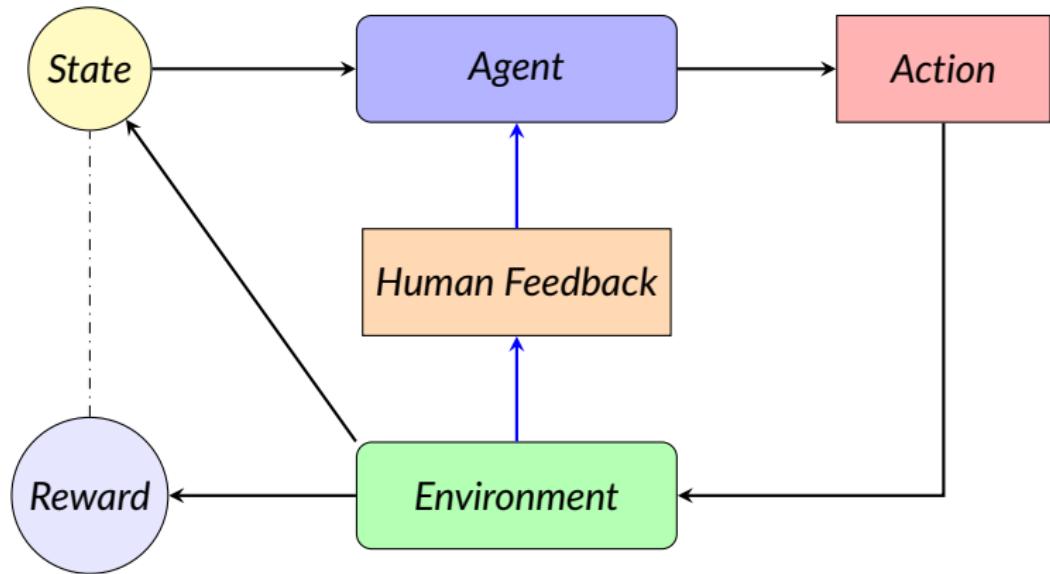
RL via Expert Demonstration

We could include *expert demonstrations* into our control loop



RLHF: RL from Human Feedback

This can be further extended to the use of **human feedback**: that **human** does **not** need to play **optimal!**



ChatGPT also extensively uses RLHF: see [this lecture](#)

Multi-Agent RL



Up to now, we considered other agents as a part of **environment**

- But, they can have their own policies
 - ↳ They could collaborate to play **jointly optimal**
 - ↳ There might be **adversary** agents trying to impact **negatively**

Take a look at [**this manuscript**](#) to learn about multi-agent RL

The End

Many thanks for attending the lectures ...

- Be confident and trust on what you have learned 
- ↳ You are now experts in Deep RL!
- Feel free to reach out
 - ↳ Would be more than happy to help!
 - ↳ I would appreciate your feedback on Quercus 
- Always track the progress in the field
 - ↳ It's fast growing, so you should keep yourself posted 
 - ↳ You are ready to follow any advanced topics in RL 

Next Summer, there will be a **Generative AI course**

- ↳ We may use **some Deep RL** there as well 