

# Reinforcement Learning

## Chapter 4: Function Approximation

Ali Bereyhi

`ali.bereyhi@utoronto.ca`

Department of Electrical and Computer Engineering  
University of Toronto

Fall 2025

## Back to Model-free RL: Control

- + But, we have in general *prediction* and *control* problems. In which one are we going to use function *approximation*?
- Well, we can use in both

### Recall

We have two major problems in *model-free* RL

- *Prediction* in which for a given policy  $\pi$  we evaluate values by sampling the environment
- *Control* in which after each interaction, we improve our policy aiming to converge to the *optimal policy*

Let's now go to the *control*

## Recap: Online Control $\epsilon$ -Greedy

We have seen a typical control loop via  $\epsilon$ -greedy algorithm

```

X_Control():
1: Initiate two random policies  $\pi$  and  $\bar{\pi}$ 
2: while  $\pi \neq \bar{\pi}$  do
3:    $\hat{q}_\pi = \text{X\_QUpdate}(\pi)$  and  $\pi \leftarrow \bar{\pi}$    via function approximation
4:    $\bar{\pi} = \epsilon\text{-Greedy}(\hat{q}_\pi)$ 
5: end while
  
```

We can realize this loop using function approximation

- ① Start with an initial **approximator**
- ② Use the **approximator** to **improve policy** via  $\epsilon$ -Greedy
- ③ Use **SGD** to update the **approximation** model **from observation**


We need an **action-value** approximator in this case

let's formally define the **Q-Network** then

## Q-Net $\equiv$ Action-Value Approximator – Form II

### Q-Network

*In the context of RL, the action-value approximation model that maps features to the vector of all action values is often referred to as Q-Net*




### Deep Q-Network

*If we set Q-Net to be a DNN; then, it is usually called*

*Deep Q-Network  $\equiv$  DQN*

## Example: MLP as DQN

DQN can be simply an MLP



- Give the feature as an **input vector**
- The **MLP** estimates all **action-values**



# Building Control Loop with Q-Net

We can use Q-Net to build a control loop: *similar to tabular RL, we can have on-policy or off-policy approaches*

- Deep on-policy RL
  - ↳ We can use a DQN to realize an on-policy control loop, e.g., SARSA
  - ↳ They are less in use as compared to off-policy versions
- Deep off-policy RL
  - ↳ We may use a DQN to realize an off-policy control loop, e.g., Q-learning
  - ↳ Due to some reasons deep off-policy RL is more popular
    - ↳ We can use experience replay in this case
    - ↳ They are therefore more sample-efficient
    - ↳ No worries! We will see these reasons 😊

Let's start with on-policy algorithms

## Remark: Value Network

- + But, when I look at internet, I usually see the term *value network*! Don't we have this concept?!
- Sure! We have already worked with *value networks*

### Value Network

Any approximation model that gets *features* and returns *values* is a value function; this value could be a *state value* or an *actio-value*

- Q-Net is a *value network*
- $v_{\mathbf{w}}(\cdot)$  that we used for *prediction* was also *value network*

- + Then, do we have any other sort of networks?!
- Yes! We could have *policy networks*: we will see them in the next chapter



## Recap: Going On-Policy

Let's look back at our TD-based prediction algorithm

SGD\_TD\_QEval( $\lambda$ ):

```

1: Initiate with some initial  $\mathbf{w}$  and learning rate  $\alpha$ 
2: for episode  $k = 1 : K$  do
3:   Sample a trajectory  $S_0, A_0 \xrightarrow{R_1} S_1, A_1 \xrightarrow{R_2} \dots \xrightarrow{R_{T-1}} S_{T-1}, A_{T-1} \xrightarrow{R_T} S_T$ 
4:   for  $t = 0 : T - 1$  do
5:     Compute  $\Delta_t = R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - Q_{\mathbf{w}}(S_t, A_t)$       # forward propagation
6:     Compute  $\nabla_t = \nabla Q_{\mathbf{w}}(S_t, A_t)$                                 # backpropagation
7:      $E_{\mathbf{w}} \leftarrow \lambda \gamma E_{\mathbf{w}} + \nabla_t$ 
8:     Update weights as
                                      $\mathbf{w} \leftarrow \mathbf{w} + \alpha \Delta_t E_{\mathbf{w}}$ 
9:   end for
10: end for

```

The key point in going on policy is to evaluate  $v_{\mathbf{w}}(S_{t+1})$  using our *actual policy*

## SARSA: Going On-Policy

We can modify line 5: we compute  $v_{\mathbf{w}}(S_{t+1})$  as

$$v_{\mathbf{w}}(S_{t+1}) = \sum_{m=1}^M \pi(a^m | S_{t+1}) Q_{\mathbf{w}}(S_{t+1}, a^m)$$

In *on-policy* approach, we *act before we update*

our policy leads us to next action  $A_{t+1}$

So, we could *move on our policy* and write

$$\pi(a | S_{t+1}) = \begin{cases} 1 & a = A_{t+1} \\ 0 & a \neq A_{t+1} \end{cases} \rightsquigarrow v_{\mathbf{w}}(S_{t+1}) = Q_{\mathbf{w}}(S_{t+1}, A_{t+1})$$

# SARSA with Q-Net

SGD\_SARSA() :

```

1: Initiate with  $\mathbf{w}$  and learning rate  $\alpha$ 
2: for  $\text{episode} = 1 : K$  or until  $\pi$  stops changing do
3:   Initiate with a random state-action pair  $(S_0, A_0)$ 
4:   for  $t = 0 : T - 1$  where  $S_T$  is either terminal or terminated do
5:     Act  $A_t$  and observe  $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$ 
6:     Update policy to  $\pi \leftarrow \epsilon\text{-Greedy}(Q_{\mathbf{w}}(S_t, A_t))$ 
7:     Draw the new action  $A_{t+1}$  from  $\pi(\cdot | S_{t+1})$  and move on policy

            $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1}$ 

8:     Set  $\Delta \leftarrow R_{t+1} + \gamma Q_{\mathbf{w}}(S_{t+1}, A_{t+1}) - Q_{\mathbf{w}}(S_t, A_t)$       # forward propagation
9:     Update  $\mathbf{w} \leftarrow \mathbf{w} + \alpha \Delta \nabla Q_{\mathbf{w}}(S_t, A_t)$                   # backpropagation
10:   end for
11: end for

```

## SARSA with Q-Net and Eligibility Tracing

We have seen that with *function approximation eligibility tracing* reduces to

$$E_{\mathbf{w}} \leftarrow \gamma \lambda E_{\mathbf{w}} + \nabla Q_{\mathbf{w}}(S_t, A_t)$$

Let's fit it into our SARSA control loop!

### Eligibility Tracing $\propto$ Backpropagation

If we use a DQN, we should *backpropagate* to compute the *eligibility tracing*: this point is *intuitive* as both approaches naturally follow the same logic

# SARSA( $\lambda$ ) with Function Approximation

SGD\_SARSA( $\lambda$ ) :

- 1: Initiate with  $\mathbf{w}$  and learning rate  $\alpha$
- 2: **for**  $\text{episode} = 1 : K$  or until  $\pi$  stops changing **do**
- 3:   Initiate with a random state-action pair  $(S_0, A_0)$
- 4:   **for**  $t = 0 : T - 1$  where  $S_T$  is either terminal or terminated **do**
- 5:     Act  $A_t$  and observe  $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$
- 6:     Update policy to  $\pi \leftarrow \epsilon\text{-Greedy}(Q_{\mathbf{w}}(S_t, \mathbf{a}))$
- 7:     Draw the new action  $A_{t+1}$  from  $\pi(\cdot | S_{t+1})$  and move on policy
- $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1}$
- 8:     Set  $\Delta \leftarrow R_{t+1} + \gamma Q_{\mathbf{w}}(S_{t+1}, A_{t+1}) - Q_{\mathbf{w}}(S_t, A_t)$    # forward propagation
- 9:      $E_{\mathbf{w}} \leftarrow \lambda \gamma E_{\mathbf{w}} + \nabla Q_{\mathbf{w}}(S_t, A_t)$    # backpropagation
- 10:     Update  $\mathbf{w} \leftarrow \mathbf{w} + \alpha \Delta E_{\mathbf{w}}$
- 11:   **end for**
- 12: **end for**

## Recap: Going Off-Policy

In off-policy control: we *behave* with a policy  $\pi$  but *update* by a *target policy*  $\bar{\pi}$

- We could use *importance sampling* to evaluate *target policy*
- We mainly focused on *Q-learning* approach

### Q-Learning

Q-learning is an *off-policy* TD *control* algorithm, where we sample with  $\epsilon$ -*greedy policy* but *update* with *greedy policy*

Key property of *Q-learning* is that we don't really need *importance sampling*: we could directly *update* as

$$\hat{q}_{\bar{\pi}}(S_t, A_t) \leftarrow \hat{q}_{\bar{\pi}}(S_t, A_t) + \alpha \left( R_{t+1} + \gamma \max_m \hat{q}_{\bar{\pi}}(S_{t+1}, a^m) - \hat{q}_{\bar{\pi}}(S_t, A_t) \right)$$

## Q-Learning with Q-Net

We can therefore extend out Q-learning algorithm to

SGD\_Q-Learning():

```

1: Initiate with  $\mathbf{w}$  and learning rate  $\alpha$ 
2: for episode = 1 :  $K$  or until  $\pi$  stops changing do
3:   Initiate with a random state  $S_0$ 
4:   for  $t = 0 : T - 1$  where  $S_T$  is either terminal or terminated do
5:     Update policy to  $\pi \leftarrow \epsilon\text{-Greedy}(Q_{\mathbf{w}}(S_t, \mathbf{a}))$ 
6:     Draw action  $A_t$  from  $\pi(\cdot | S_t)$  and observe  $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}$ 
7:      $\Delta \leftarrow R_{t+1} + \gamma \max_m Q_{\mathbf{w}}(S_{t+1}, a^m) - Q_{\mathbf{w}}(S_t, A_t)$  # forward propagation
8:     Update  $\mathbf{w} \leftarrow \mathbf{w} + \alpha \Delta \nabla Q_{\mathbf{w}}(S_t, A_t)$  # backpropagation
9:   end for
10: end for

```

We could potentially use *eligibility tracing* as well!

# Incremental Algorithms: Challenges

What we have developed by now are the so-called *incremental approaches*

## Incremental Algorithms

*Incremental algorithms* use the *actual sample* at each time to *update* the Q-Net

Though easy to implement, *incremental approaches* are *not* efficient

① They are extremely *sample-inefficient*

- ↳ Once we use a sample, we are *over* with it
- ↳ But, in deep learning we used to use the same samples *several times*
  - ↳ We make a training loop with *multiple epochs*
  - ↳ In each epoch, we go through the *whole dataset*




# Incremental Algorithms: Challenges

Though easy to implement, *incremental approaches* are *not* efficient

② They use samples that are *strongly correlated*

- ↳ In sample  $S_t, A_t \xrightarrow{R_{t+1}} S_{t+1}, A_{t+1}$  states  $S_t$  and  $S_{t+1}$  are closely related
  - ↳ If state is the location; then, locations in time  $t$  and  $t + 1$  are very close
- ↳ But, we know that we need *independent* samples
  - ↳ With correlated samples we can easily *stick to a local fitting*
  - ↳ Our Q-Net *does not generalize* very well



## Solution: *Batch Training*

The solution to these challenges is to use *batch training*

- 1 Save every sample in a *database*
  - ↳ We can only save the values we need, i.e., estimator values
- 2 In each iteration *sample* from *database*
  - ↳ We can reuse (reply) our *sampled experiences*
  - ↳ We can *reduce* the *correlation* between succeeding samples

This is what we call *experience reply*

- This needs us to control *off-policy*
  - ↳ We are using samples of *previous non-improved policies*
  - ↳ We want to update the value of a *different target policy*
- This is why deep *off-policy* algorithms are *more sample-efficient*

We next study this in details 😊